

# Few-Shot Learning on Graphs: From Meta-Learning to LLM-empowered Pre-Training and Beyond

Lecture-style Tutorial

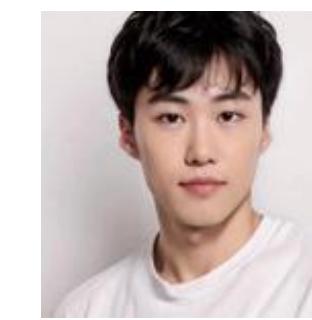
The Web Conference 2025



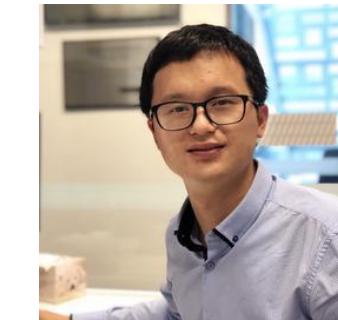
**Yuan Fang**



**Yuxia Wu**



**Xingtong Yu**



**Shirui Pan**

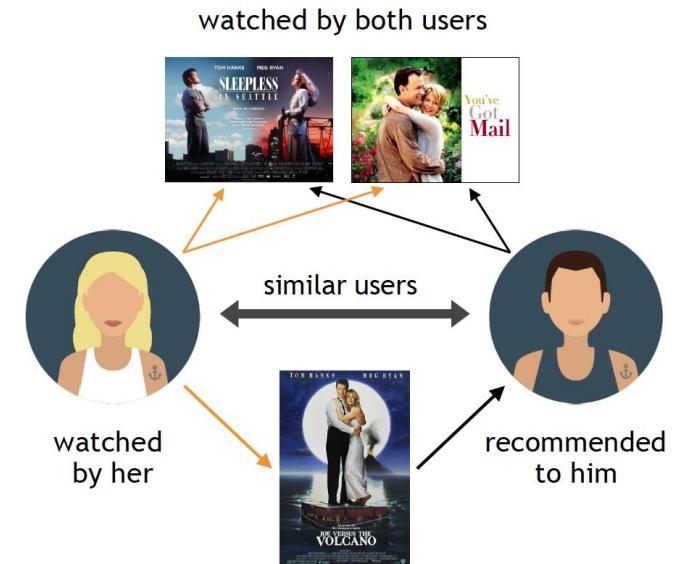


# Graphs

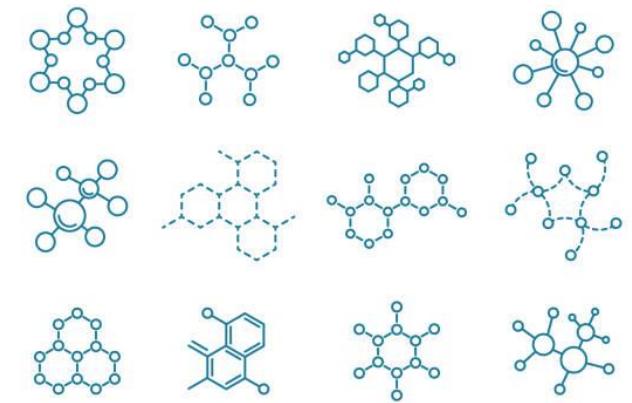
- Graphs model the interactions among various objects



Social network



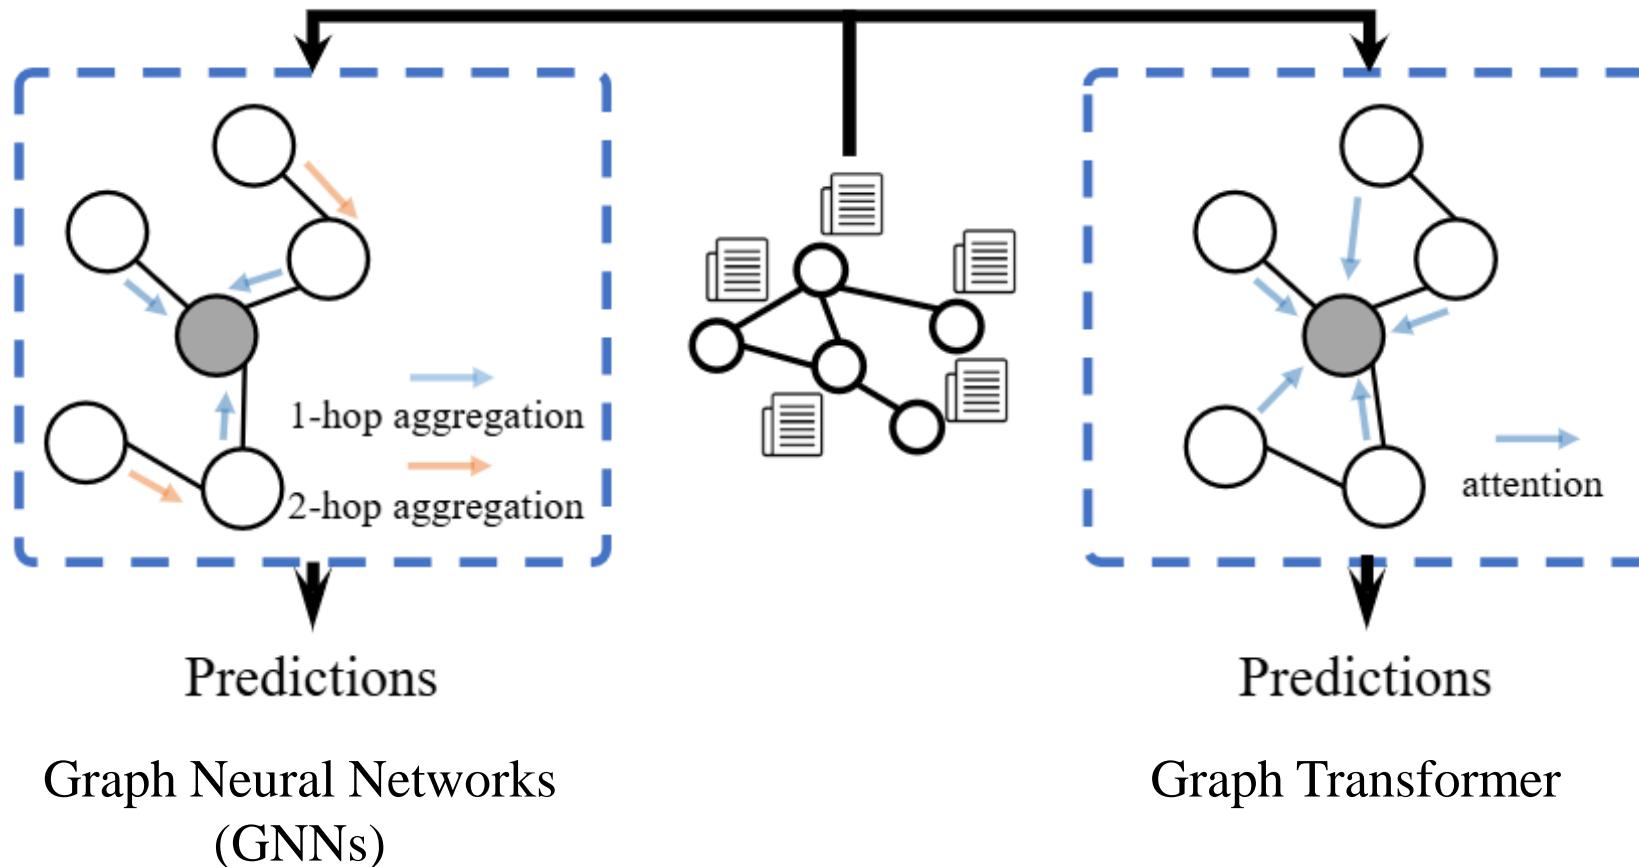
Recommendation System



Molecular graphs

# End to End Graph Learning

- (Semi-)Supervised graph representation learning methods

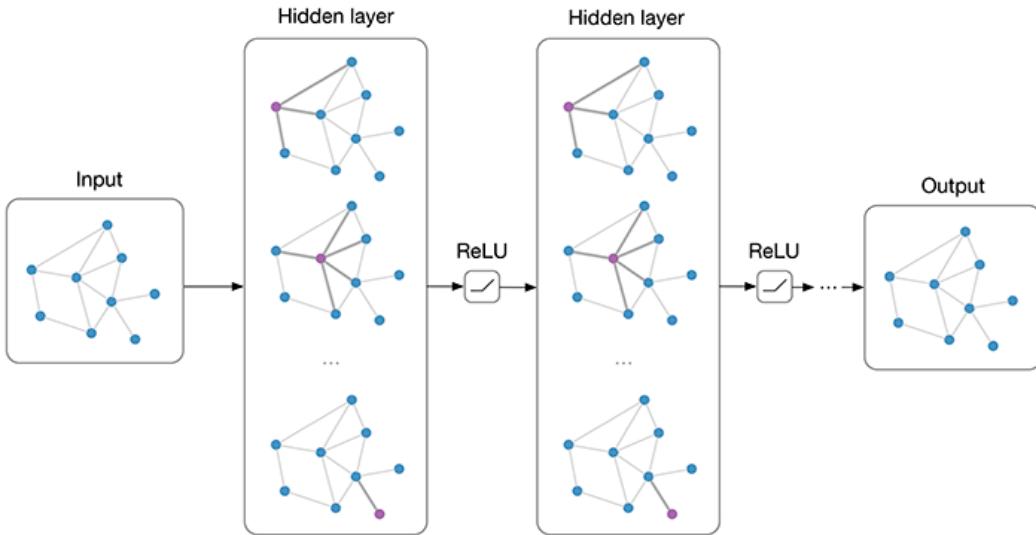


Graph Neural Networks  
(GNNs)

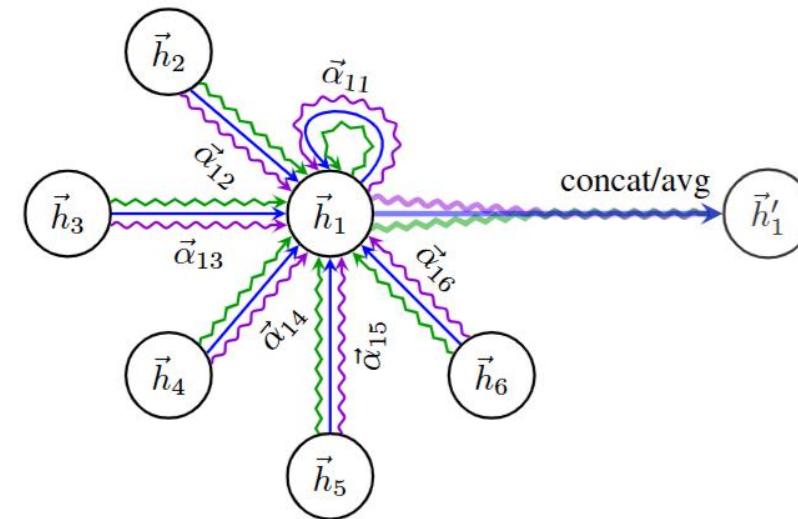
Graph Transformer

# Graph Neural Networks

- GNNs typically leverage message-passing framework



GCN



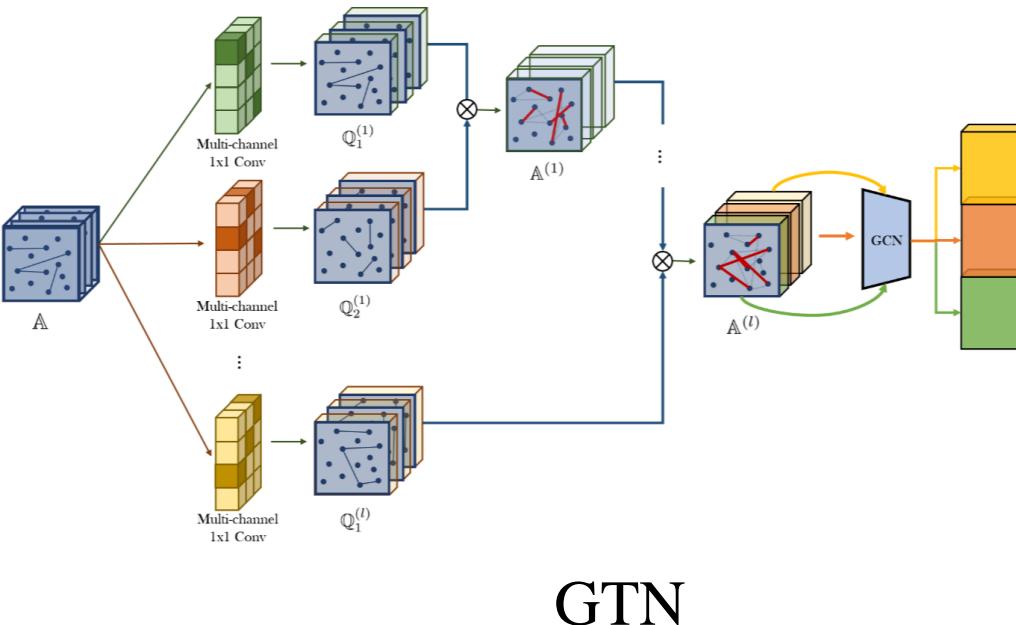
GAT

Kipf, et al. "Semi-supervised classification with graph convolutional networks." ICLR'17.

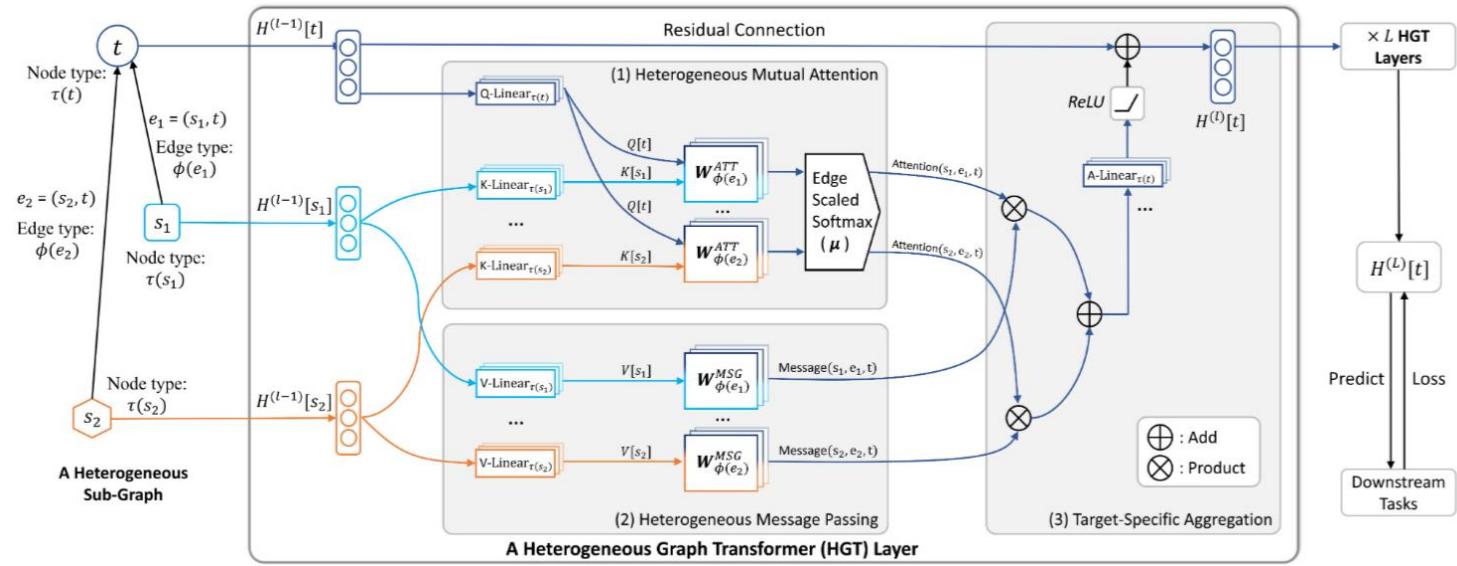
Velickovic, et al. "Graph attention networks". ICLR'18.

# Graph Transformers

- Transformers are widely used in graph learning



GTN



HGT

Yun, et al. "Graph transformer networks." NeurIPS'19.

Hu, et al. "Heterogeneous graph transformer." WWW'20.

# Few-shot Learning Problems

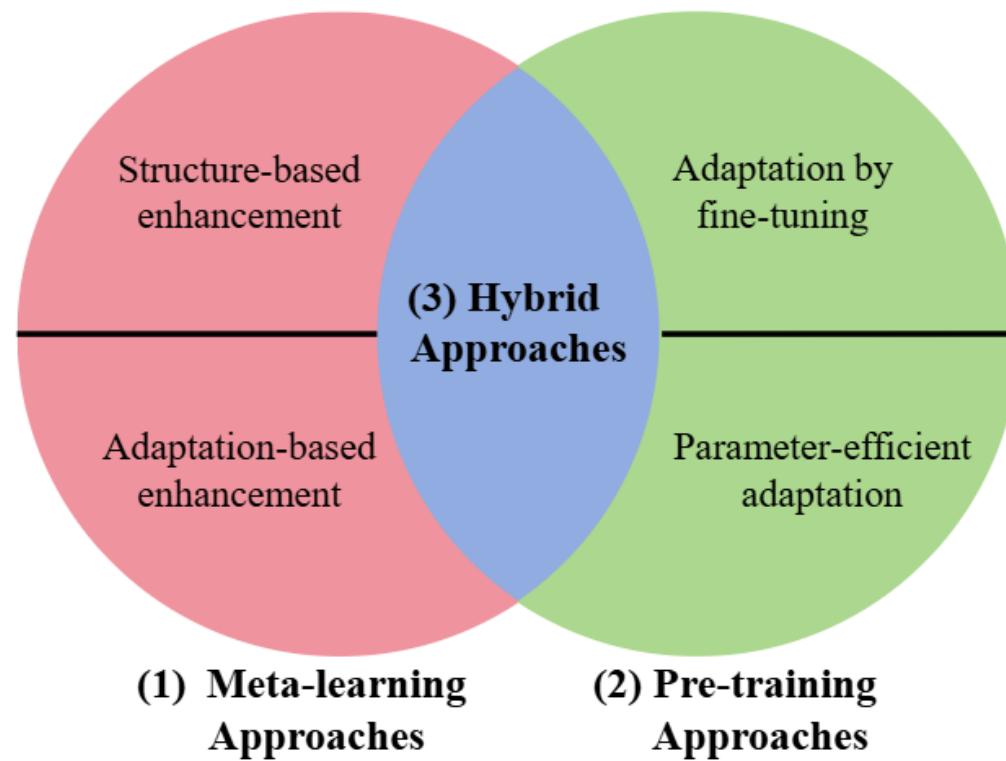
- Performance highly depends on
  - Abundant labeled data
  - Rich Structure
  - Challenging or expensive to obtain labels, leads to
  - Graph structure may be sparse, leads to

**Label scarcity**

**Structure scarcity**

# Few-shot Learning Methods

- Learn prior knowledge and adapt to downstream applications



# Few-shot Learning Methods

- **Meta learning methods**
  - Derive prior knowledge from a series of “meta-training” task
- **Pre-training methods**
  - Utilize unlabeled data to optimize self-supervised pretext tasks
  - Employ fine-tuning or parameter-efficient adaption
- **Hybrid methods**
  - Integrate both paradigms

# Outline



**Shirui Pan** Griffith University  
09:00-09:10 Introduction (10mins)



**Yuan Fang** Singapore Management University  
09:10-09:40 Problem and Applications (30mins)



**Yuxia Wu** Singapore Management University  
09:40-10:05 Meta-Learning Approaches (25mins)



**Xingtong Yu** Singapore Management University  
10:05-10:30 Pre-LLM (25mins)



10:30-11:00 Break (30mins)

**Yuxia Wu** Singapore Management University  
10:30-11:00 LLM Era (30mins)

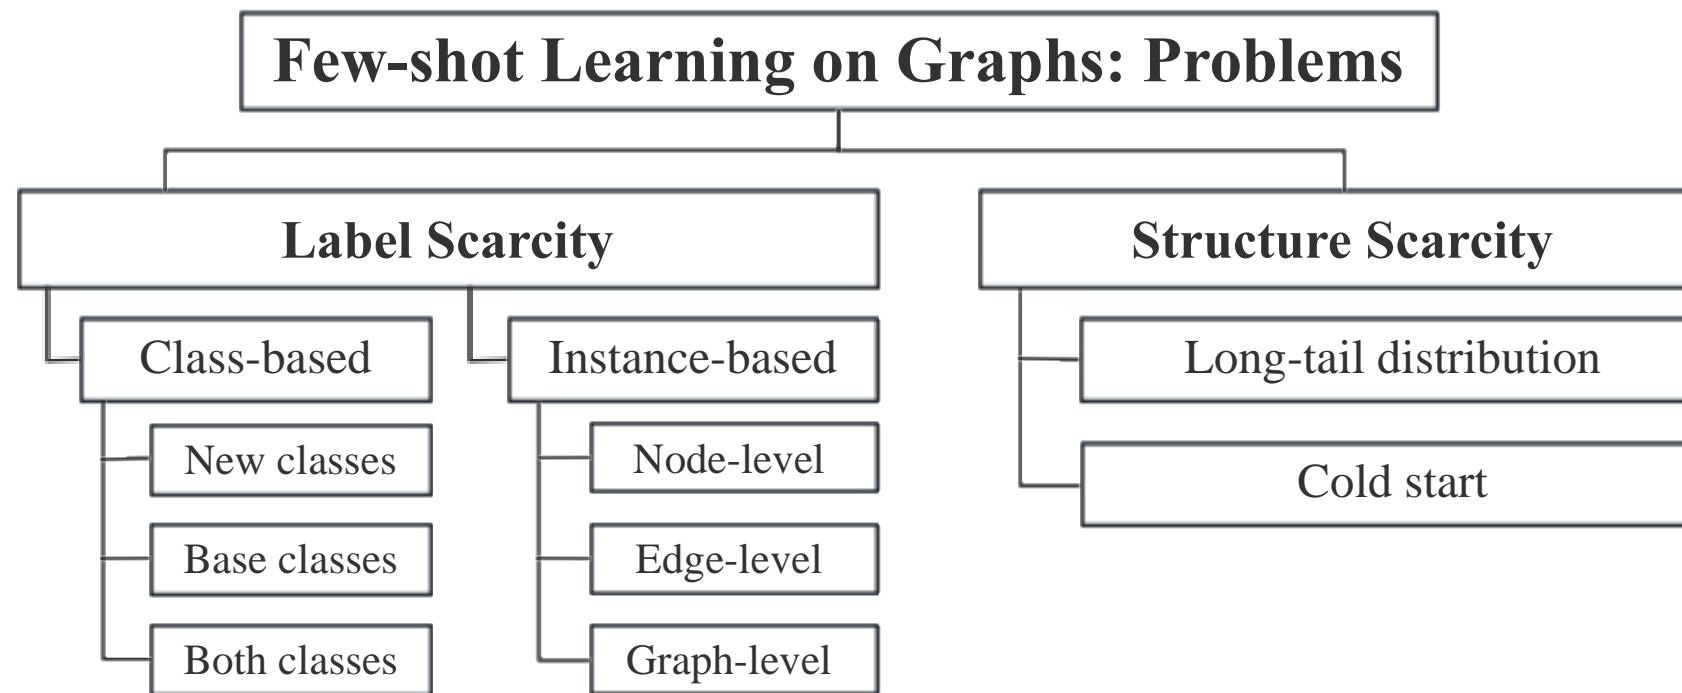
**Xingtong Yu** Singapore Management University  
11:00-11:20 Hybrid Approaches (20mins)

**Yuan Fang** Singapore Management University  
11:20-11:40 Problem and Applications (20mins)

**Host: Yuan Fang**  
11:40-12:30 Panel (50mins)

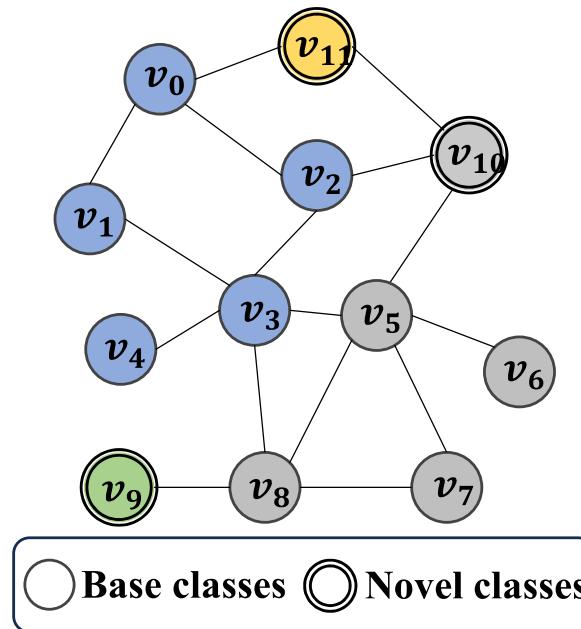
# Few-Shot Learning Problems on Graphs

- Label scarcity: the lack of labeled data
- Structure scarcity:



# Label Scarcity Problems on Graphs

- Class-based Label Scarcity



The entire set of classes ( $C$ ) on a graph

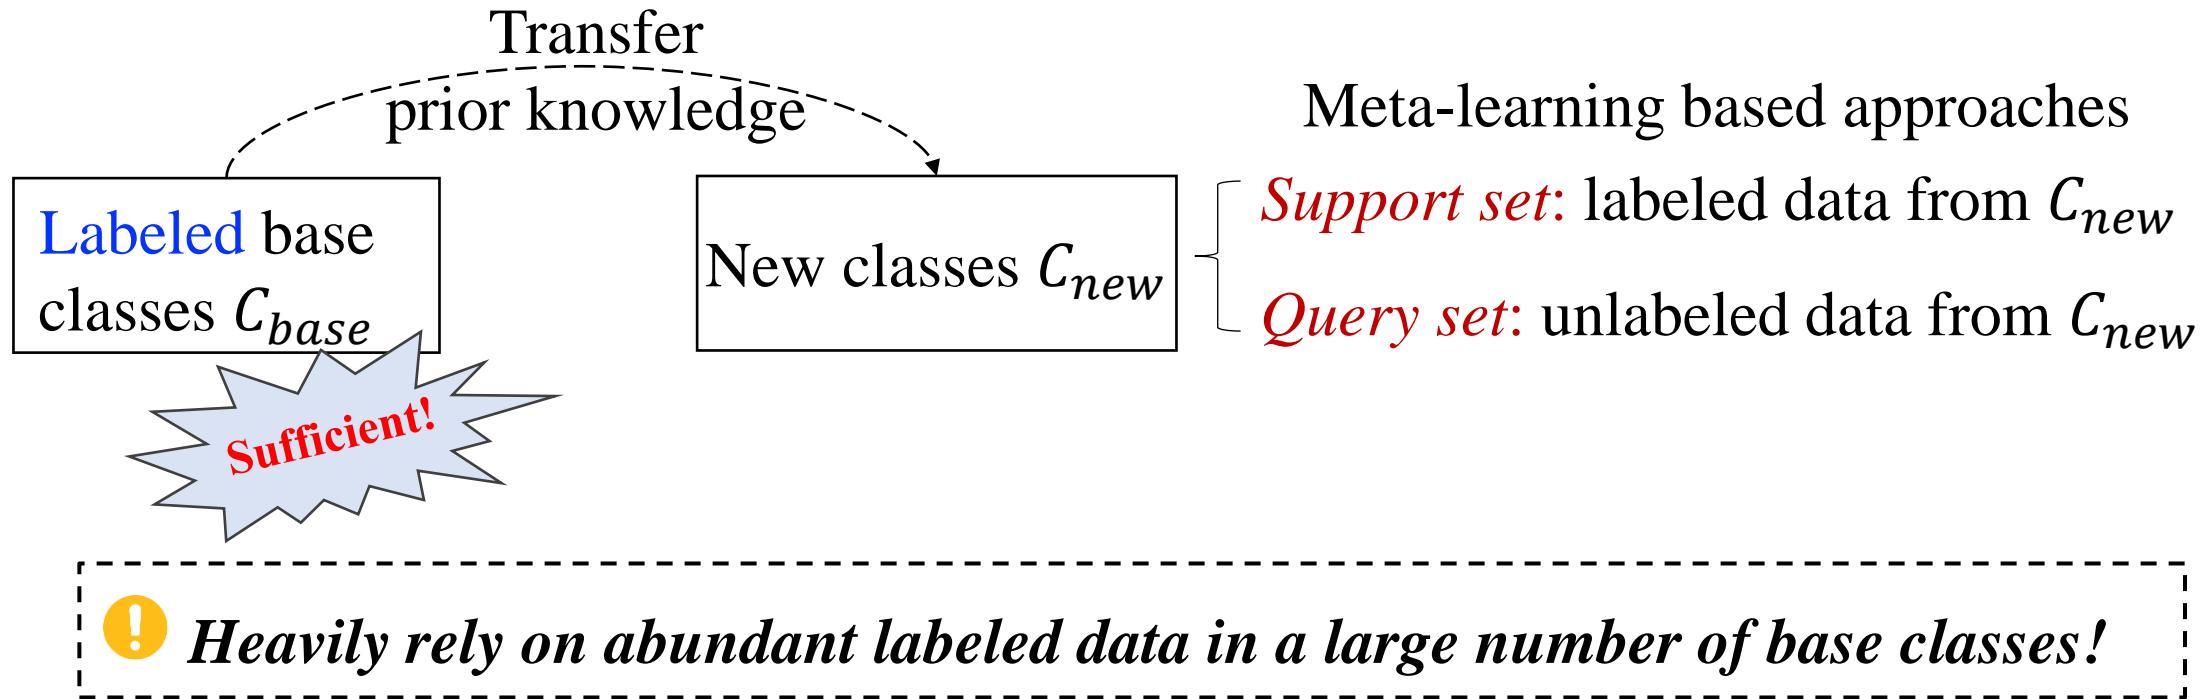
- Base class set  $C_{base}$  for model training
- New class set  $C_{new}$  for testing
- $C = C_{base} \cup C_{new}$
- $C_{base} \cap C_{new} = \emptyset$

Label scarcity could happen in either subsets or both

# Label Scarcity Problems on Graphs

- Class-based Label Scarcity

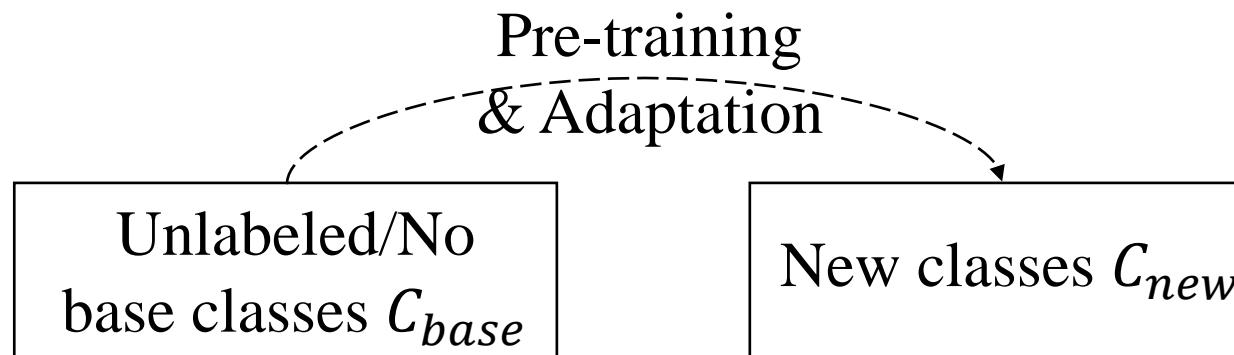
➤ Label scarcity in new classes  $C_{new}$



# Label Scarcity Problems on Graphs

- Class-based Label Scarcity

➤ Label scarcity in base classes  $C_{base}$



- ✗ **Meta-learning based methods**
- ✓ **Self-supervised methods**
  - ❑ Pre-train graph encoders on  $C_{base}$
  - ❑ Fine-tuning on novel tasks

# Label Scarcity Problems on Graphs

- Class-based Label Scarcity

➤ Label scarcity in both classes: labeled data are limited in both  $C_{base}$  and  $C_{new}$

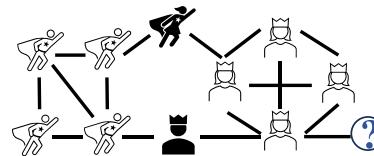
## Self-supervised methods

- **Pre-train** graph encoders
- **Fine-tuning** on novel downstream tasks: **parameter-efficient** adaptation

# Label Scarcity Problems on Graphs

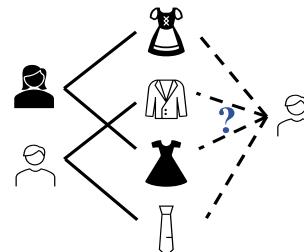
- Instance-based Label Scarcity

➤ Node-level label scarcity



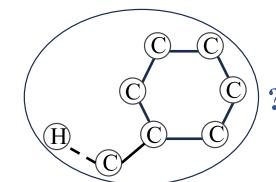
Social network

➤ Edge-level label scarcity



Recommender system

➤ Graph-level label scarcity



Molecular graph

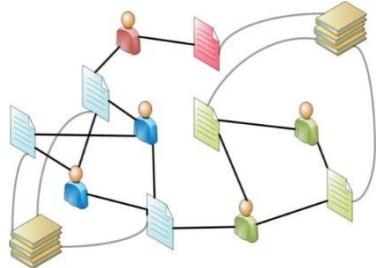
Instance	Application domain
Node	Academic network
	Social network
	E-commerce network
	Protein-protein interaction
	Traffic flow
Edge	Drug-drug interaction
	Protein multimer structure
	E-comm./academic network
	Knowledge graphs
Graph	Molecular graph
	Protein graph
	Social network

# Label Scarcity Problems on Graphs

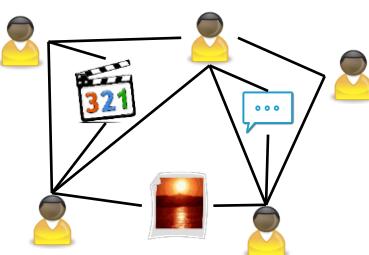
- Instance-based Label Scarcity

- Node-level label scarcity

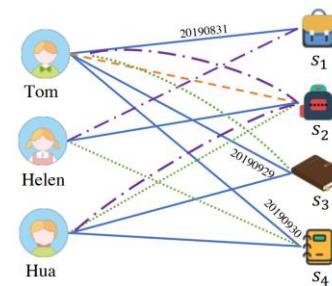
Academic network



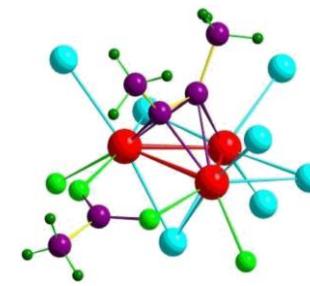
Social network



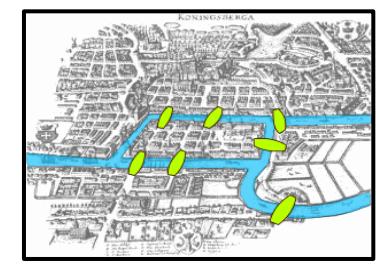
Recommender system



Molecular graph



Traffic Flow

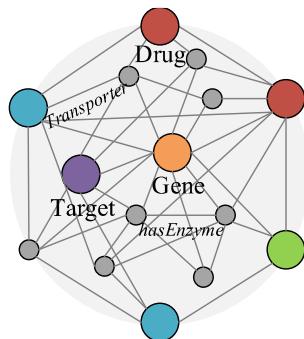


# Label Scarcity Problems on Graphs

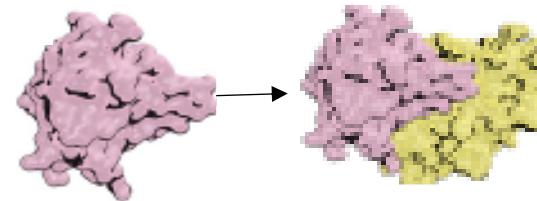
- Instance-based Label Scarcity

- Edge-level label scarcity

Drug-drug interaction



Multimer structure prediction

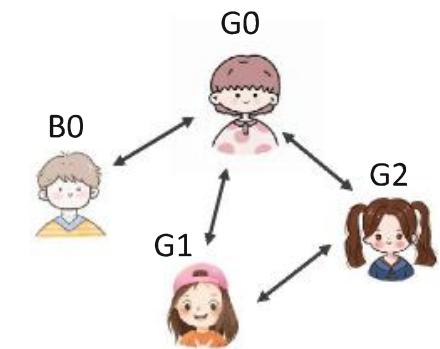


Knowledge graph



[Image from Microsoft]

E-commerce



Product sharing

Baek, et al. "Learning to extrapolate knowledge: Transductive few-shot out-of-graph link prediction ." NeurIPS'20

Gao, et al. "Protein multimer structure prediction via prompt learning." ICLR'24

Zhu, et al. "Few-shot link prediction for event-based social networks via meta- learning." DASFAA'23

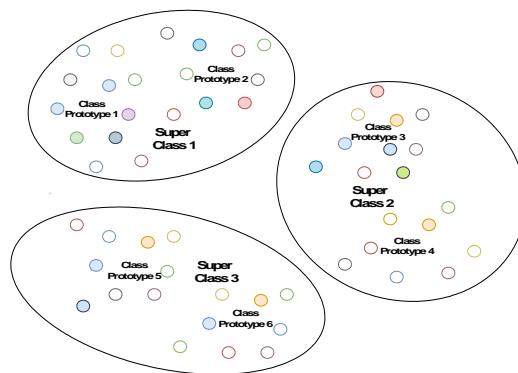
# Label Scarcity Problems on Graphs

- Instance-based Label Scarcity

- Graph-level label scarcity:

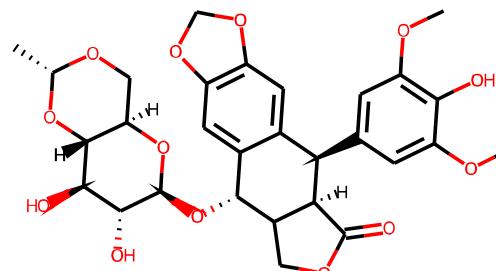
Predicting properties/categories for subgraphs/whole graphs with limited labeled data

Social Network



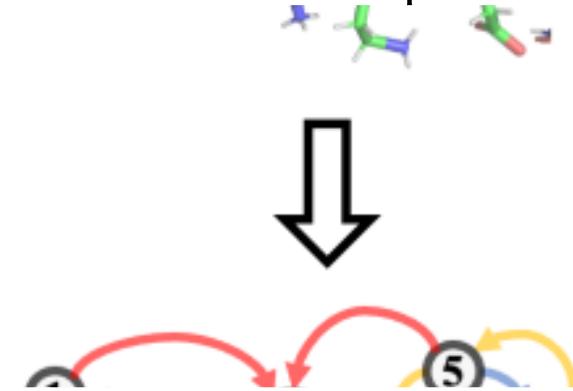
User communities classification

Molecular Graph



Property prediction

Protein Graph



Property prediction

Chauhan, et al. "Few-shot learning on graphs via super-classes based on graph spectral measures." ICLR'20

Zhu, et al. "Dual-view Molecular Pre-training." KDD'23

Zhang, et al. "Protein Representation Learning by Geometric Structure Pretraining." ICLR'23

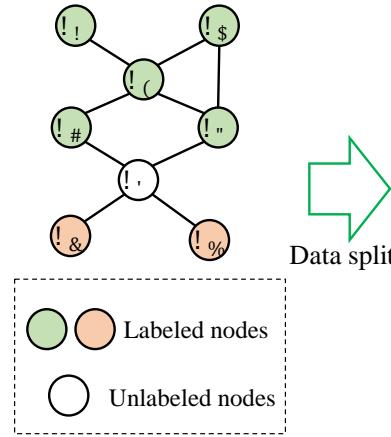
# Structure Scarcity Problems on Graphs

- Long-tailed distribution
  - learn from an imbalanced distribution : a large number of nodes have few connections
- Cold-start
  - Learn representations for new nodes with no or very few connections

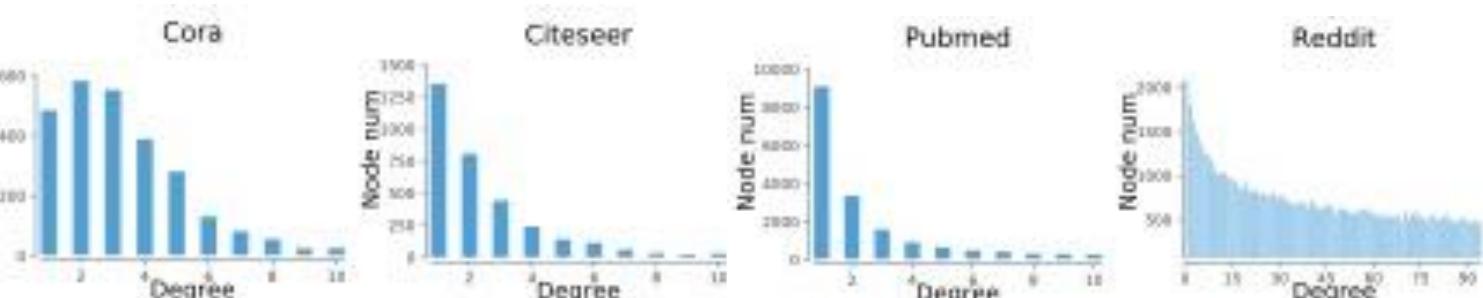
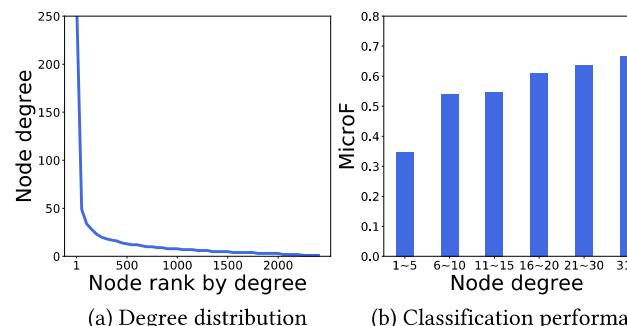
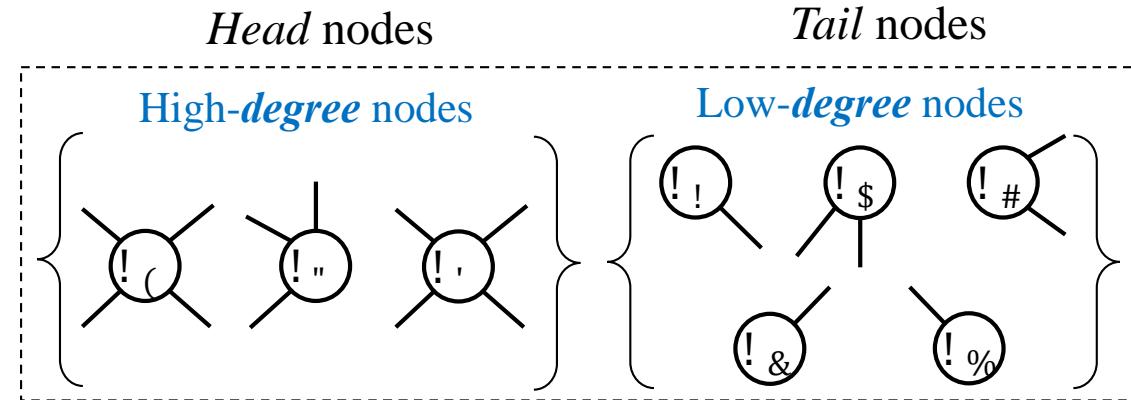
Goal	Application domain
Long-tailed distribution	Academic network
	Social network
	E-commerce network
	Protein-protein interaction
	Air traffic control
Cold start	Social network
	E-commerce network

# Structure Scarcity Problems on Graphs

- Long-tailed distribution



*Example:*  
Node representation learning with different degrees  
(High:  $\text{degree} \geq 3$ )



Liu, et al. "A Survey of Imbalanced Learning on Graphs: Problems, Techniques, and Future Directions." arXiv'23

Tang, et al. "Investigating and Mitigating Degree-Related Biases in Graph Convolutional Networks." CIKM'20

# Structure Scarcity Problems on Graphs

- Cold-start learning: new nodes with few connections

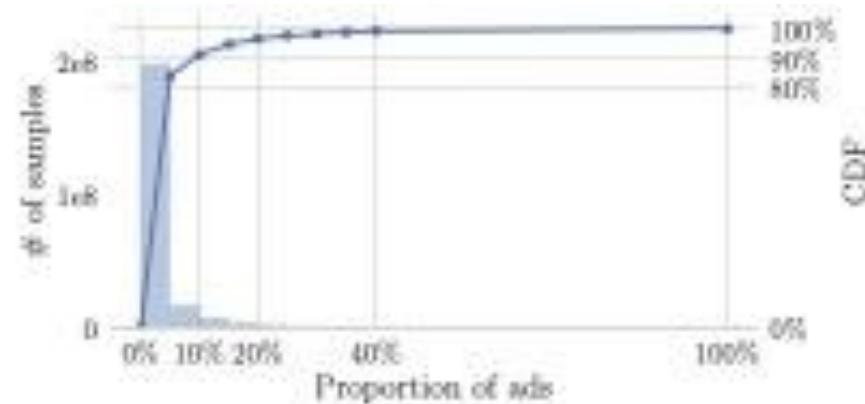


Figure 1: Histogram of the number of samples over different proportions of ads of the KDD Cup 2012 search ads dataset.

5% of ads accounted for over 80% of samples;  
95% ads had a very small amount of data.

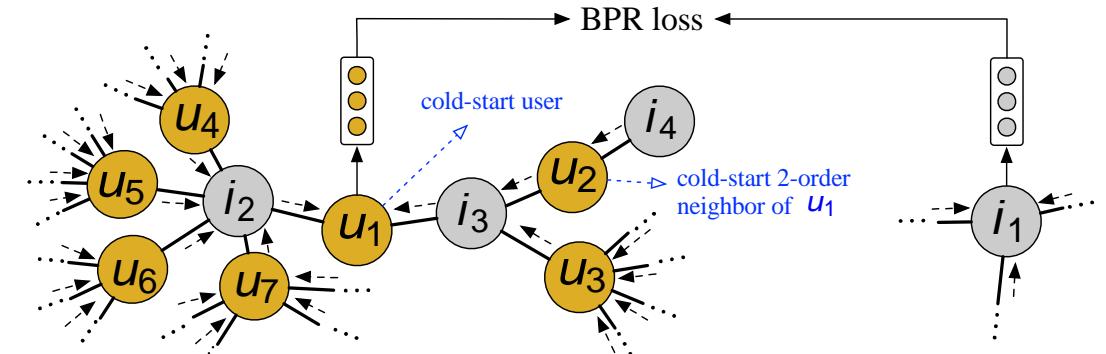


Figure 1: A GNN model for recommendation.

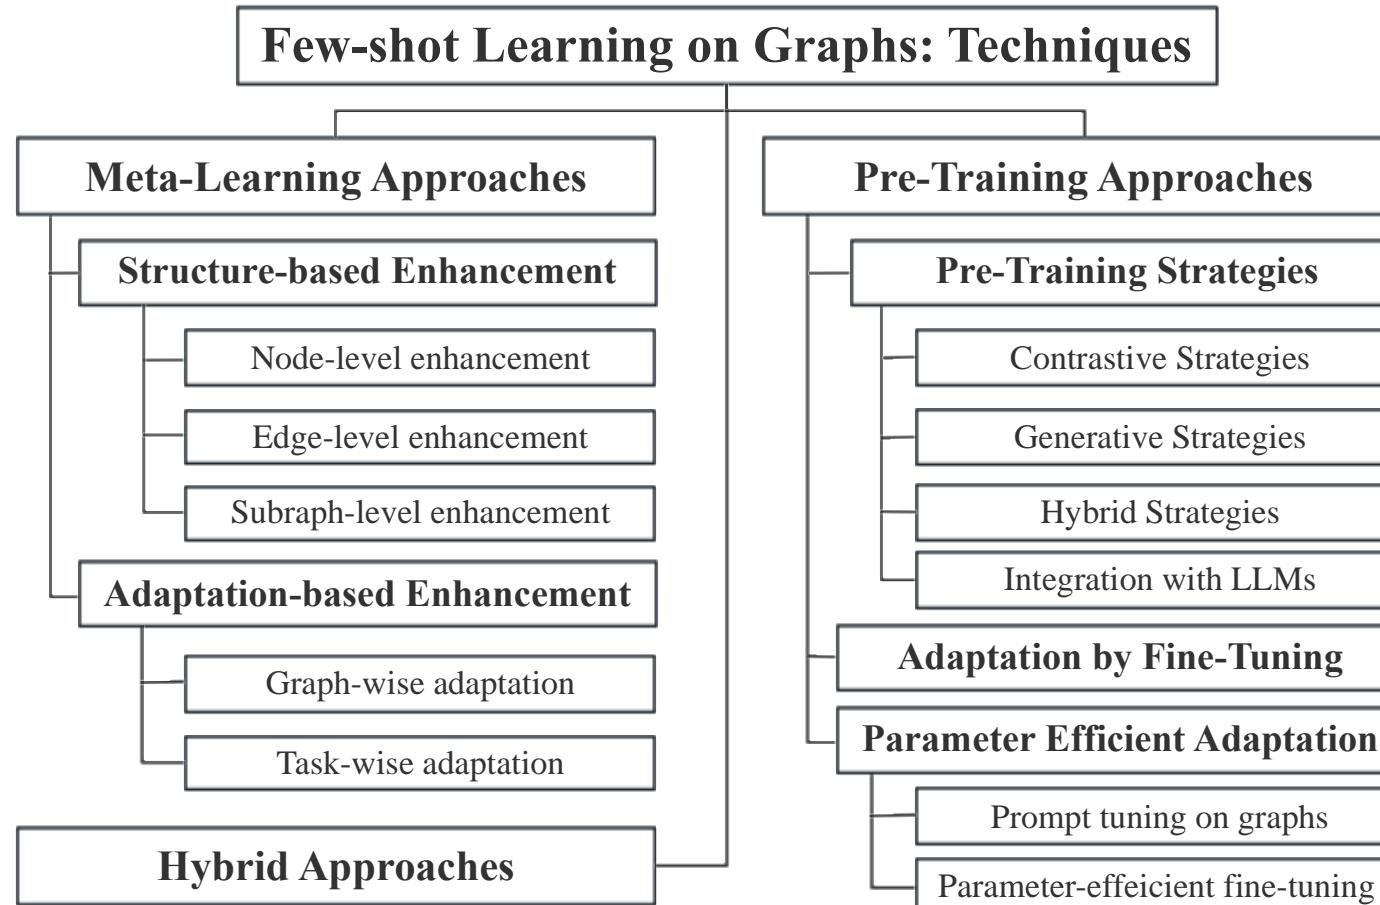
Classic GNNs may have limited effectiveness in addressing cold-start problems

Pan, et al. "Warm Up Cold-start Advertisements- Improving CTR Predictions via Learning to Learn ID Embeddings." SIGIR'19

Hao, et al. "Pre-Training Graph Neural Networks for Cold-Start Users and Items Representation." WSDM'21

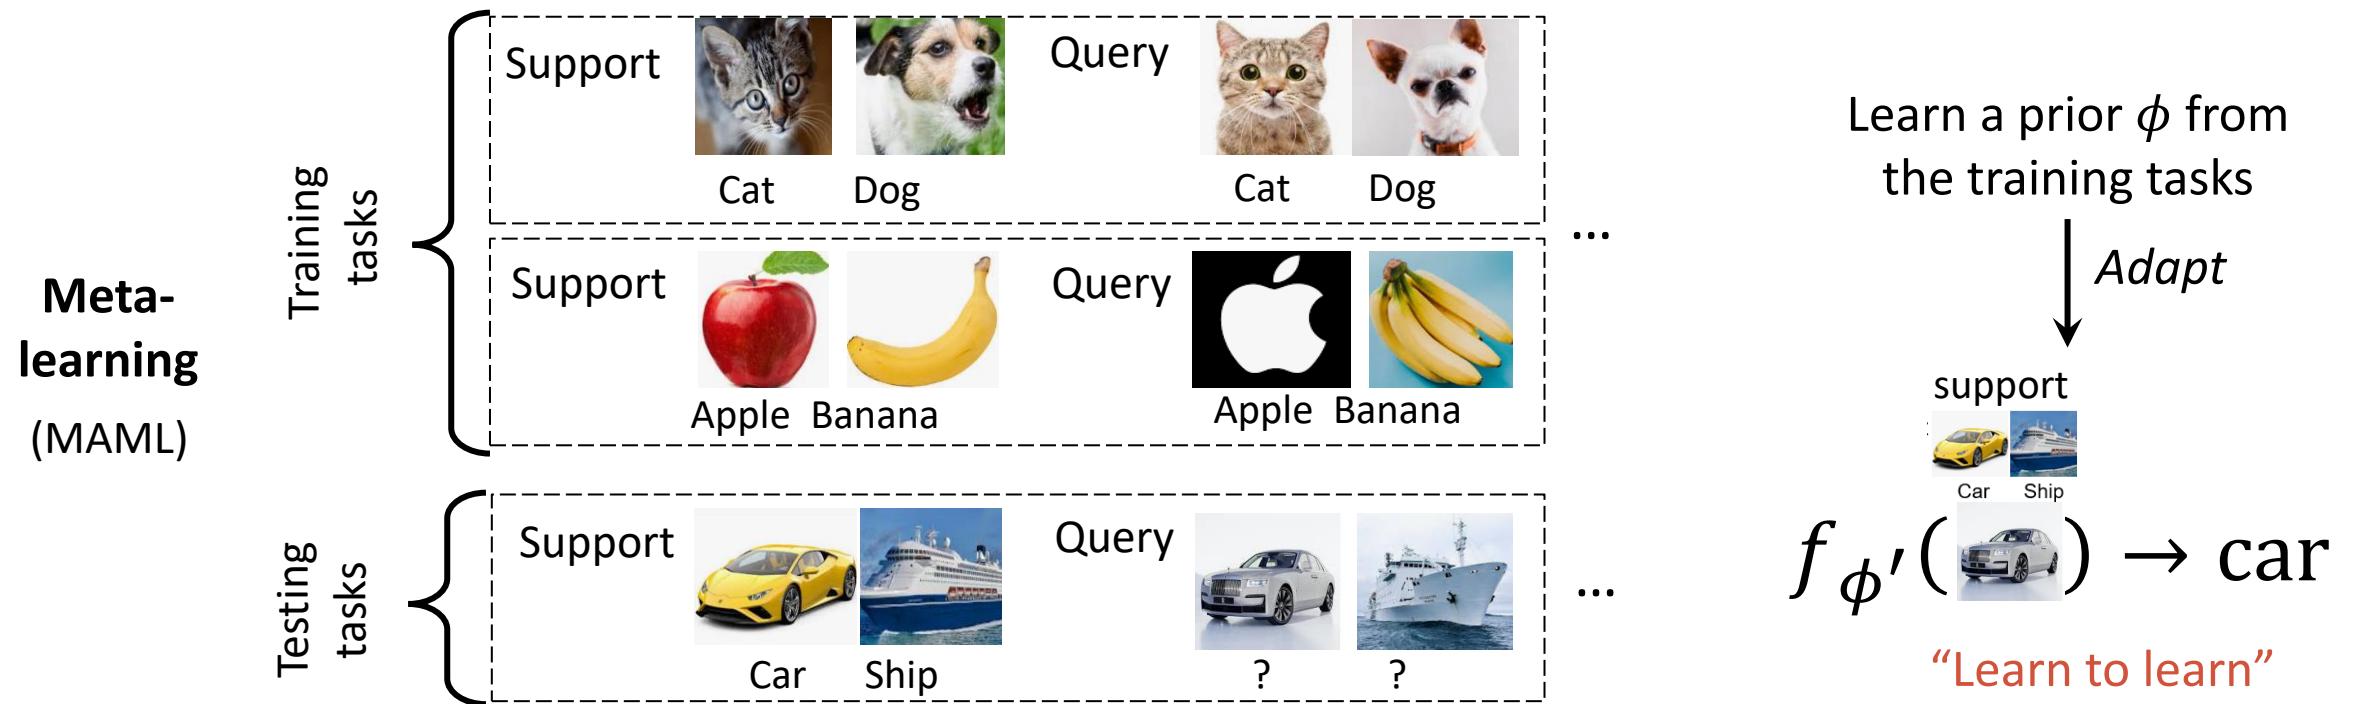
# Overall Taxonomy

- Taxonomy of few-shot learning techniques on graphs



# Meta-learning techniques on graphs

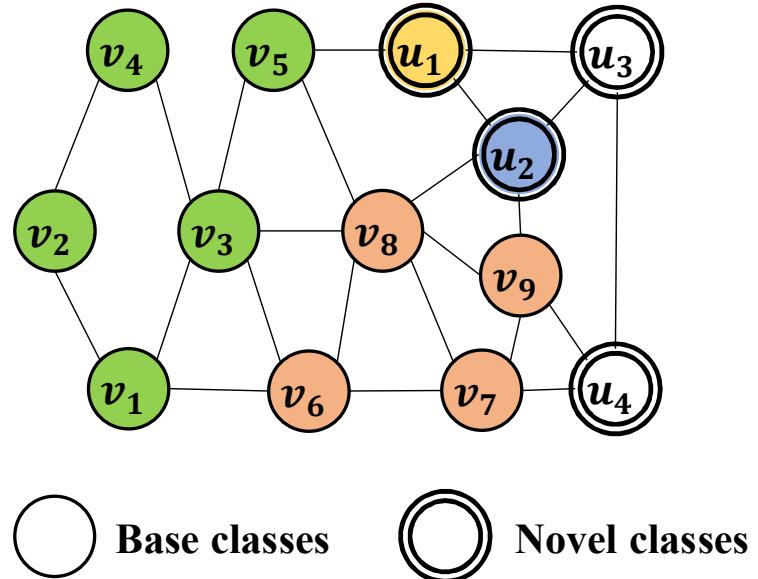
- Standard meta-learning techniques



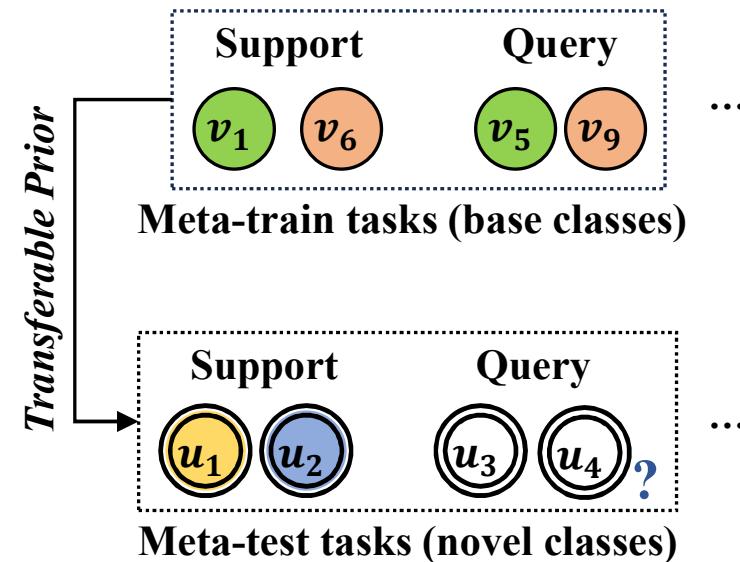
C. Finn et al. “Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks.” ICML 2017.

# Meta-learning techniques on graphs

- Standard meta-learning on graph



(a) Toy graph with base and novel classes



(b) Few-shot node classification

# Structure-based Enhancement on Graphs

- Node-level enhancement: **GPN**

Differentiating node weights in a task to reflect their varying structural importance

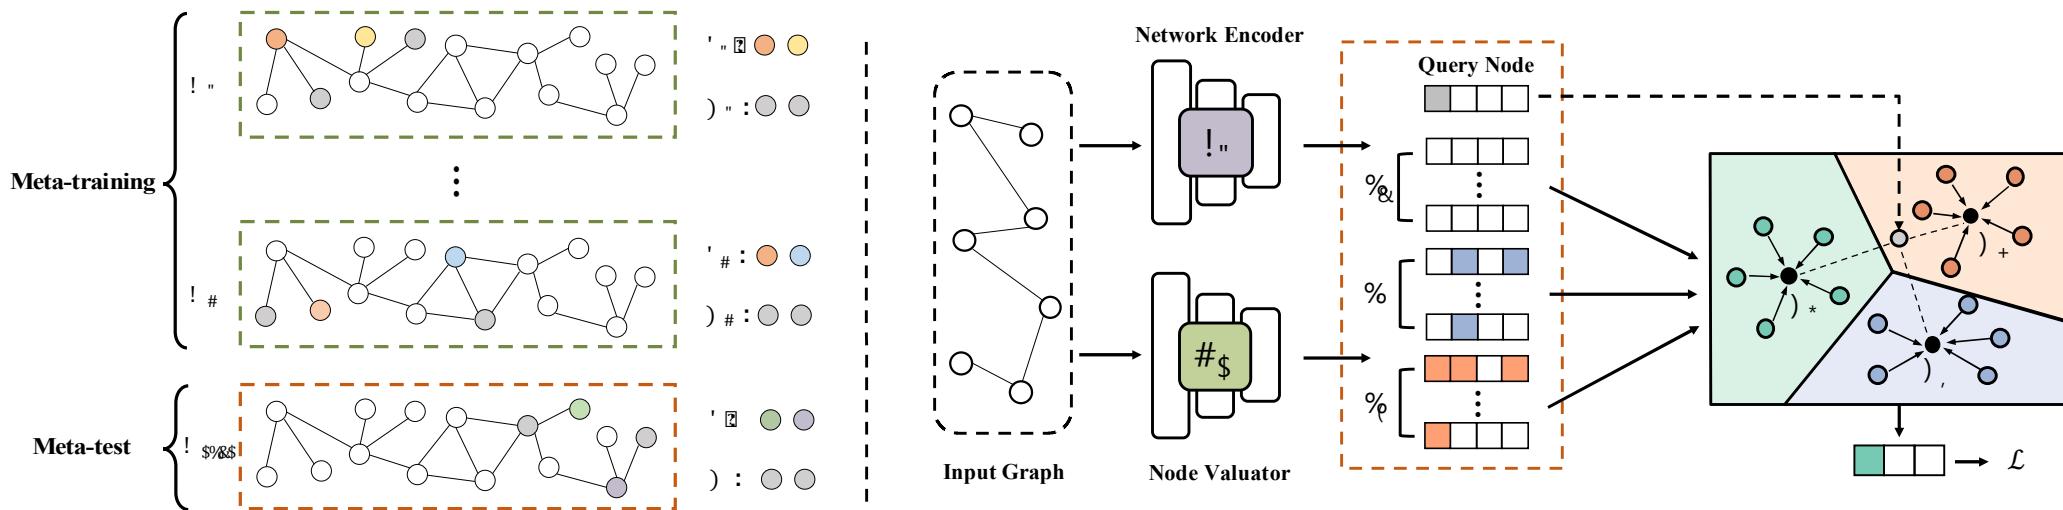


Figure 2: (Left) Episodic training on attributed networks. In each episode, we create a semi-supervised few-shot node classification task by random sampling; (Right) The architecture of the proposed framework Graph Prototypical Networks (GPN).

# Structure-based Enhancement on Graphs

- Node-level enhancement: FAAN

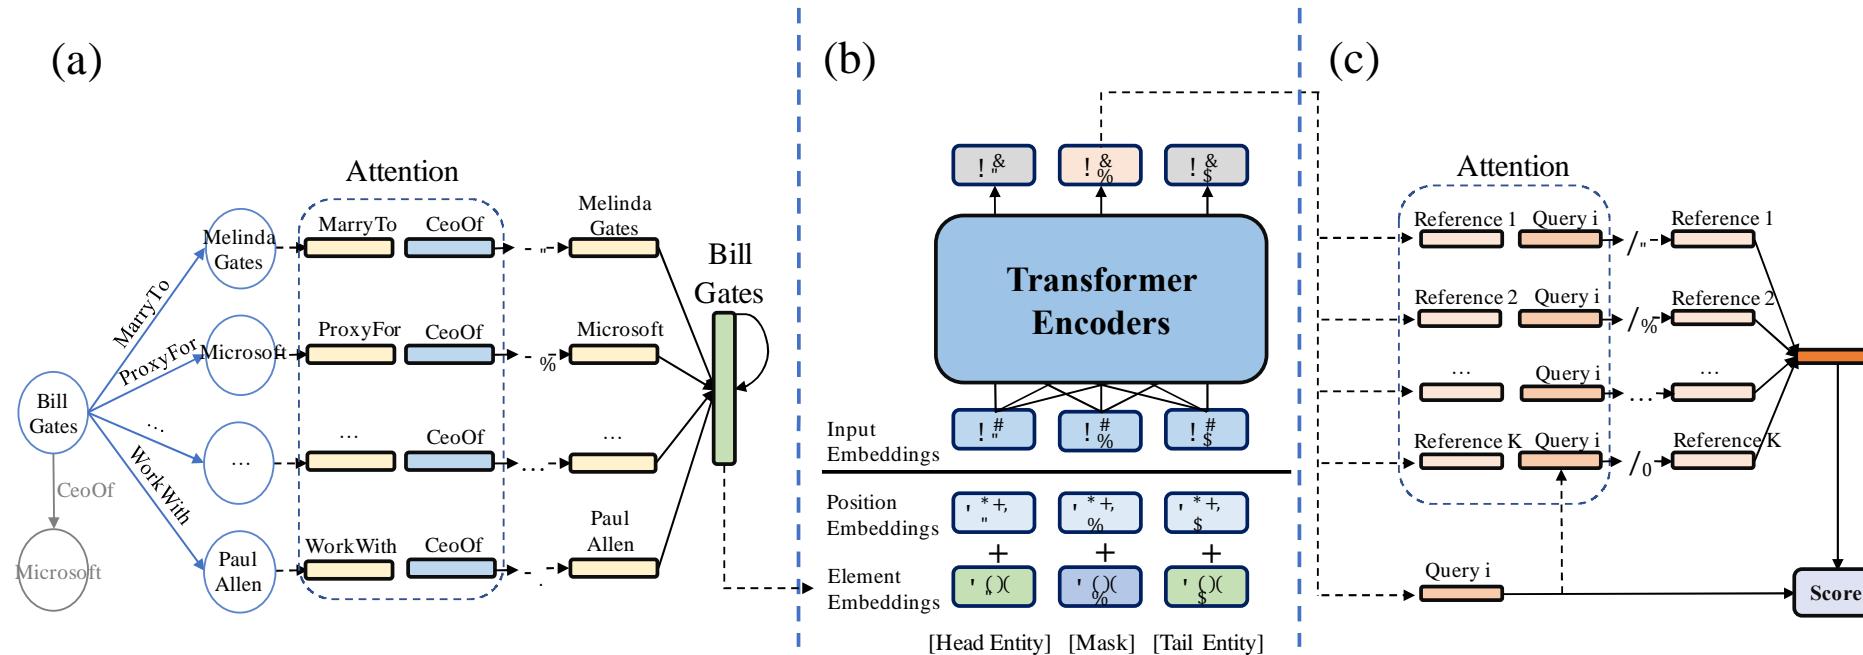


Figure 2: The framework of FAAN: (a) Adaptive neighbor encoder for entities; (b) Transformer encoder for entity pairs; (c) Adaptive matching processor to match  $K$ -shot references and the query.

# Structure-based Enhancement on Graphs

- Edge-level enhancement: HMNet, TCVAE
  - Leverage auxiliary information associated with edges

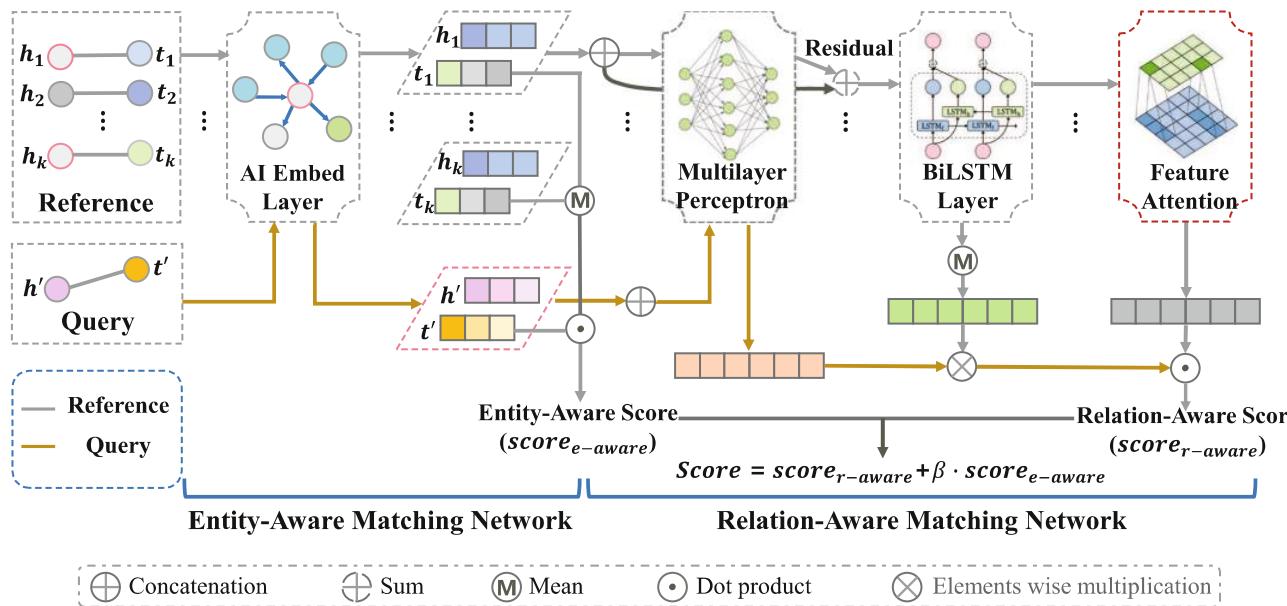


Fig. 2. Illustration of the proposed HMNet model.

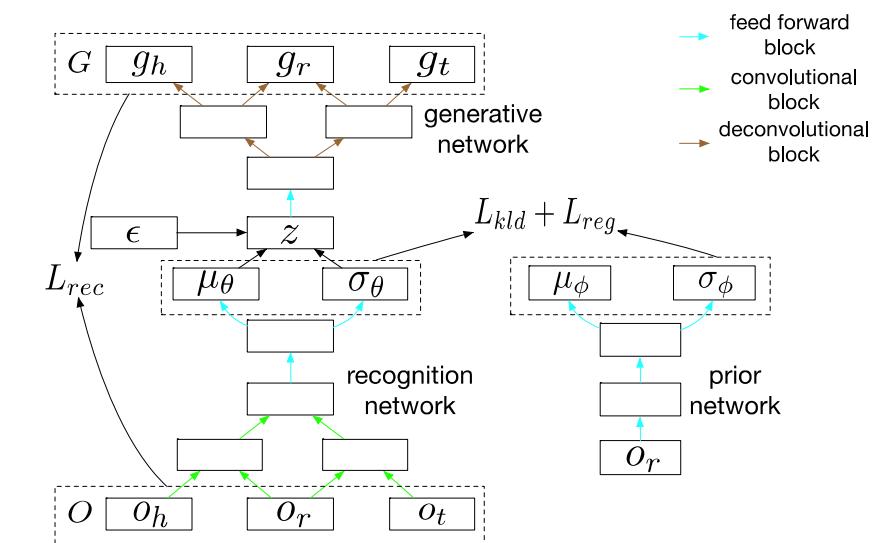


Figure 3: Structure of TCVAE. This figure should be viewed in color.

S. Xiao, et al. "HMNet- Hybrid Matching Network for Few-Shot Link Prediction." DASFAA'21

Z. Wang, et al "Tackling Long-Tailed Relations and Uncommon Entities in Knowledge Graph Completion." EMNLP'19

# Structure-based Enhancement on Graphs

- Edge-level enhancement: **RALE**
  - Leverage **paths** to capture long-range dependencies between distant node

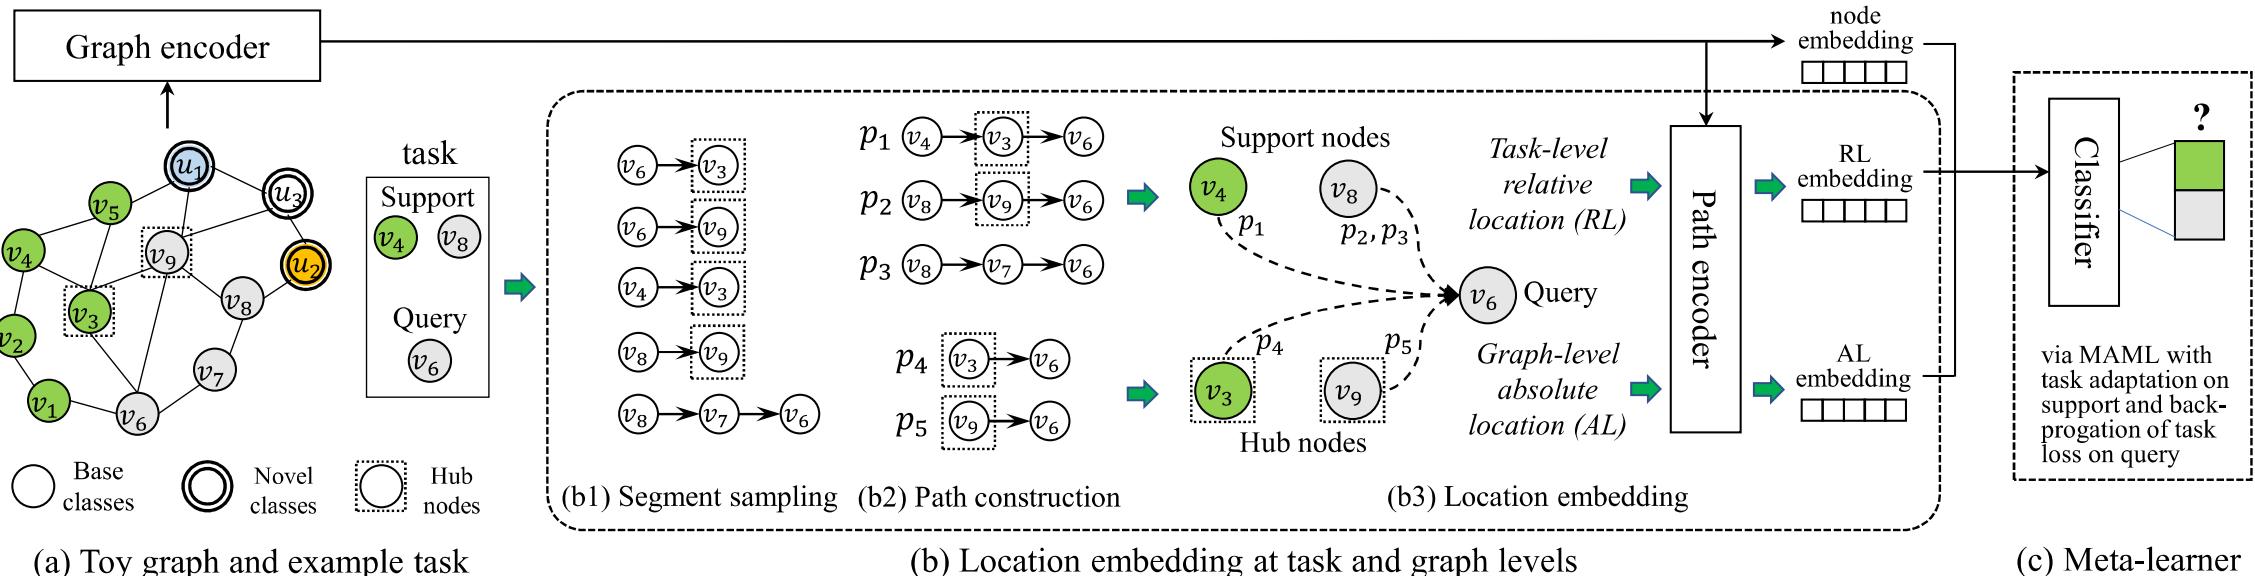


Figure 3: Overview of the proposed model RALE.

- Paths between each query node and the support nodes: Task-level dependencies
- Paths between each target node and the hub nodes: Graph-level dependencies

# Structure-based Enhancement on Graphs

- Edge-level enhancement: MetaHIN

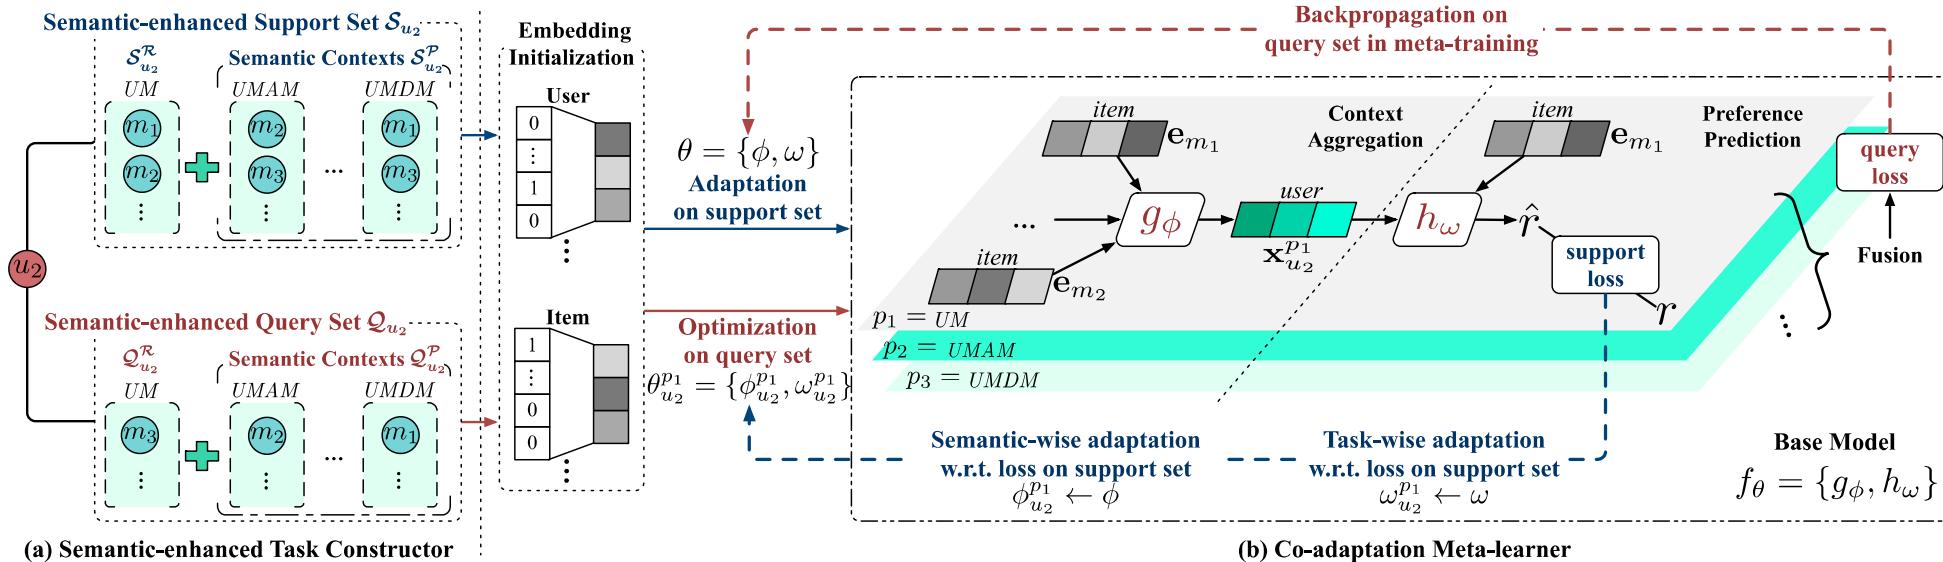
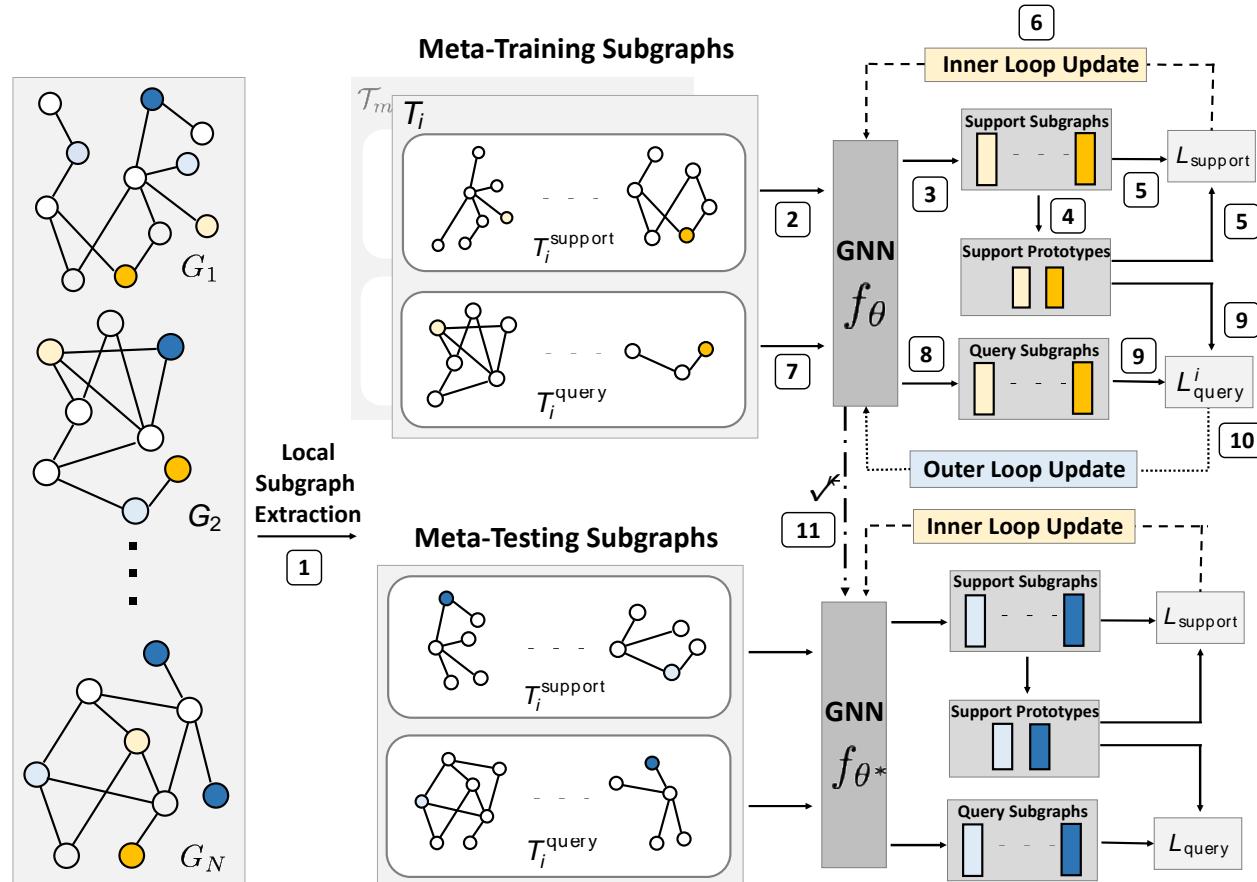


Figure 2: Illustration of the meta-training procedure of a task in MetaHIN. (a) Semantic-enhanced task constructor, where the support and query sets are augmented with meta-path based heterogeneous semantic contexts. (b) Co-adaptation meta-learner, with semantic- and task-wise adaptations on the support set, while the global prior  $\theta$  is optimized on the query set. During meta-testing, each task follows the same procedure except updating the global prior.

- Meta-paths: heterogeneous semantic relationships

# Structure-based Enhancement on Graphs

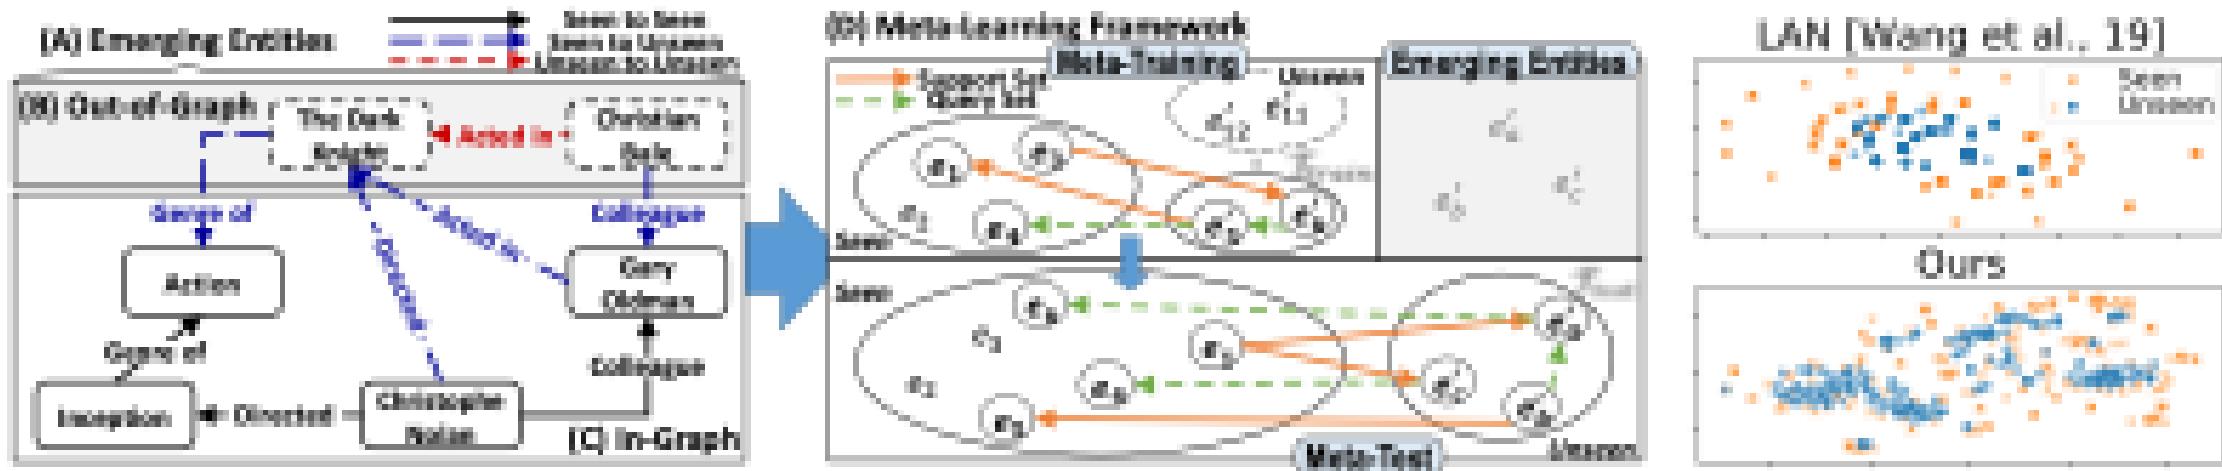
- Subgraph-level enhancement: G-Meta



- Generate class prototypes from subgraph
- Expand query node to its subgraph

# Structure-based Enhancement on Graphs

- Subgraph-level enhancement: GEN



- Extrapolate knowledge through the neighbors (one-hop subgraph) of the support set
- Embeds the unseen entities on the manifold of seen entities

# Structure-based Enhancement on Graphs

- Subgraph-level enhancement: Meta-tail2vec

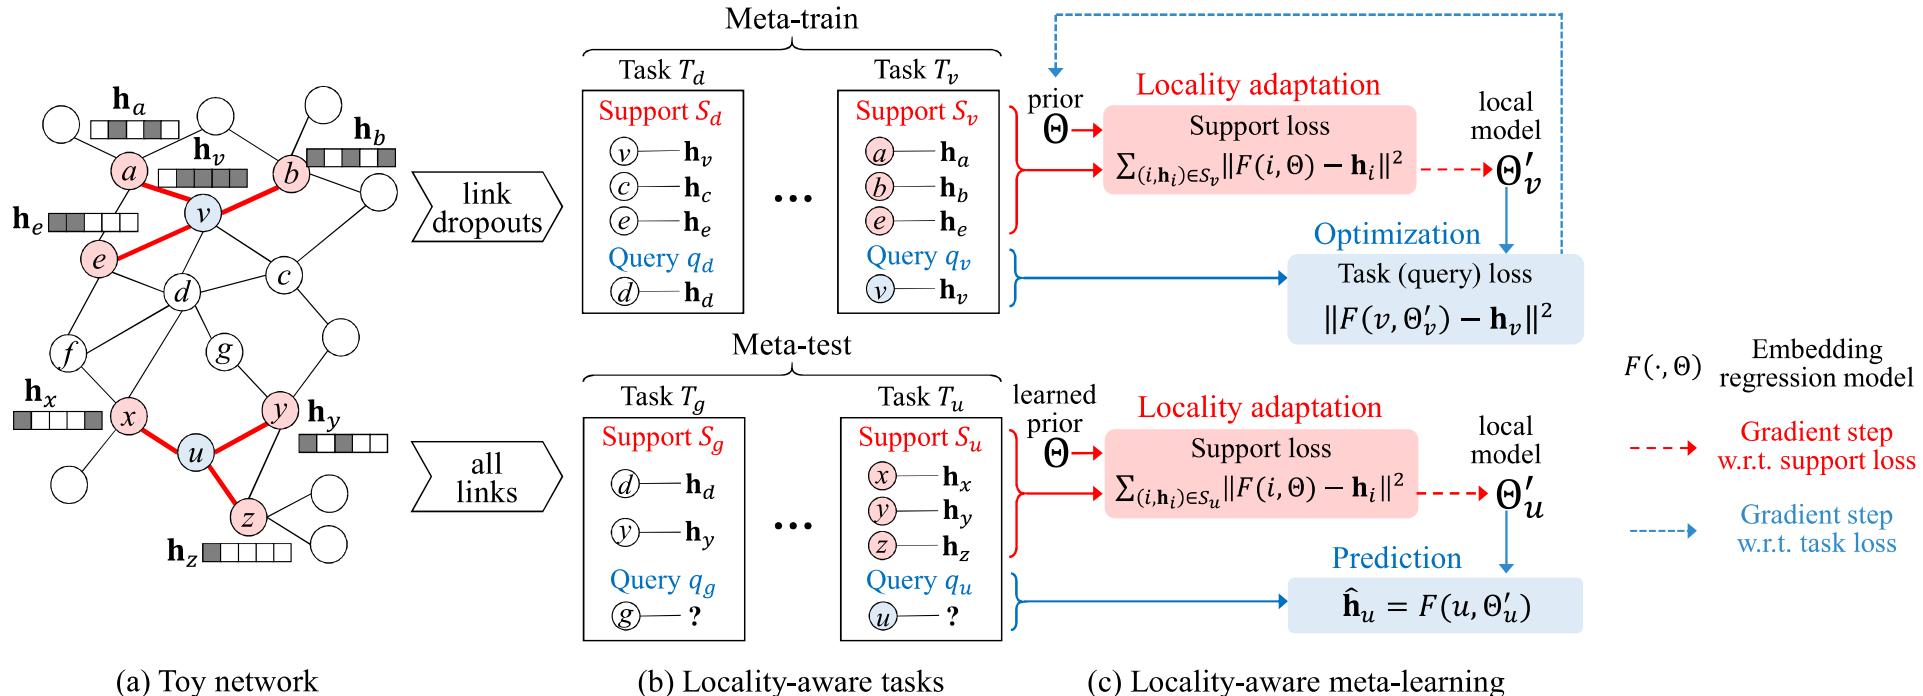


Figure 3: Overall framework of our locality-aware tail node embedding model meta-tail2vec. (Best viewed in color.)

- Locality-aware tasks: support set sampled from the neighborhood subgraph of the query node

# Adaptation-based Enhancement on Graphs

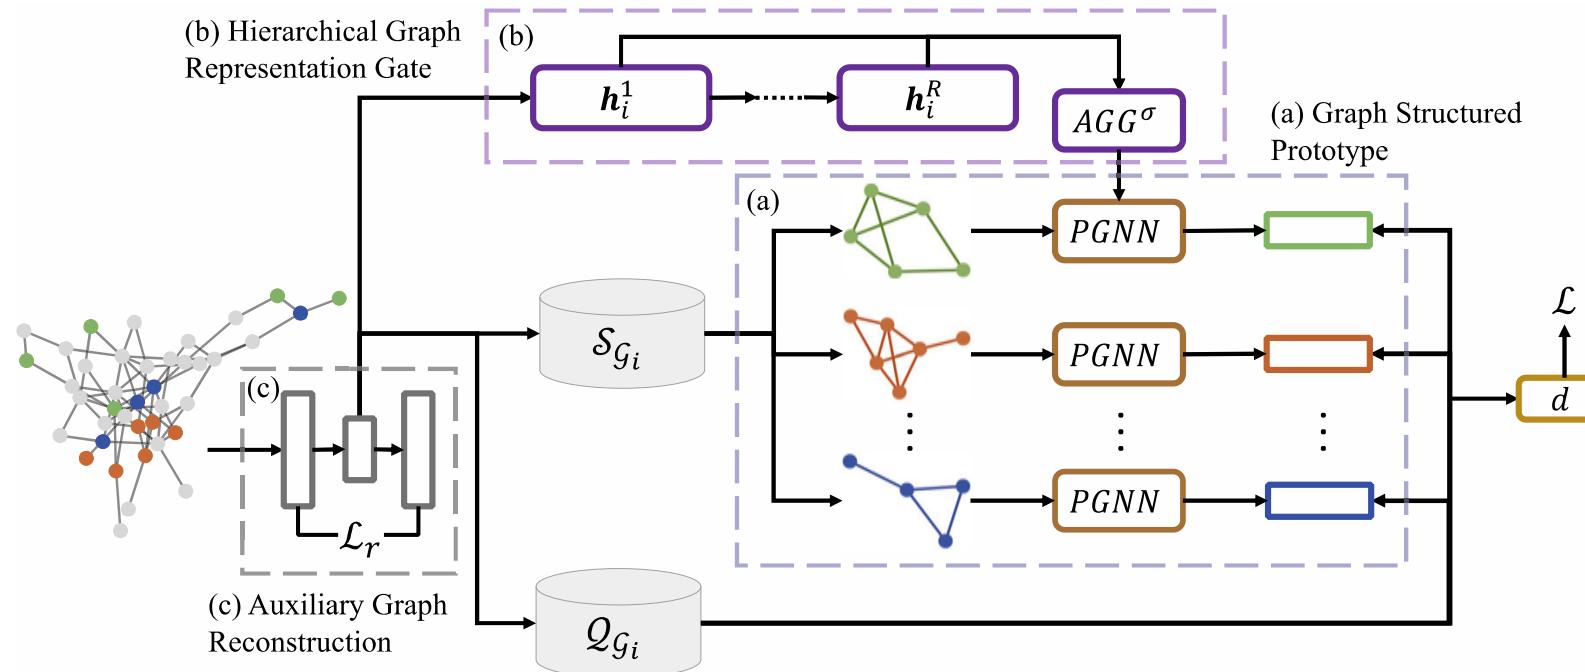
- Customization of a globally shared prior into a localized or specialized model for each task

TABLE V: Adaptation-based meta-learning enhancement for few-shot learning on graphs.

Method	Adaptation enhancement	Meta learner	Task		
			Node	Edge	Graph
GFL [36]	graph	MAML	✓	✗	✗
MI-GNN [145]	graph	hybrid	✓	✗	✗
MetaTNE [32]	task	Protonets	✓	✗	✗
AMM-GNN [65]	task	MAML	✓	✗	✗
AS-MAML [148]	step size	MAML	✗	✗	✓
MetaDyGNN [137]	hybrid	MAML	✗	✓	✗

# Adaptation-based Enhancement on Graphs

- Graph-wise adaptation: **GFL**



- Recognize the topological variances across different graphs
- Customize a global prior for each individual graph (class prototypes tailored to each graph)
- Apply gate function to the global prior

# Adaptation-based Enhancement on Graphs

- Graph-wise adaptation: MI-GNN

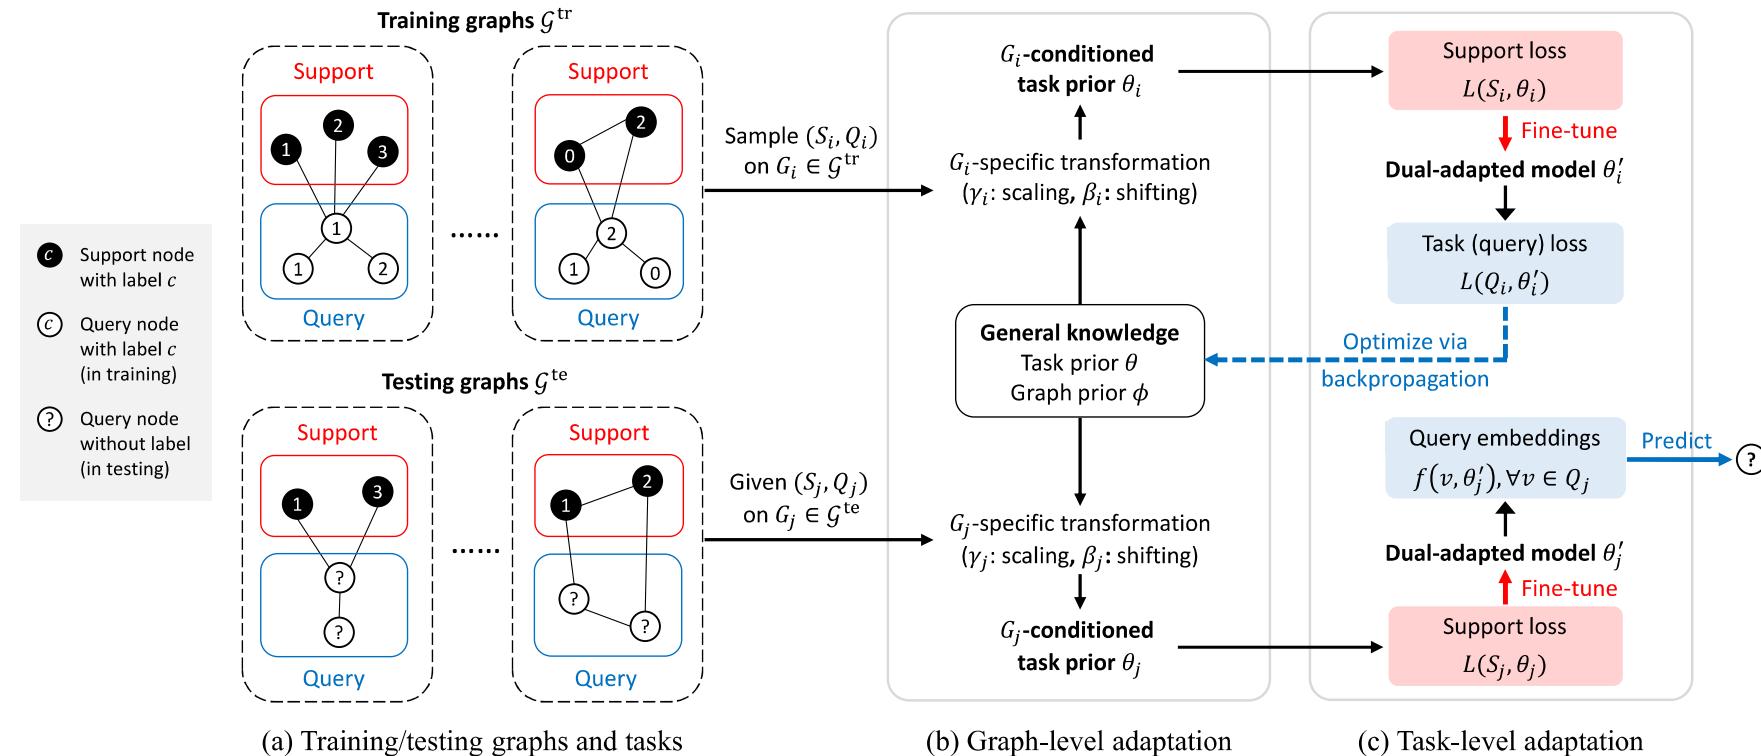
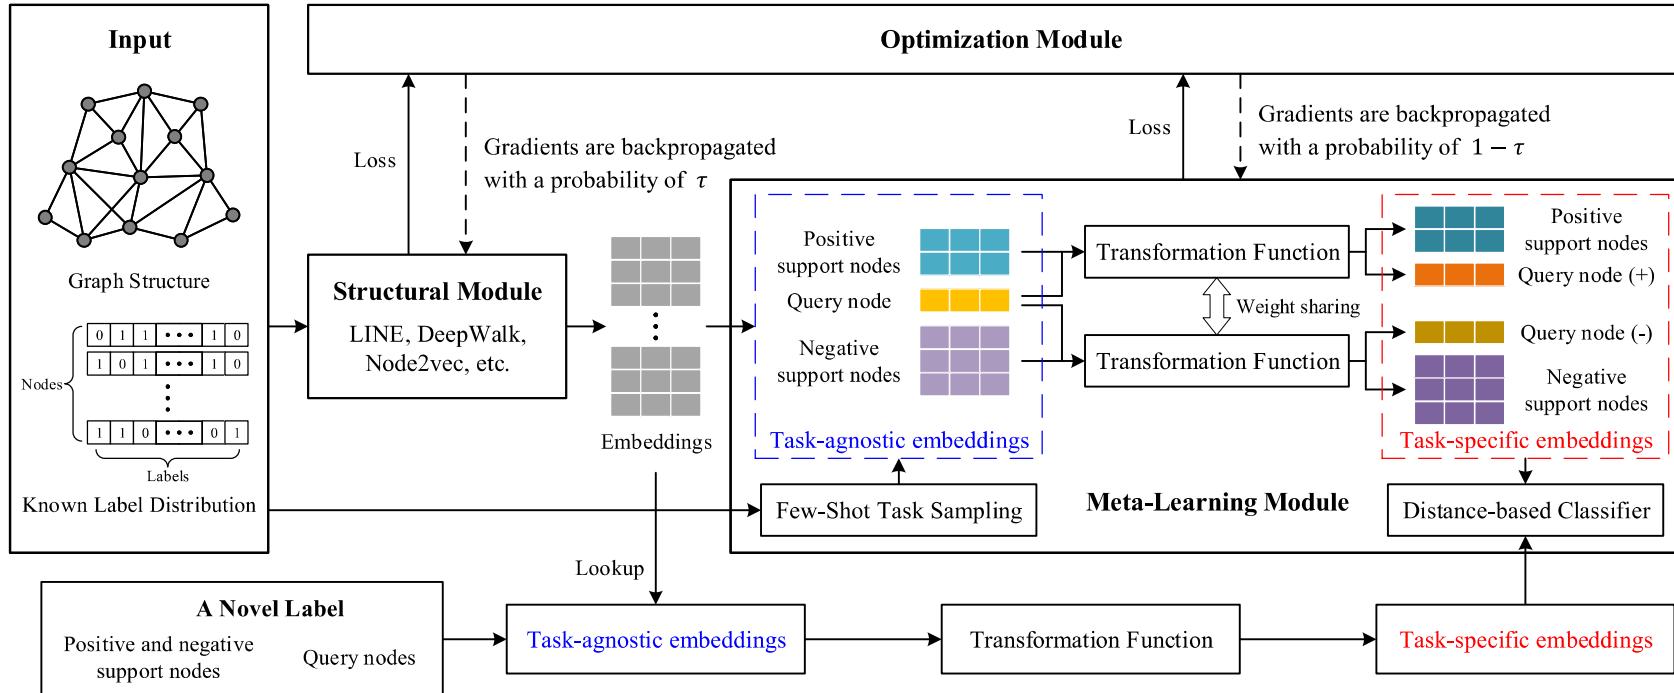


Figure 2: Overall framework of MI-GNN, illustrating the pipeline on a training graph  $G_i$  and a testing graph  $G_j$ .

- Employ a Feature-wise Linear Modulation (FiLM) to modulate the global prior for each graph

# Adaptation-based Enhancement on Graphs

- Task-wise adaptation: MetaTNE



- Multi-label few-shot classification: same node could be associated with different labels in different tasks
- Adaptation for the node embeddings (the query set in each task)

# Adaptation-based Enhancement on Graphs

- Task-wise adaptation: AMM-GNN

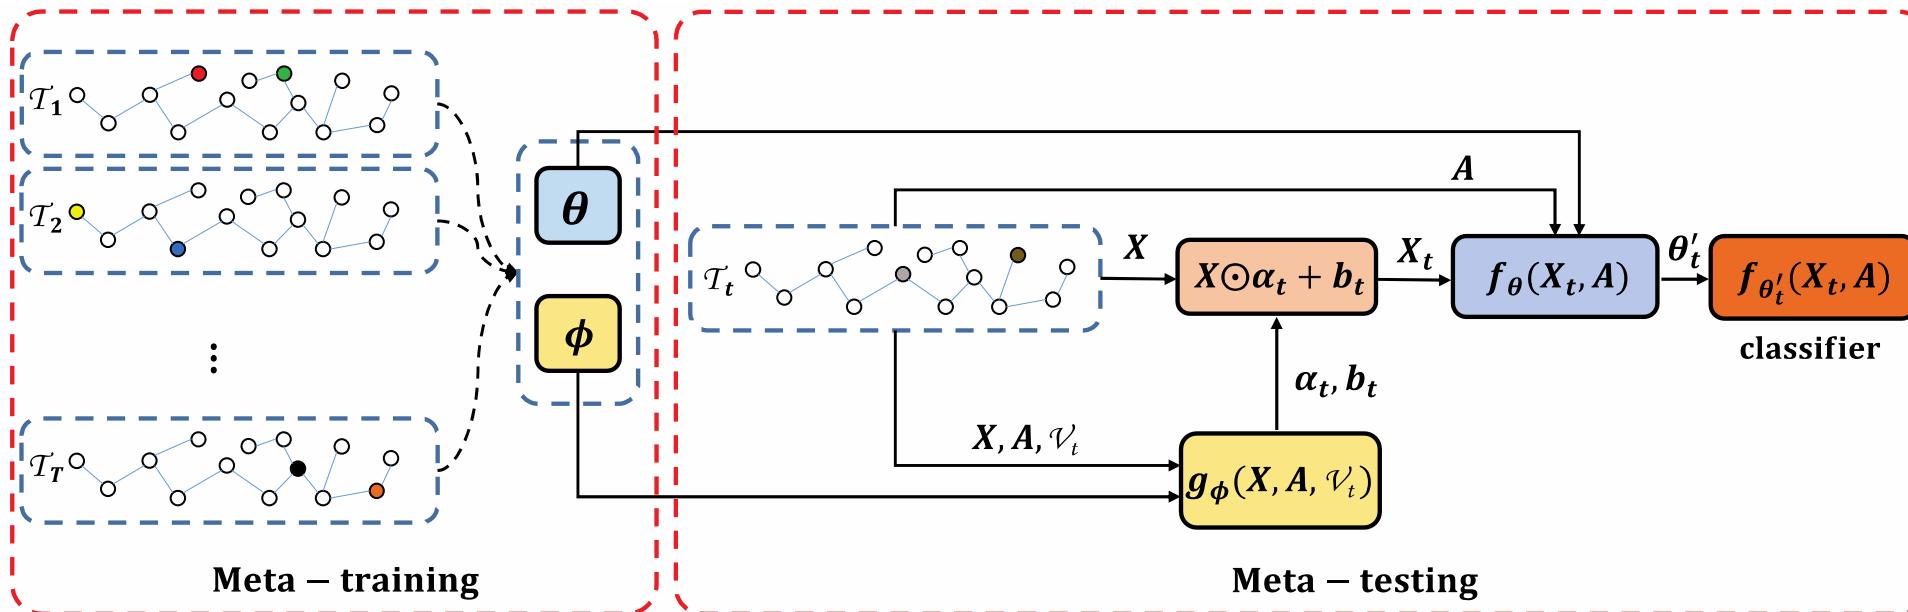
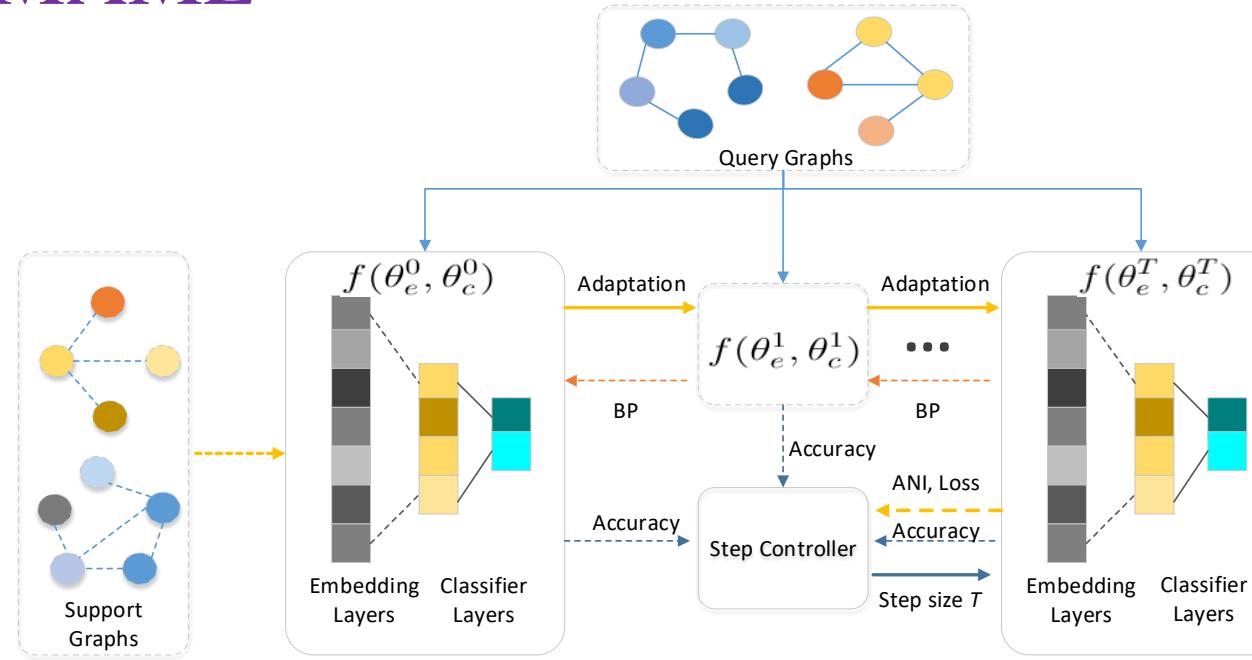


Figure 1: The overview of the proposed AMM-GNN framework. *Left:* In the meta-training phase, multiple tasks are sampled to train the meta-learning model, and we obtain two parameter sets  $\theta$  and  $\phi$ . *Right:* In the meta-testing phase, we use parameter sets  $\phi$  and  $\theta$  for attribute matching and gradient descent respectively, and obtain the classifier  $f_{\theta'}(\cdot)$  for a new sampled task  $T_t$ .

- Customize a task-specific feature matrix for adaptation

# Adaptation-based Enhancement on Graphs

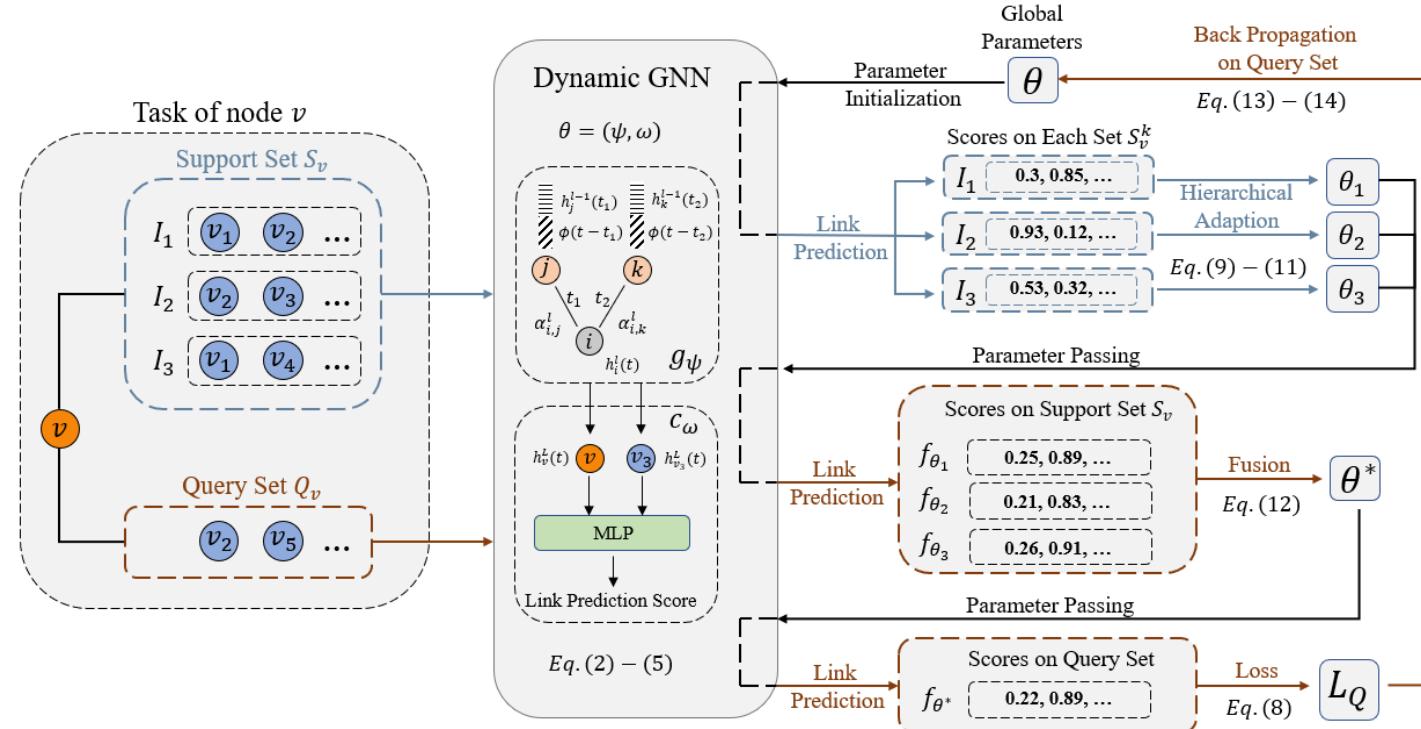
- Others: AS-MAML



- Improve adaptation from an optimization standpoint
- Reinforcement learning-based controller to determine the optimal step size for the adaptation process

# Adaptation-based Enhancement on Graphs

- Others: MetaDyGNN



- Adaptation for dynamic graphs: time- and node-wise

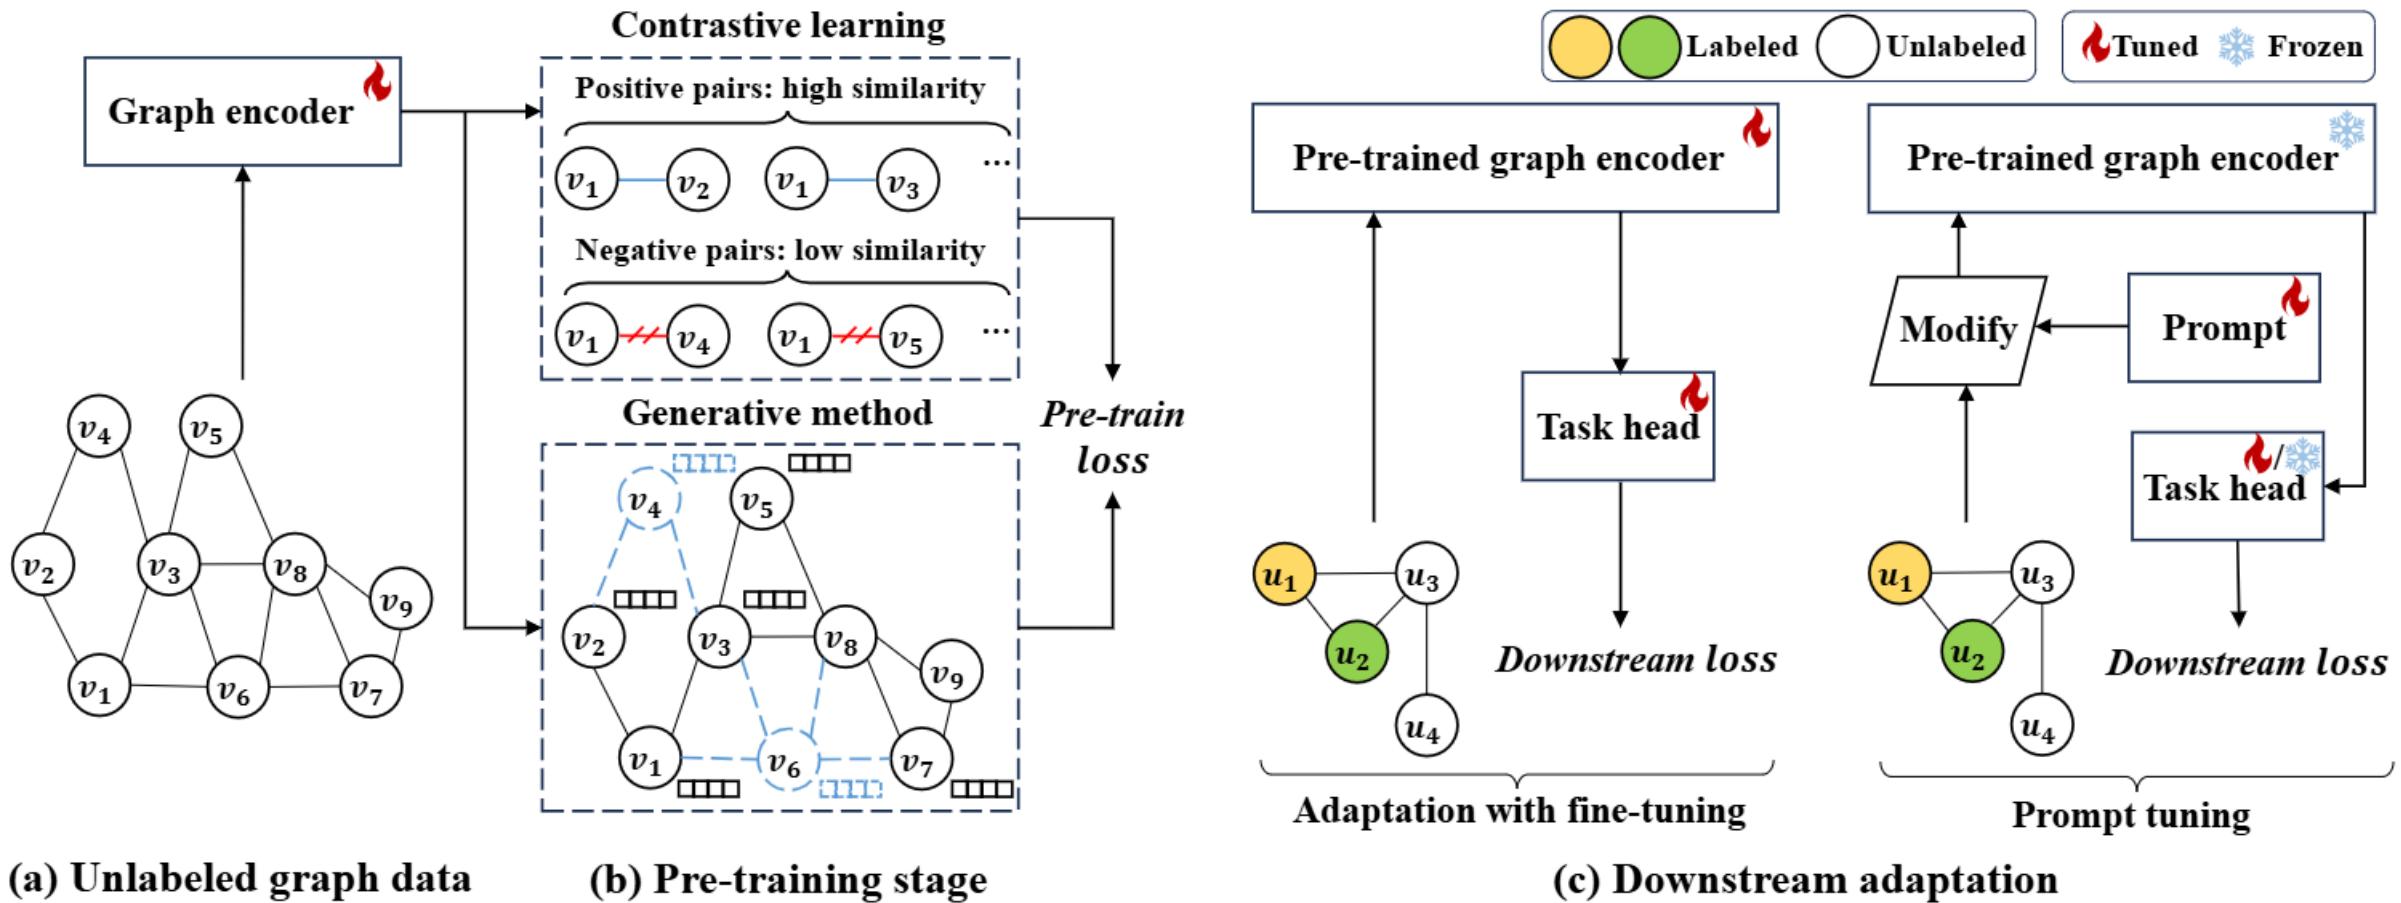
# Summary

- Existing research often enhances a **standard meta-learner**: structural augmentation or refining the adaptation process
- Drawbacks:
  - Require **abundant labels** for a base set during the meta-training phase
  - Fail to leverage the vast amount of **unlabeled data** to learn a more comprehensive prior
  - Limited by the i.i.d. **assumption** in task distribution, and cannot handle **different types of downstream tasks**

Can we address a *diverse* range of few-shot tasks on graphs *without an extensively annotated base set*, while *utilizing abundant unlabeled graphs*?

# Pre-training on Graphs

- Pre-training stage utilizes self-supervised method
- Prior knowledge are then adapted to downstream tasks



# Pre-training Strategies

- Graph pre-training strategies mainly fall into:

- Contrastive strategies

- Generative strategies

# Contrastive Strategies

- Contrasting instances at various scales within a graph

➤ Sample positive and negative instances

➤ Positive instances closer to the target

➤ Negative instances further to the target

Pre-training data  $\mathcal{T}_{\text{pre}}$       Target instance  $o$

Positive samples  $\mathcal{P}_o$       Negative samples  $\mathcal{N}_o$

Contrastive loss:

$$-\sum_{o \in \mathcal{T}_{\text{pre}}} \ln \frac{\sum_{a \in \mathcal{P}_o} \exp\left(\frac{\text{sim}(\mathbf{h}_a, \mathbf{h}_o)}{\tau}\right)}{\sum_{a \in \mathcal{P}_o} \exp\left(\frac{\text{sim}(\mathbf{h}_a, \mathbf{h}_o)}{\tau}\right) + \sum_{b \in \mathcal{N}_o} \exp\left(\frac{\text{sim}(\mathbf{h}_b, \mathbf{h}_o)}{\tau}\right)}$$

Method	Instance	Augmentation	Graph types
GRACE [72]	node	uniform	general
GCC [30]	graph	uniform	general
GraphCL [40]	graph	uniform	general
SimGRACE [74]	graph	perturbing encoder	general
GraphLoG [73]	dataset	uniform	general
DGI [29]	cross-scale	uniform	general
InfoGraph [42]	cross-scale	uniform	general
Subg-Con [71]	cross-scale	uniform	general
MVGRL [149]	cross-scale	diffusion	general
JOAO [41]	graph	adaptive to loss	general
GCGM [150]	node	adaptive to loss	general
You <i>et al.</i> [151]	graph	view generator	general
GCA [152]	node	adaptive to instance	general
HeCo [153]	node	uniform	hetero.
CPT-HG [154]	cross-scale	uniform	hetero.
PT-HGNN [155]	cross-scale	uniform	hetero.
SelfRGNN [76]	node	curvature over time	dynamic
DDGCL [156]	graph	uniform	dynamic
CPDG [75]	cross-scale	temporal-aware sampling	dynamic
GearNet [157]	graph	uniform	3D

# Generative Strategies

- Reconstruct parts of the graph

➤ Structure reconstruction

- Entire graph structure
- Part of graph structure

➤ Feature reconstruction

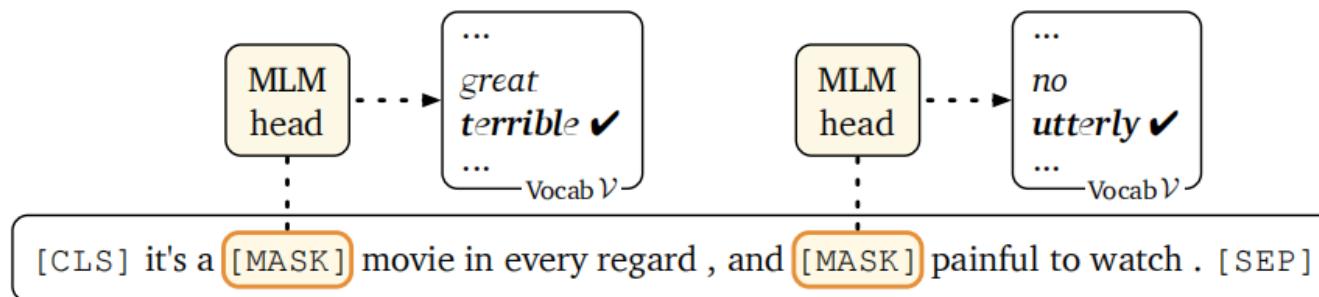
- Origin feature
- Latent embedding

Method	Reconstruction objective						Graph type
	node feat.	node deg.	edge	adj. matrix	graph feat.	other info.	
VGAE [43]	✗	✗	✗	✓	✗	✗	general
GPT-GNN [39]	✓	✗	✓	✗	✗	✗	general
MaskGAE [77]	✗	✓	✓	✗	✗	✗	general
NWR-GAE [161]	✓	✓	✗	✗	✗	✗	general
LaGraph [162]	✓	✗	✗	✗	✓	✗	general
GraphMAE [163]	✓	✗	✗	✗	✗	✗	general
GraphMAE2 [78]	✓	✗	✗	✗	✗	✗	general
Liu <i>et al.</i> [164]	✓	✗	✗	✗	✗	✗	KG
Wen <i>et al.</i> [79]	✓	✗	✓	✗	✗	✗	KG
MPKG [165]	✓	✗	✓	✗	✗	✓	KG
PT-DGNN [166]	✗	✗	✓	✗	✗	✗	dynamic
STEP [167]	✓	✗	✗	✗	✗	✗	dynamic
PMGT [168]	✓	✗	✓	✗	✗	✓	MMG
ColdGPT [169]	✓	✗	✗	✗	✗	✓	MMG

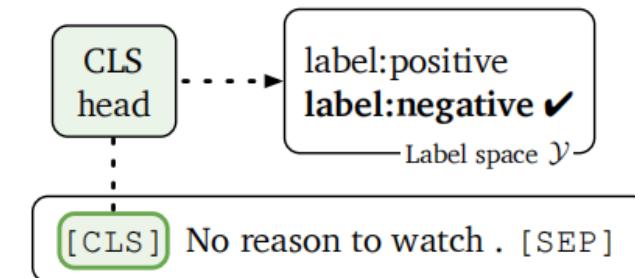
# Fine-tuning

- Prior knowledge are transferred to downstream tasks by initializing a downstream model with the pre-trained weights
  - Task-specific projection head
  - Update the parameters in
    - Pretrained model
    - Task head
  - Objective gap between pretext and downstream tasks
  - Updating all parameters is inefficient

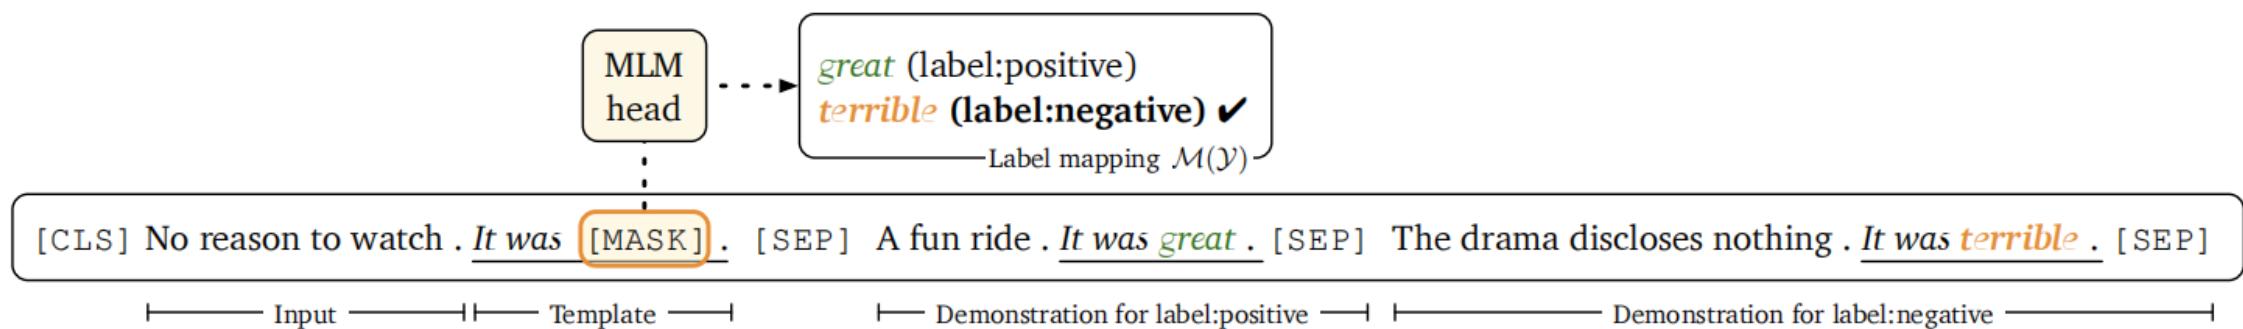
# Prompt tuning



(a) MLM pre-training



(b) Fine-tuning



# Prompt tuning

- Unified template

➤ Aligns the pretext and downstream losses

- Prompt

➤ Modify the original input/embedding for the pre-trained model

Paper	Template	Feature prompt	Structure prompt	Multiple pretext tasks	Prompt Initialization	Downstream Task	Node	Edge	Graph
GPPT [48]	subgraph-token similarity: $\text{sim}(\mathbf{s}_v, \mathbf{t}_y)$	input	✗	✗	random	✓	✗	✗	
VPGNN [178]	node-token matching: $\text{match}(\mathbf{h}_v, \mathbf{t}_y)$	✗	✓	✗	random	✓	✗	✗	
GraphPrompt [19] MOP [179]	subgraph similarity: $\text{sim}(\mathbf{s}_u, \mathbf{s}_v)$	readout	✗	✗	random	✓	✓	✓	
GraphPrompt+ [80] ProNoG [180]		readout	✗	✗	random	✗	✓	✗	
MDGPT [181]		all layers	✗	✗	random	✓	✓	✓	
MultiGPrompt [84] HetGPT [116]		readout	✗	✗	conditional	✓	✓	✓	
GPF [182] IGAP [115]	node similarity: $\text{sim}(\mathbf{h}_u, \mathbf{h}_v)$	readout all layers input	✗ ✗ ✗	✗ ✓ ✗	pretext tokens pretext tokens random	✓ ✓ ✓	✓ ✓ ✗	✓ ✓ ✗	
SGL [90]	dual-template: $\text{CL}(\mathbf{h}_u, \mathbf{h}_v), \text{GL}(\mathbf{x}_v, \tilde{\mathbf{x}}_v)$	✗	✓	✓	random	✗	✗	✓	
HGPrompt [83]	dual-template: $\text{sim}(\mathbf{s}_u, \mathbf{s}_v)$ , graph template	readout	✗	✗	random	✓	✓	✓	
SAP [85]	view similarity: $\text{sim}(\text{MLP}(X), \text{GNN}(X, A))$	✗	✓	✓	random	✓	✗	✓	
ULTRA-DP [86]	node-node/group similarity: $\text{sim}(\mathbf{h}_u, \mathbf{h}_v)$	input	✓	✓	random	✓	✓	✗	
VNT [87]	node attribute reconstruction: $\text{MSE}(\mathbf{x}_v, \tilde{\mathbf{x}}_v)$ structure recovery: $\text{MSE}(\{\mathbf{h}_u, \mathbf{h}_v\})$	input	✗	✗	meta-trained	✓	✗	✗	
ProG [49]	subgraph classification: $\text{CLS}(\mathbf{s})$	✗	✓	✗	meta-trained	✓	✓	✓	
DyGPrompt [183] TIGPrompt [184]	temporal node similarity: $\text{sim}(\mathbf{h}_{t,u}, \mathbf{h}_{t,v})$	input input	✗ ✗	✗ ✗	conditional time-based	✓ ✓	✓ ✓	✗ ✗	

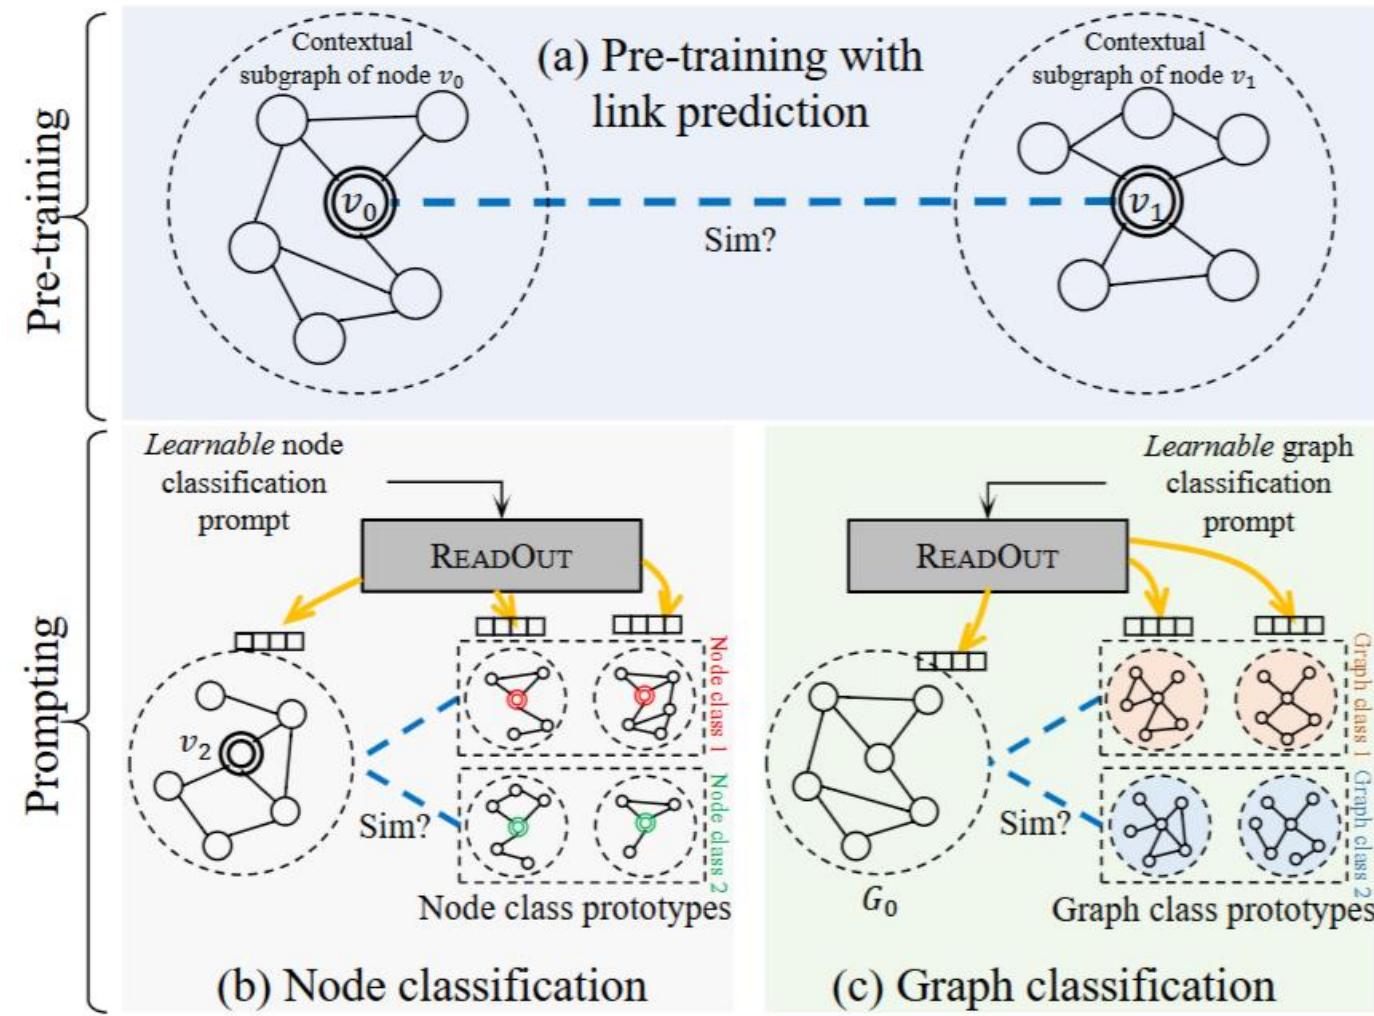
# GraphPrompt

- Motivation

- Gap between graph pre-training and downstream tasks

- Challenges

- What is the unified task template?
  - How to design task-specific prompts?



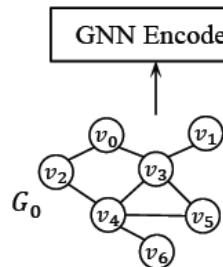
Liu, et al. "Graphprompt: Unifying pre-training and downstream tasks for graph neural networks." WWW'23.

# GraphPrompt

Unified task template

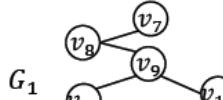
Link Prediction

$$\text{sim}(\mathbf{s}_v, \mathbf{s}_a) > \text{sim}(\mathbf{s}_v, \mathbf{s}_b)$$

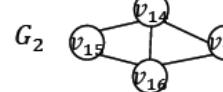


Node Classification(NC)

$$\tilde{\mathbf{s}}_c = \frac{1}{k} \sum_{(v_i, \ell_i) \in D, \ell_i=c} \mathbf{s}_{v_i}$$



$$\ell_j = \arg \max_{c \in C} \text{sim}(\mathbf{s}_{v_j}, \tilde{\mathbf{s}}_c)$$



Graph Classification(GC)

$$\tilde{\mathbf{s}}_c = \frac{1}{k} \sum_{(G_i, L_i) \in \mathcal{D}, L_i=c} \mathbf{s}_{G_i}$$

(a) Toy graphs

$$L_j = \arg \max_{c \in C} \text{sim}(\mathbf{s}_{G_j}, \tilde{\mathbf{s}}_c)$$

mean embedding of (sub)graphs  
class label

A Notation for NC and GC

$$y = \arg \max_{c \in Y} \text{sim}(\mathbf{s}_x, \tilde{\mathbf{s}}_c)$$

$$\mathbf{s}_x = \text{READOUT}(\{\mathbf{h}_v : v \in V(S_x)\})$$

Pre-Training Objective

$$\mathcal{L}_{\text{pre}}(\Theta) = - \sum_{(v, a, b) \in \mathcal{T}_{\text{pre}}} \ln \frac{\exp(\text{sim}(\mathbf{s}_v, \mathbf{s}_a)/\tau)}{\sum_{u \in \{a, b\}} \exp(\text{sim}(\mathbf{s}_v, \mathbf{s}_u)/\tau)}$$

Liu, et al. "Graphprompt: Unifying pre-training and downstream tasks for graph neural networks." WWW'23.

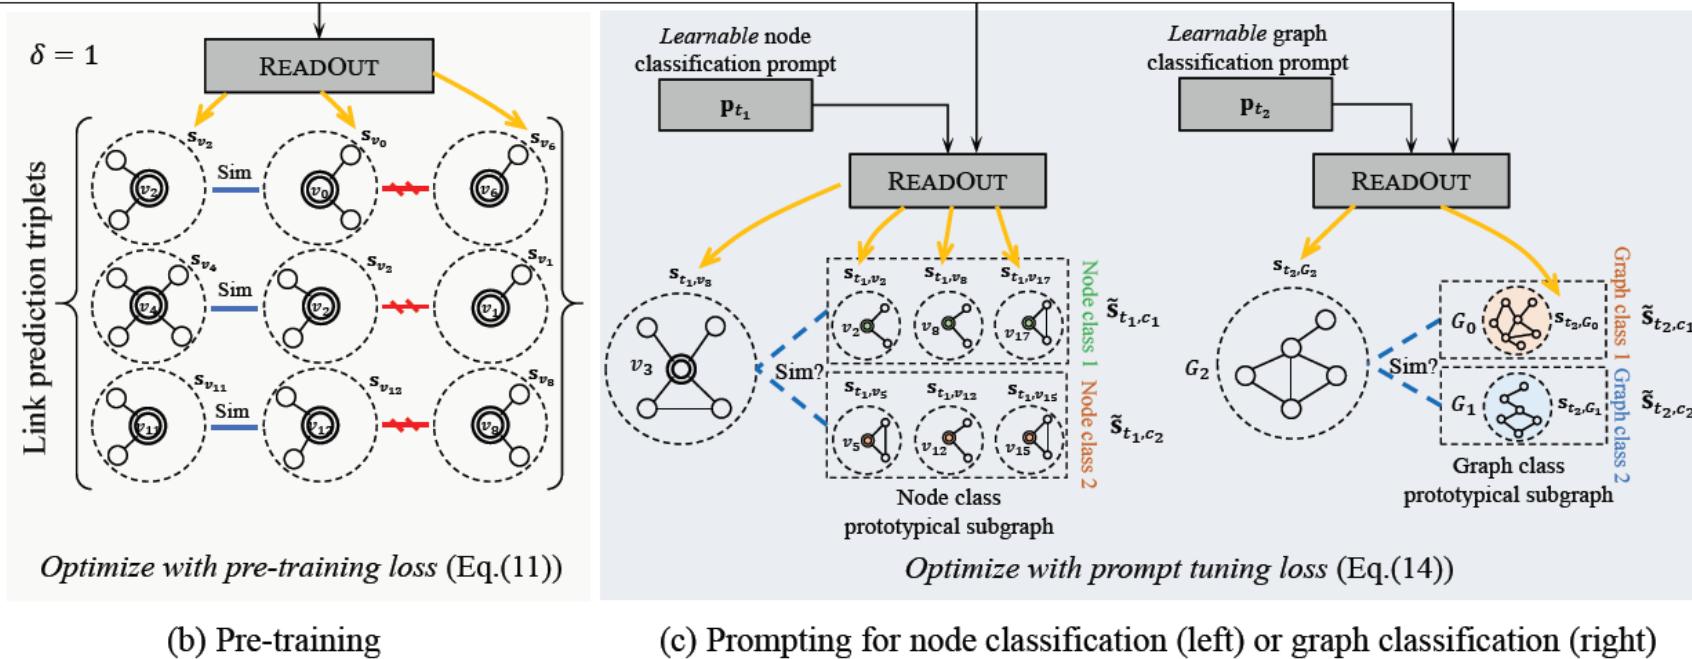


Figure 2: Overall framework of GRAPHPROMPT.

Prompt Design

$$\mathbf{s}_{t,x} = \text{READOUT}(\{\mathbf{p}_t \odot \mathbf{h}_v : v \in V(S_x)\})$$

# Generalized Graph Prompt

- Motivation
  - Can more advanced pretext tasks be unified under the subgraph similarity calculation template?
  - How to utilize hierarchical knowledge across multiple layers of the pre-trained graph encoders

# Generalized Graph Prompt

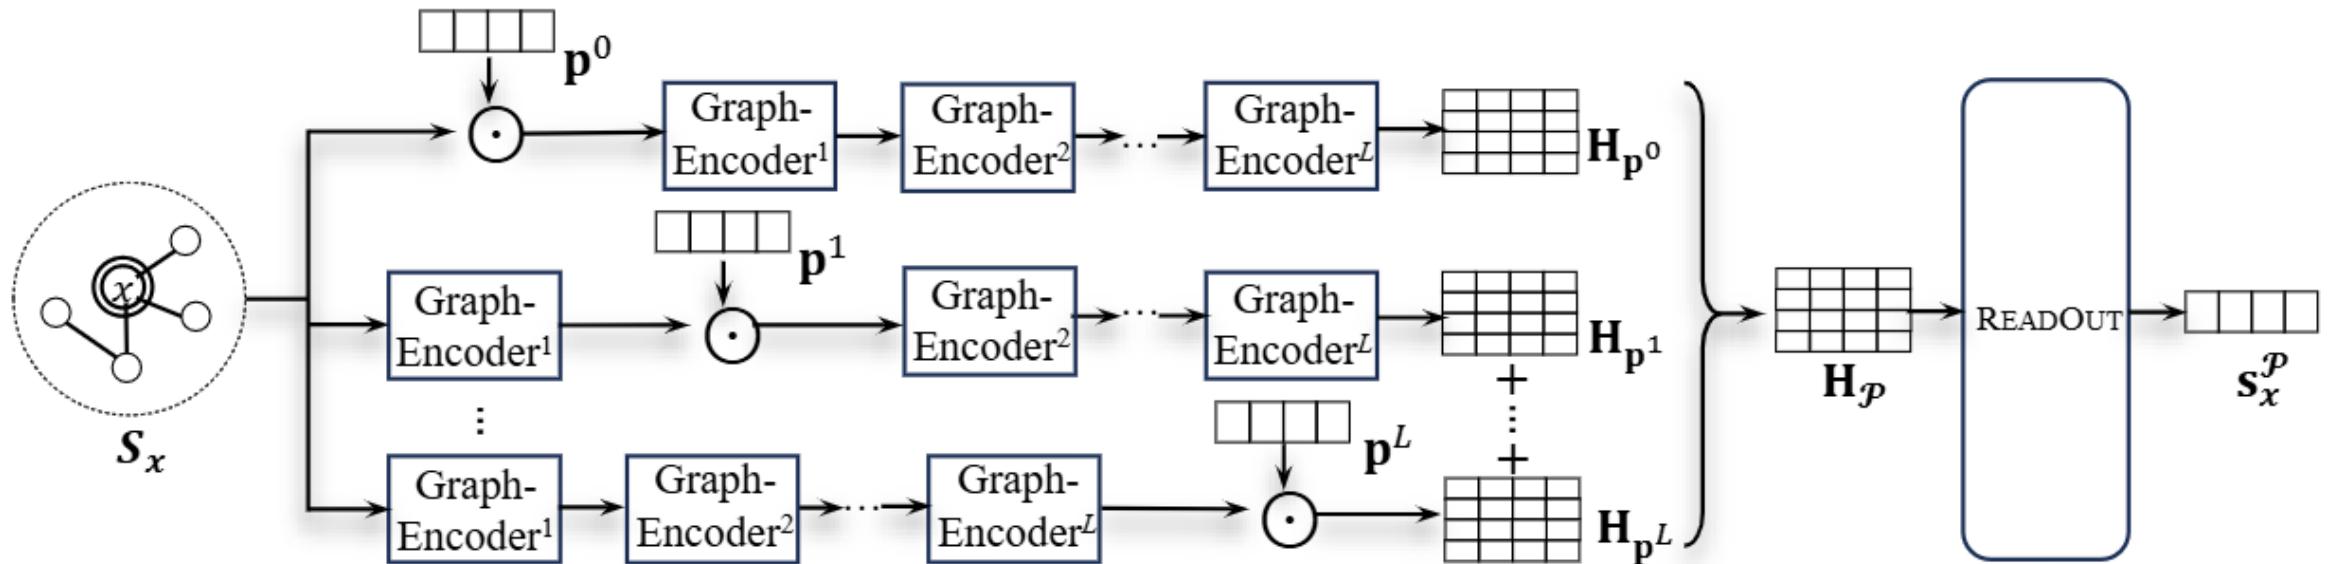
- Any standard contrastive pretext task on graphs can be unified under the loss:

$$\mathcal{L}(\Theta) = - \sum_{o \in \mathcal{T}_{\text{pre}}} \ln \frac{\sum_{a \in Pos_o} \exp(\text{sim}(\mathbf{s}_a, \mathbf{s}_o)/\tau)}{\sum_{b \in Neg_o} \exp(\text{sim}(\mathbf{s}_b, \mathbf{s}_o)/\tau)}$$

	Target instance $o$	Positive instance $a$	Negative instance $b$
LP [39]	a node $v$	a node linked to $v$	a node not linked to $v$
DGI [34]	a graph $G$	a node in $G$	a node in $G'$ , a corrupted graph of $G$
InfoGraph [36]	a graph $G$	a node in $G$	a node in $G' \neq G$
GraphCL [35]	an augmented graph $G_i$ from a graph $G$ by strategy $i$	an augmented graph $G_j$ from a graph $G$ by strategy $j$	an augmented graph $G'_j$ from a graph $G' \neq G$ by strategy $j$
GCC [22]	a random walk induced subgraph $G_v^r$ from a node $v$ 's $r$ -egonet	a random walk induced subgraph $\tilde{G}_v^r \neq G_v^r$ from $v$ 's $r$ -egonet	a random walk induced subgraph $G_v^{r'}$ from $v$ 's $r'$ -egonet, $r' \neq r$

# Generalized Graph Prompt

- Layer wise prompt design



Prompts

$$\mathcal{P} = \{p^0, p^1, \dots, p^L\}$$

Layer-wise modification

$$H_p = \text{GRAPHENCODER}_p(X, A; \Theta)$$

$$H^{l+1} = \text{AGGR}(p^l \odot H^l, A; \theta^{l+1})$$

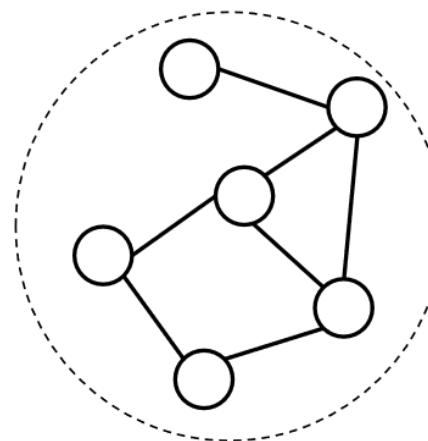
Fusion

$$H_p = \sum_{l=0}^L w^l H_{p^l}$$

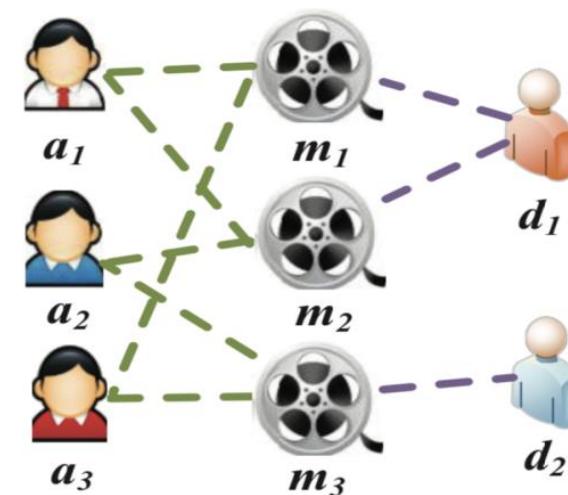
# HGPrompt

- Motivation

- How to unify homogeneous graphs and heterogeneous graphs?
- How to transfer task-specific heterogeneous knowledge?

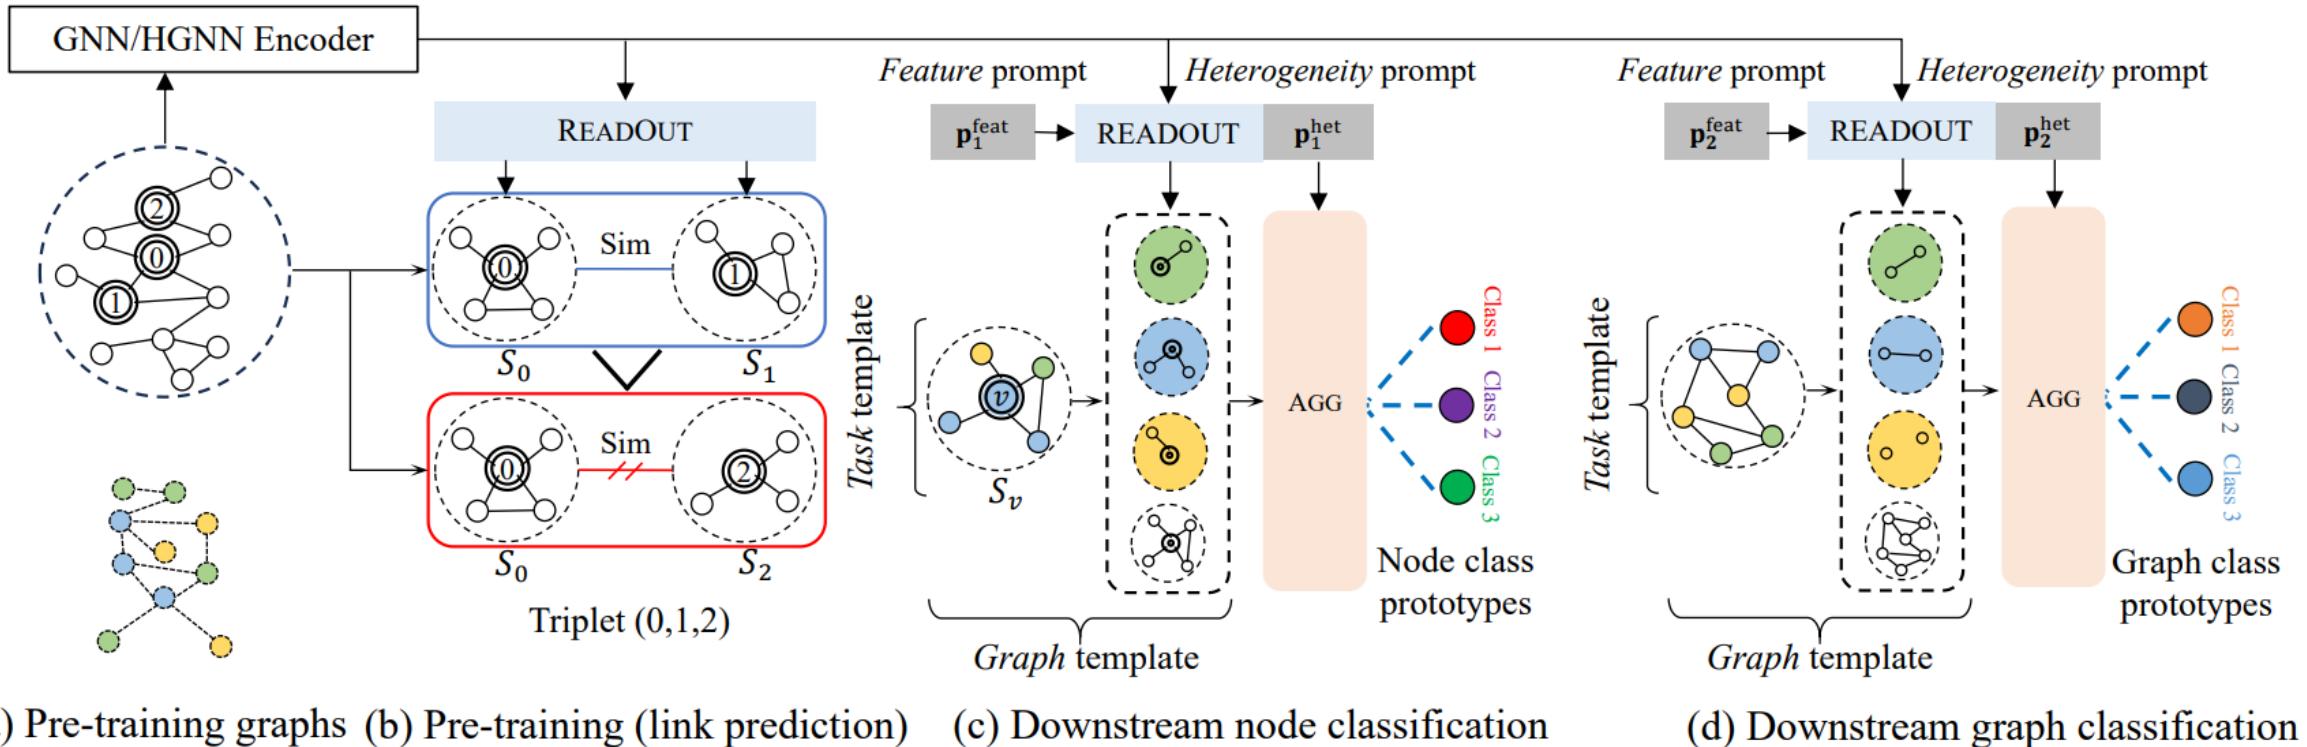


Homogeneous graph



Heterogeneous graph

# HGPrompt



(a) Pre-training graphs

## Dual templates

### Task template

$$\mathcal{GT}(G) = \{G^0\} \cup \{G^i : i \in A\}$$

### Graph template

$$\text{sim}(\mathbf{s}_v, \mathbf{s}_a) > \text{sim}(\mathbf{s}_v, \mathbf{s}_b)$$

$$\ell_j = \arg \max_{c \in C} \text{sim}(\mathbf{s}_{v_j}, \tilde{\mathbf{s}}_c)$$

$$L_j = \arg \max_{c \in C} \text{sim}(\mathbf{s}_{G_j}, \tilde{\mathbf{s}}_c)$$

## Dual prompts

### Feature prompt

$$\text{READOUT}(\{\mathbf{p}^{\text{feat}} \odot \mathbf{h}_v \mid v \in V(S)\})$$

### Heterogeneity prompt

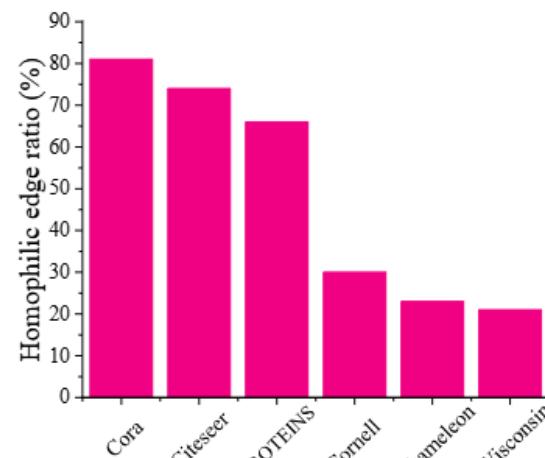
$$\text{AGG}(\{(1 + p_i^{\text{het}}) \odot \text{READOUT}(S^i) \mid S^i \in \mathcal{GT}(S)\})$$



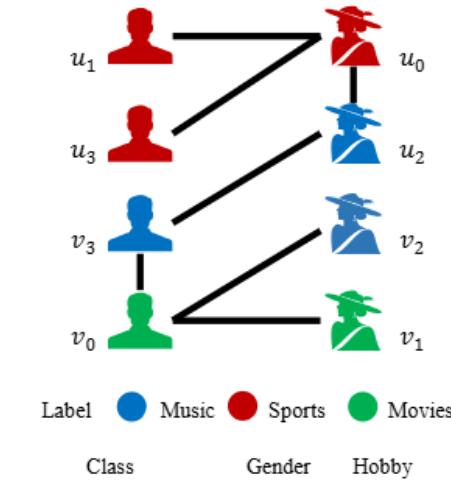
# ProNoG

- Motivation

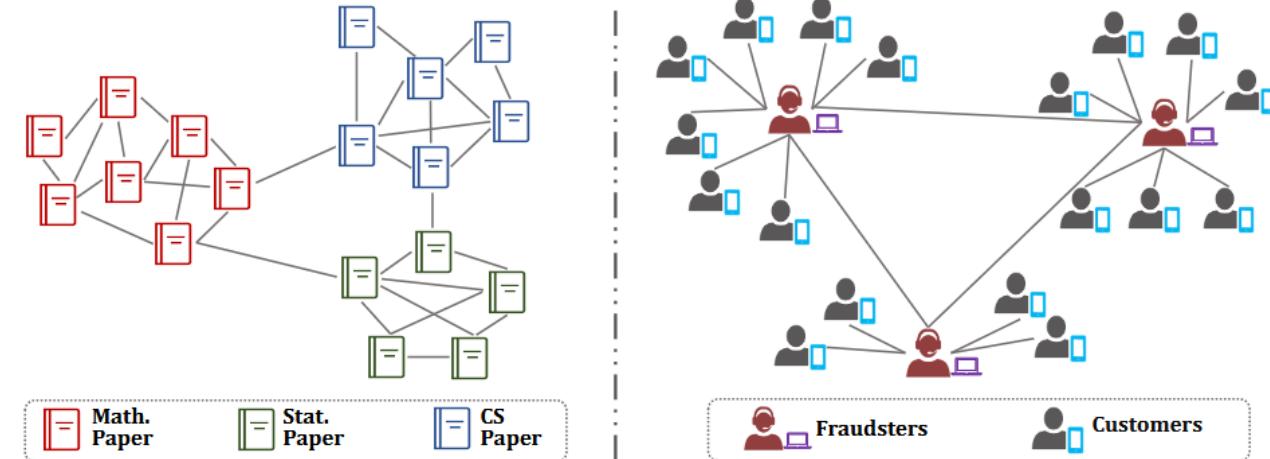
- Graphs exhibit different homophily ratio depending on nodes label
- How to capture node specific homophily pattern?



(a) Varying non-homophilic patterns across different graphs

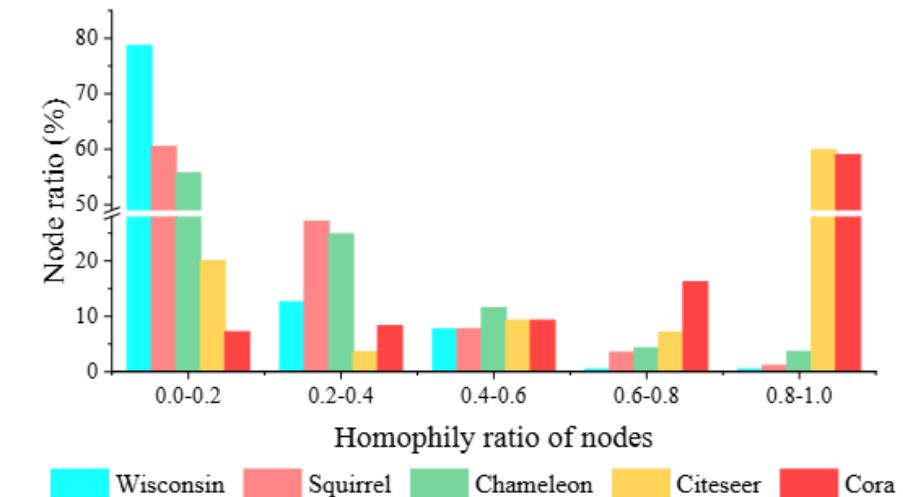


(b) Dependence of homophily ratio on the target label



Homophily graph

Heterophily graph



(c) Diverse non-homophilic patterns across nodes in the same graph

Yu, et al. "Non-homophilic graph pre-training and prompt learning." SIGKDD'25.

# ProNoG

## Contrastive pre-training method loss function

$$\mathcal{L}_T = - \sum_{u \in V} \ln P(u, \mathcal{A}_u, \mathcal{B}_u), \quad (4)$$

$$P(u, \mathcal{A}_u, \mathcal{B}_u) \triangleq \frac{\sum_{a \in \mathcal{A}_u} \text{sim}(\mathbf{h}_u, \mathbf{h}_a)}{\sum_{a \in \mathcal{A}_u} \text{sim}(\mathbf{h}_u, \mathbf{h}_a) + \sum_{b \in \mathcal{B}_u} \text{sim}(\mathbf{h}_u, \mathbf{h}_b)}, \quad (5)$$

## Theorems

**THEOREM 1.** *For a homophily task  $T$ , adding a homophily sample always results in a smaller loss than adding a non-homophily sample.*

**THEOREM 2.** *Consider a graph  $G = (V, E)$  with a label mapping function  $V \rightarrow Y$ , and let  $y_v \in Y$  denote the label mapped to  $v \in V$ . Suppose the label mapping satisfies that*

$$\forall u, a, b \in V, y_u = y_a \wedge y_u \neq y_b \Rightarrow \text{sim}(u, a) > \text{sim}(u, b).$$

*Let  $\mathbb{E}_T$  denote the expected number of homophily samples for a homophily task  $T$  on the graph  $G$ . Then,  $\mathbb{E}_T$  increases monotonically as the homophily ratio  $\mathcal{H}(G)$  defined w.r.t.  $Y$  increases.*

## Definition of homophily task

**DEFINITION 1 (HOMOPHILY TASK).** *On a graph  $G = (V, E)$ , a pre-training task  $T = (\{\mathcal{A}_u : u \in V\}, \{\mathcal{B}_u : u \in V\})$  is a homophily task if and only if,  $\forall u \in V, \forall a \in \mathcal{A}_u, \forall b \in \mathcal{B}_u, (u, a) \in E \wedge (u, b) \notin E$ . A task that is not a homophily task is called a non-homophily task.  $\square$*

## Insights

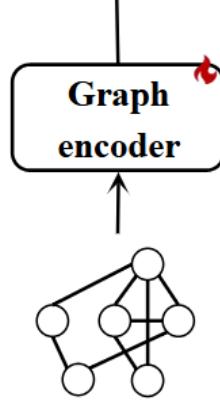
For non-homophilic graphs, especially those with low homophily ratio, non-homophily tasks are a better choice compared to homophily tasks when optimizing the training loss.

**Table 6: Positive and negative samples for homophily and non-homophily methods.**

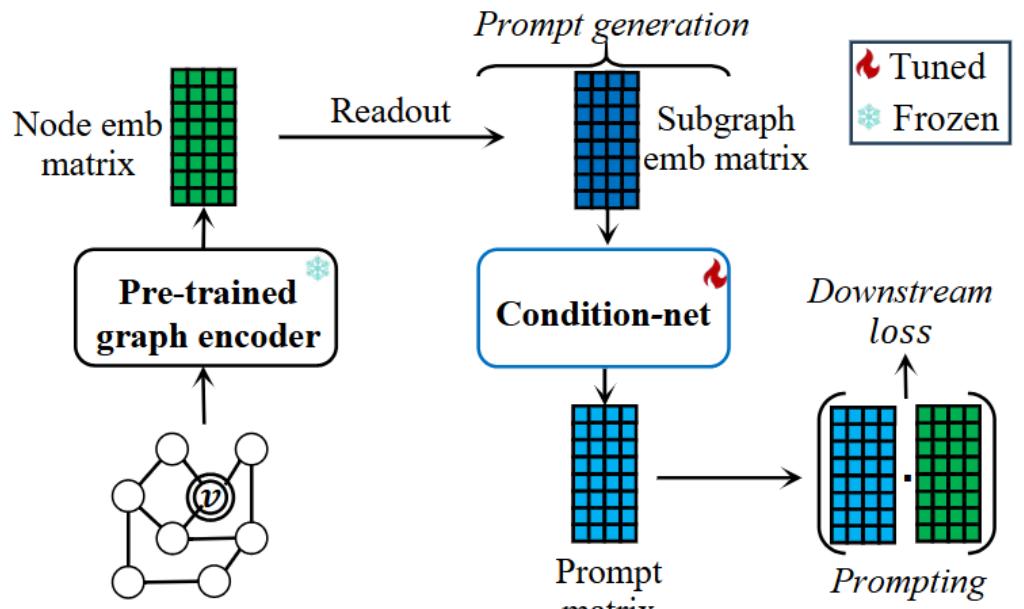
Pre-training task	Positive instances $\mathcal{A}_u$	Negative instances $\mathcal{B}_u$	Homophily task
Link prediction [26, 62, 64] DGI [48]	a node connected to node $u$ nodes in graph $G$	nodes disconnected to node $u$ nodes in corrupted graph $G'$	Yes No
GraphCL [60]	an augmented graph from graph $G$	augmented graphs from $G' \neq G$	No
GraphACL [55]	nodes with similar ego-subgraph to node $u$	nodes with dissimilar ego-subgraph to node $u$	No

# ProNoG

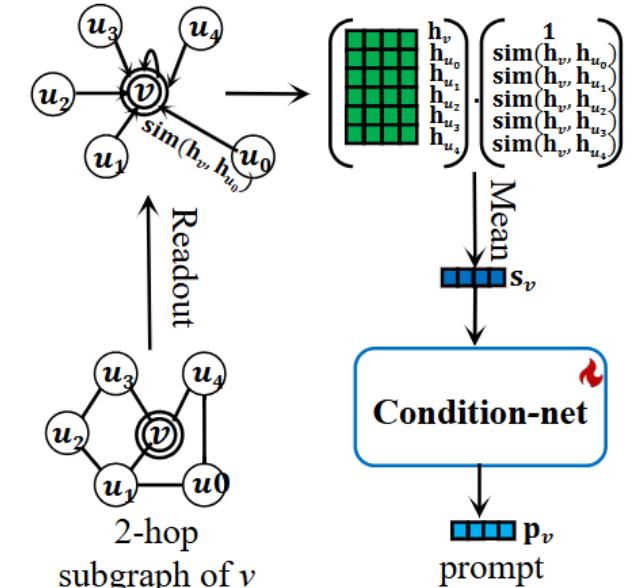
*Non-homophily pre-training loss*



(a) Pre-training



(b) Downstream adaptation with conditional prompting



(c) Details of prompt generation

Figure 2: Overall framework of ProNoG.

## Prompt generation

$$s_v = \frac{1}{|S_v|} \sum_{u \in S_v} \mathbf{h}_u \cdot \text{sim}(\mathbf{h}_u, \mathbf{h}_v),$$

$$\mathbf{p}_{t,v} = \text{CondNet}(s_v; \phi_t),$$

## Prompt tuning

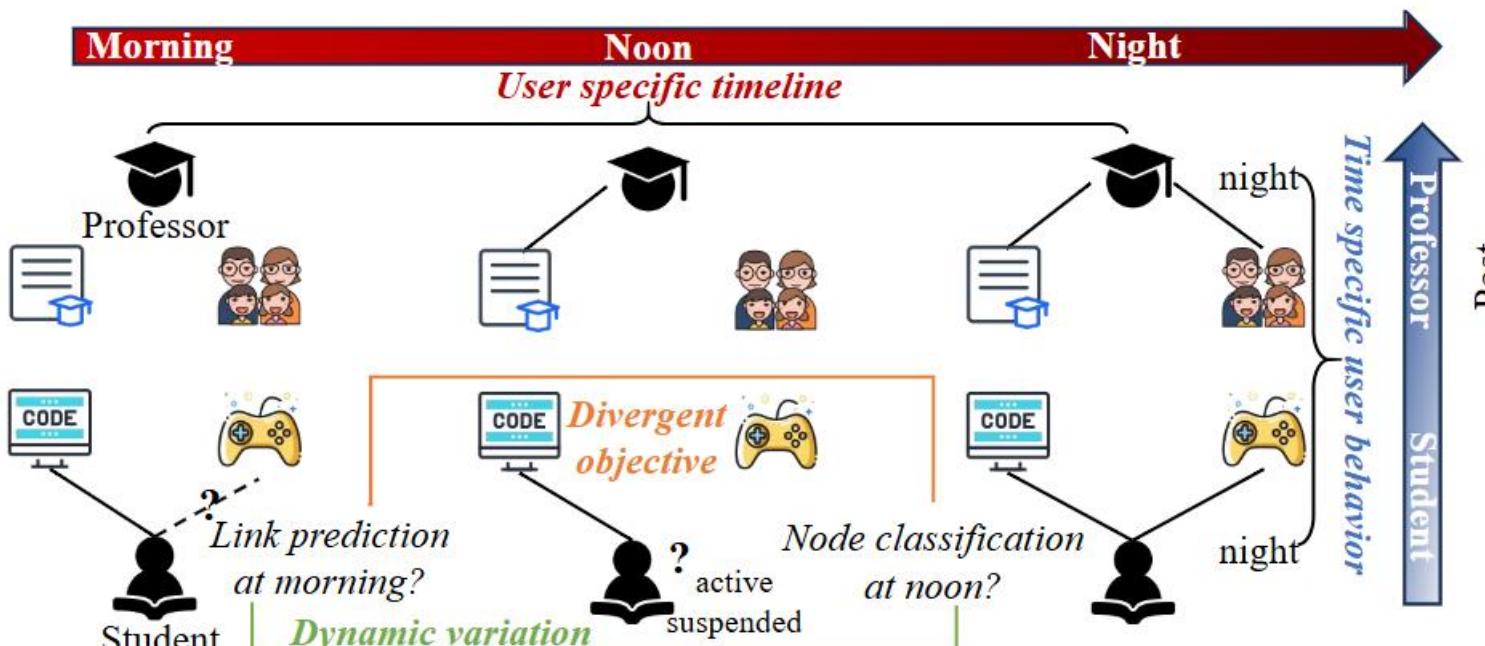
$$\tilde{\mathbf{h}}_{t,v} = \mathbf{p}_{t,v} \odot \mathbf{h}_v,$$

$$\mathcal{L}_{\text{down}}(\phi_t) = - \sum_{(x_i, y_i) \in \mathcal{D}_t} \ln \frac{\exp\left(\frac{1}{\tau} \text{sim}(\tilde{\mathbf{h}}_{t,x_i}, \tilde{\mathbf{h}}_{t,y_i})\right)}{\sum_{c \in Y} \exp\left(\frac{1}{\tau} \text{sim}(\tilde{\mathbf{h}}_{t,x_i}, \tilde{\mathbf{h}}_{t,c})\right)},$$

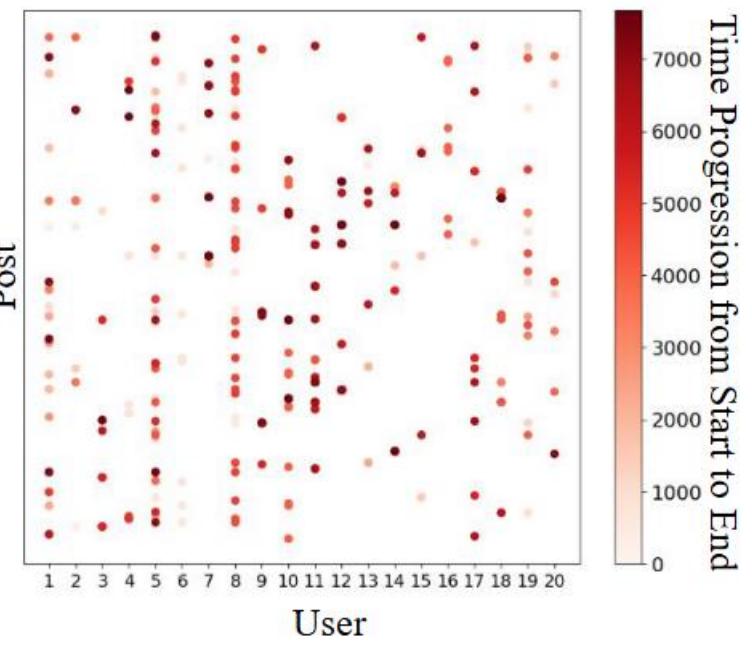
Yu, et al. "Non-homophilic graph pre-training and prompt learning." SIGKDD'25.

# DyGPrompt

- Motivation
  - How to design bridge temporal variations across time and different task objectives
  - How to capture evolving patterns across different nodes and time points

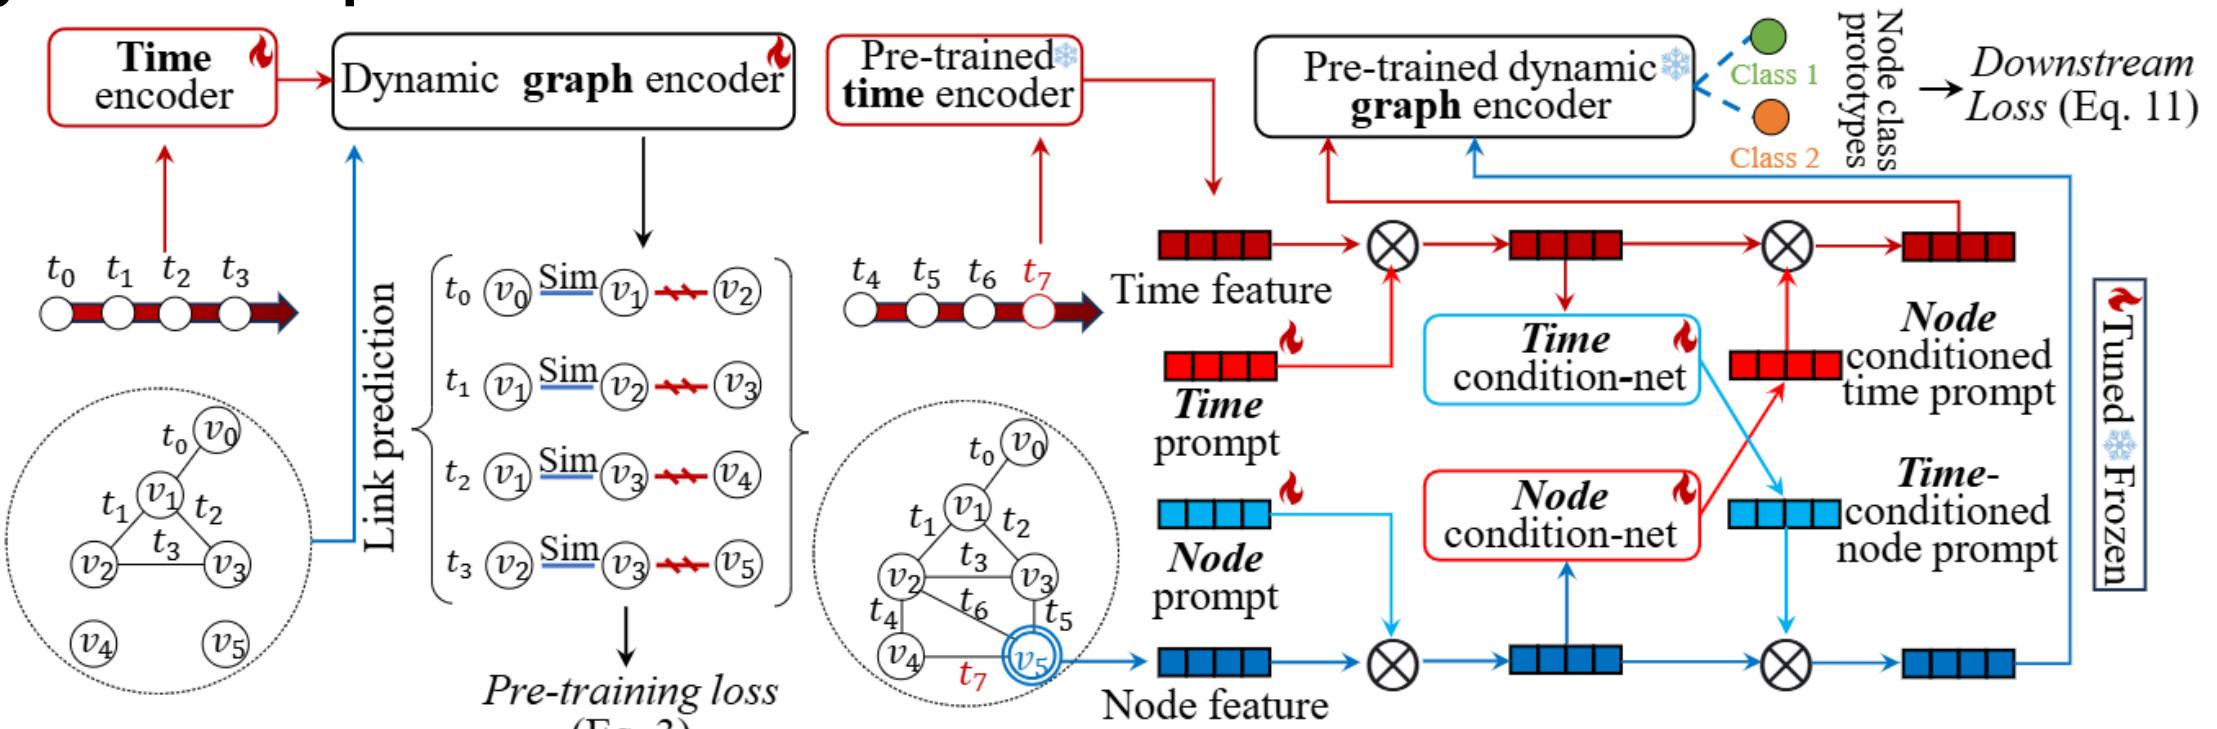


(a) Users' comments over time



(b) Envolving node-time patterns

# DyGPrompt



Node prompt

$$\mathbf{x}_{t,v}^{\text{node}} = \mathbf{p}^{\text{node}} \odot \mathbf{x}_{t,v}$$

Time prompt

$$\mathbf{f}_t^{\text{time}} = \mathbf{p}^{\text{time}} \odot \mathbf{f}_t$$

Time conditioned node prompts

$$\tilde{\mathbf{p}}_t^{\text{node}} = \text{TCN}(\mathbf{f}_t^{\text{time}}; \kappa)$$

$$\tilde{\mathbf{x}}_{t,v}^{\text{node}} = \tilde{\mathbf{p}}_t^{\text{node}} \odot \mathbf{x}_{t,v}^{\text{node}}$$

Node conditioned time prompts

$$\tilde{\mathbf{p}}_{t,v}^{\text{time}} = \text{NCN}(\mathbf{x}_{t,v}^{\text{node}}; \phi)$$

$$\tilde{\mathbf{f}}_{t,v}^{\text{time}} = \tilde{\mathbf{p}}_{t,v}^{\text{time}} \odot \mathbf{f}_t^{\text{time}}$$

- Motivation
  - Chain-of-Thought (CoT) mimics the logical process a person may employ to solve a task
  - Current graph models inference in single step
  - *Would introducing additional inference steps in a CoT style enhance the ability of pre-trained graph models to refine their predictions*

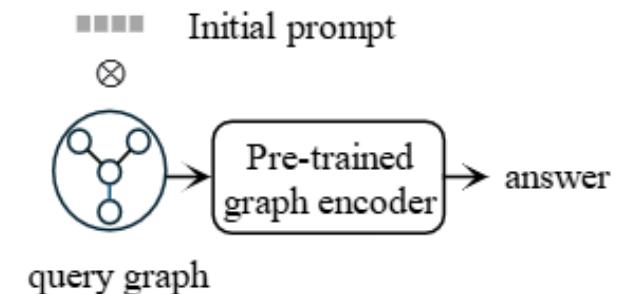
Query: If there are 9 apples, I eat 5 apples, how many apples are left?

Answer: There are 9 apples initially. 5 are eaten. Now there are  $9 - 5 = 4$  apples. The answer is 4.

Query: If there are 2 peaches, I buy 3 peaches, how many peaches are there?

Answer: ...

(a) Chain-of-Thought prompting in NLP



(b) Standard graph prompting

- CoT in natural language processing
  - Handcrafted before the learning phase and typically consists of a structured text in the form  $\langle \text{input}, \text{chain of thought}, \text{output} \rangle$ .
  - Prompt serves as an example to guide the model in generating intermediate thoughts that lead to the final answer.
- CoT for graphs
  - What should be the inference steps and thoughts for a graph task?
  - How can we leverage a “thought” to learn prompts and guide the next-step inference

# GCoT

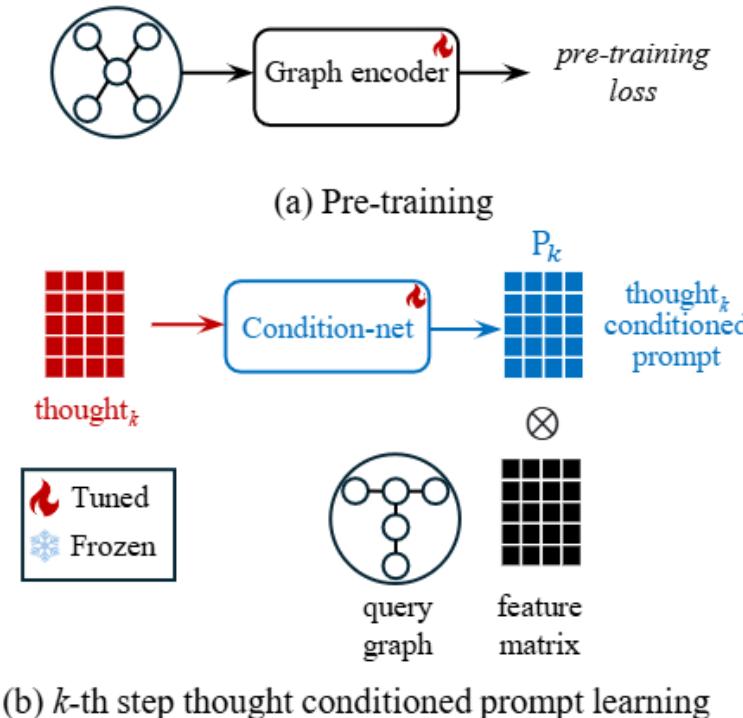


Figure 2: Overall framework of GCoT.

## Prompt based inference

$$\{\mathbf{H}_k^1, \mathbf{H}_k^2, \dots, \mathbf{H}_k^L\} = \text{GRAPHENCODER}(\mathbf{X}_k, \mathbf{G}; \Theta_0)$$

## Thought construction

$$\mathbf{T}_k = \text{Fuse}(\mathbf{H}_k^1, \mathbf{H}_k^2, \dots, \mathbf{H}_k^L),$$

Yu, et al. "GCoT: Chain-of-Thought Prompt Learning for Graphs." arXiv preprint.

## Thought conditioned prompting

$$\begin{aligned} \mathbf{P}_k &= \text{CondNet}(\mathbf{T}_k; \phi), \\ \mathbf{X}_{k+1} &= \mathbf{P}_k \odot \mathbf{X}, \end{aligned}$$

# MultiGPrompt

- Motivation
  - How to leverage diverse pretext tasks for graph models in a synergistic manner?
  - How to transfer both task specific and global pre-trained knowledge to downstream tasks?

# MultiGPrompt

## Multi-task pre-training

Pretext tokens

$$\mathcal{T}_{\langle k \rangle} = \{t_{\langle k \rangle,0}, t_{\langle k \rangle,1}, \dots, t_{\langle k \rangle,L}\}$$

Add token to each layer of graph encoder

$$H^{l+1} = MP(t_{\langle k \rangle,l} \odot H^l, A; \theta^l)$$

Graph encoder output embedding

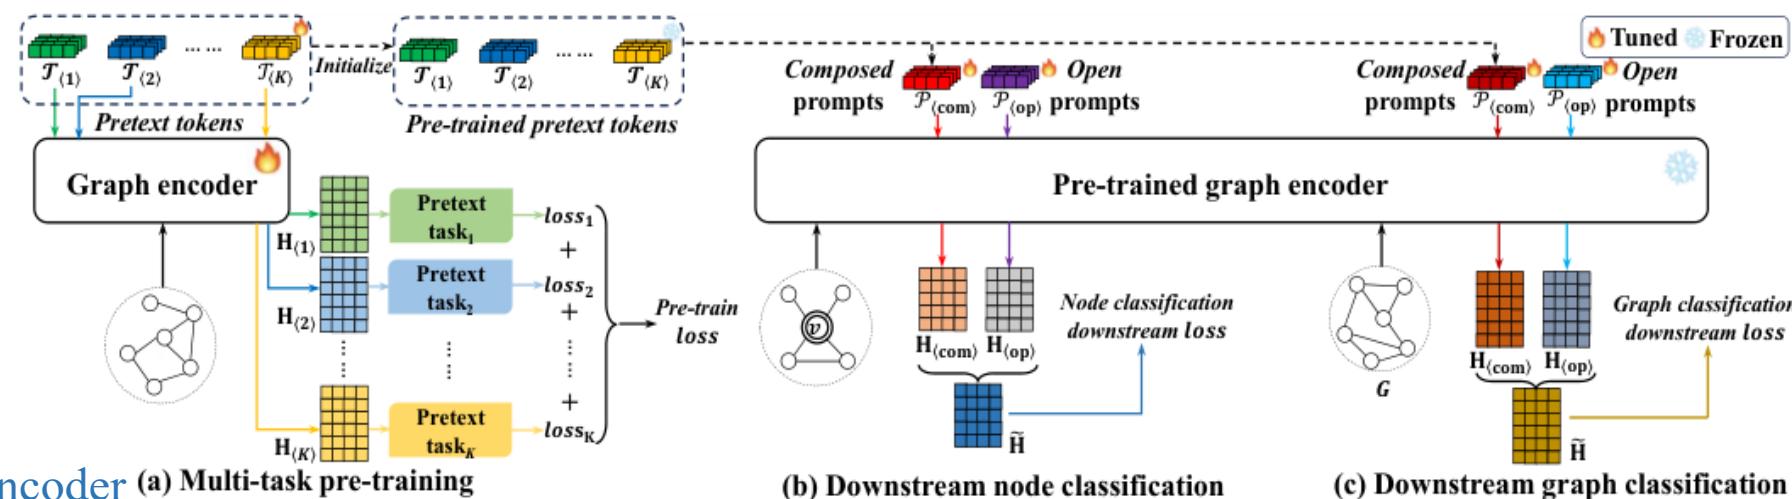
$$H_t = \text{GRAPHENCODER}_t(X, A; \Theta)$$

Overall embedding

$$H_{\langle k \rangle} = \sum_{l=0}^L \alpha_l H_{t_{\langle k \rangle,l}}$$

Pre-Training Objective

$$\mathcal{L}_{\text{pre}}(\mathcal{H}; \mathcal{T}, \Theta) = \sum_{k=1}^K \beta_k \mathcal{L}_{\text{pre}_{\langle k \rangle}}(H_{\langle k \rangle}; \mathcal{T}_{\langle k \rangle}, \Theta),$$



(a) Multi-task pre-training

(b) Downstream node classification

(c) Downstream graph classification

Prompt tuning

Composed prompt

$$\mathcal{P}_{\langle \text{com} \rangle} = \{p_{\langle \text{com} \rangle,0}, p_{\langle \text{com} \rangle,1}, \dots, p_{\langle \text{com} \rangle,L}\}$$

$$p_{\langle \text{com} \rangle,l} = \text{COMPOSE}(t_{\langle 1 \rangle,l}, t_{\langle 2 \rangle,l}, \dots, t_{\langle K \rangle,l}; \Gamma).$$

Open prompt

$$\mathcal{P}_{\langle \text{op} \rangle} = \{p_{\langle \text{op} \rangle,0}, p_{\langle \text{op} \rangle,1}, \dots, p_{\langle \text{op} \rangle,L}\}$$

Add prompt to each layer of graph encoder

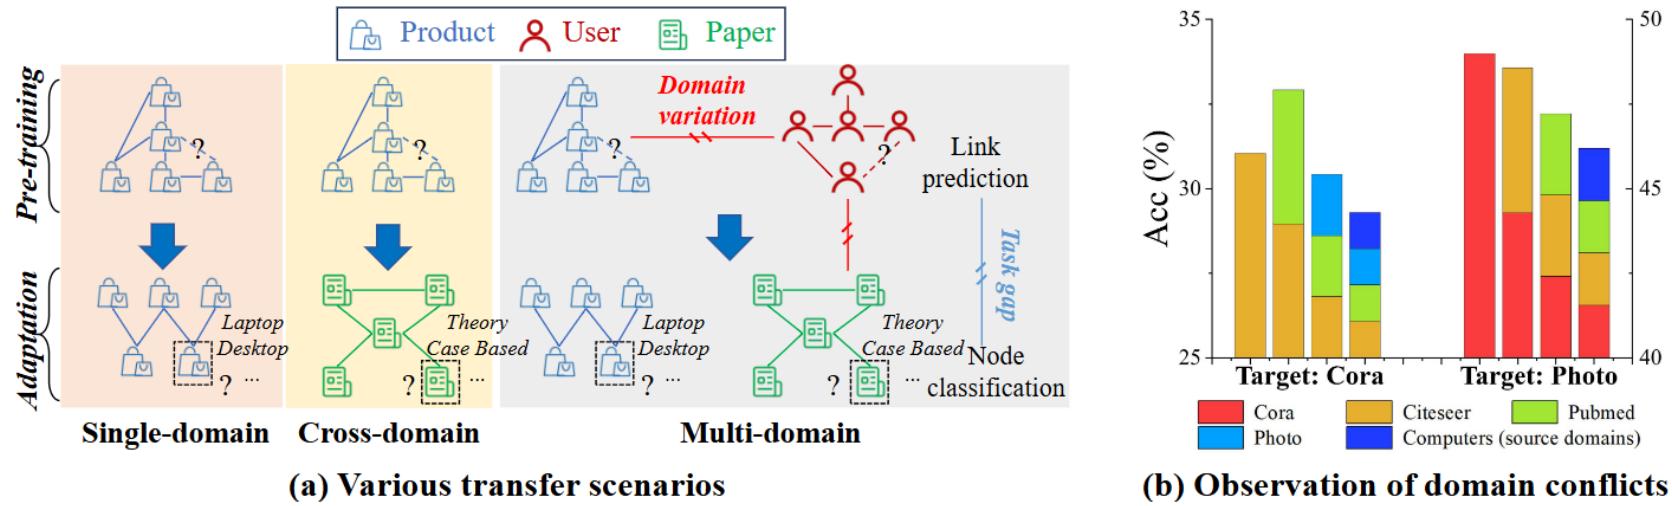
$$H_p = \text{GRAPHENCODER}_p(X, A; \Theta_{\text{pre}})$$

Aggregate dual prompt

$$\tilde{H} = \text{AGGR}(H_{\langle \text{com} \rangle}, H_{\langle \text{op} \rangle}; \Delta)$$

# MDGPT & SAMGPT

- Motivation



- How to align multi-domain graphs in the pre-training phase in both feature and structure level
- How to adapt multi-domain prior knowledge to downstream tasks in different domains?

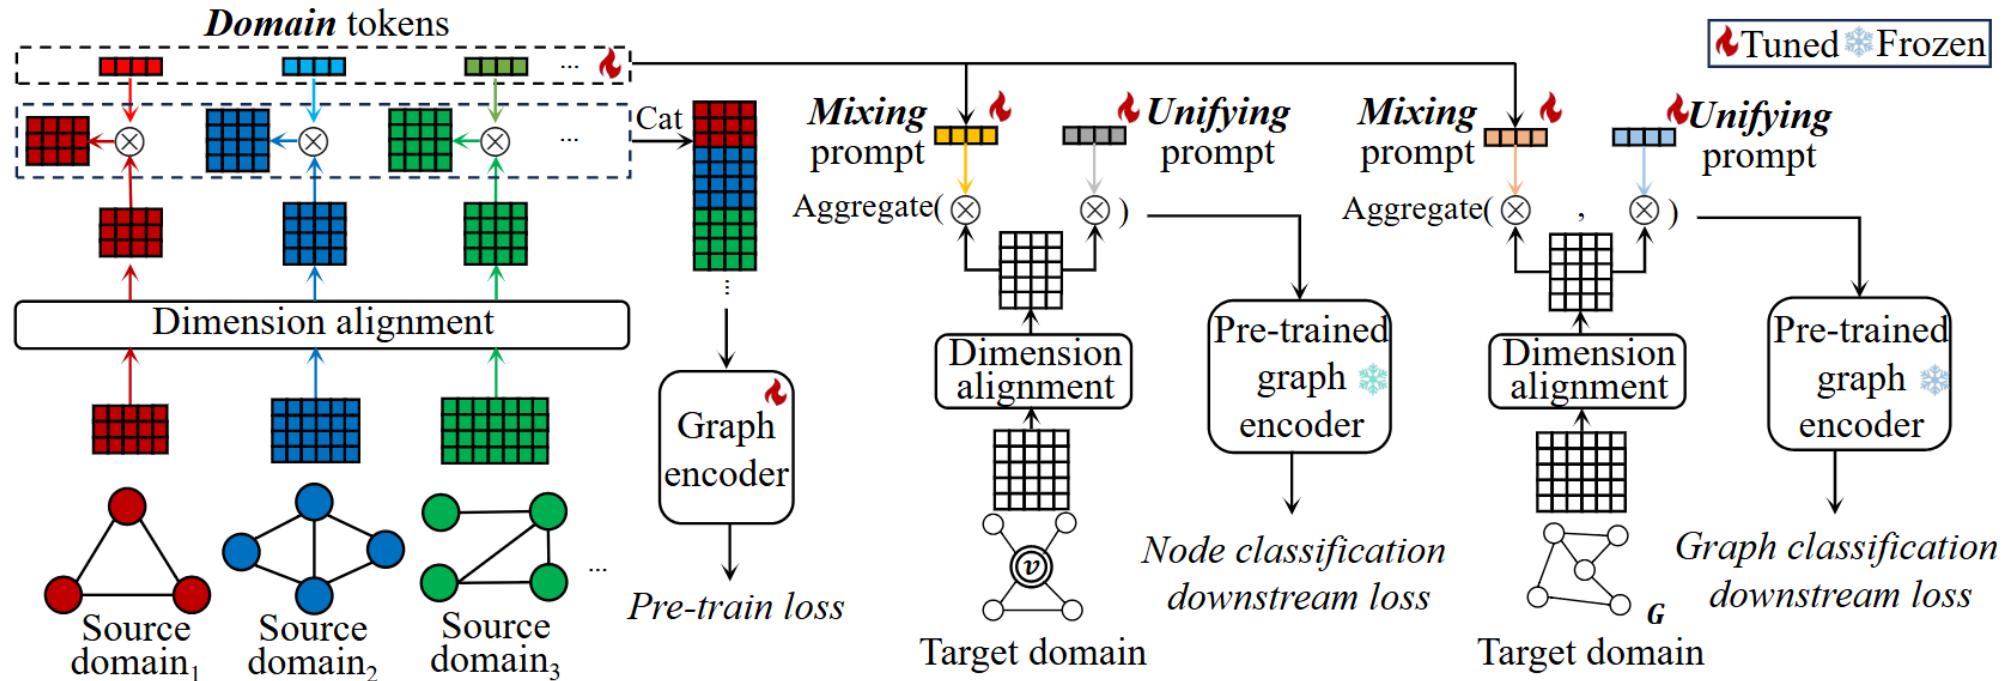
	Nodes	Edges	Feature dimension	Node classes	Avg. nd	Avg. spl	Avg. cc
Cora	2,708	10,556	1,433	7	3.89	6.30	0.24
Citeseer	3,327	9,104	3,703	6	2.73	9.31	0.14
Pubmed	19,717	88,648	500	3	4.49	6.33	0.06
Photo	7,650	238,162	745	8	31.13	4.05	0.40
Computers	13,752	491,722	767	10	35.75	3.38	0.34
Facebook	22,470	342,004	128	4	15.22	4.97	0.35
LastFM	7,624	55,612	128	18	7.29	5.23	0.21

nd: node degree, spl: shortest path length [3], cc: clustering coefficient [8].

Yu, et al. "Text-free multi-domain graph pre-training: Toward graph foundation models." arXiv preprint.

Yu, et al. "SAMGPT: Text-free Graph Foundation Model for Multi-domain Pre-training and Cross-domain Adaptation." WWW'25.

# MDGPT



## Multi-domain pre-training

### Dimension alignment

$$\tilde{\mathbf{X}}_i = \text{DA}_{S_i}(\mathbf{X}_i)$$

$$\text{DA}_{S_i}: \mathbb{R}^{|V| \times d_{S_i}} \rightarrow \mathbb{R}^{|V| \times \tilde{d}}$$

### Semantic alignment

$$\hat{\mathbf{X}}_i = \mathbf{t}_{S_i} \odot \tilde{\mathbf{X}}_i,$$

$$\mathbf{H}_S = \text{GE}(\mathcal{G}_S, \mathcal{X}_S; \Theta),$$

## Downstream adaptation

### Dimension alignment

$$\tilde{\mathbf{X}} = \text{DA}_T(\mathbf{X})$$

$$\mathbf{H} = \text{GE}(G, \mathbf{p}_{\text{uni}} \odot \tilde{\mathbf{X}}; \Theta_{\text{pre}}) + \text{GE}(G, \mathbf{p}_{\text{mix}} \odot \tilde{\mathbf{X}}; \Theta_{\text{pre}}),$$

### Unifying prompt

$$\mathbf{p}_{\text{uni}}$$

### Mixing prompt

$$\mathbf{p}_{\text{mix}} = \sum_{i=1}^K \gamma_i \mathbf{t}_{S_i}$$

Yu, et al. "Text-free multi-domain graph pre-training: Toward graph foundation models." arXiv preprint.

# SAMGPT

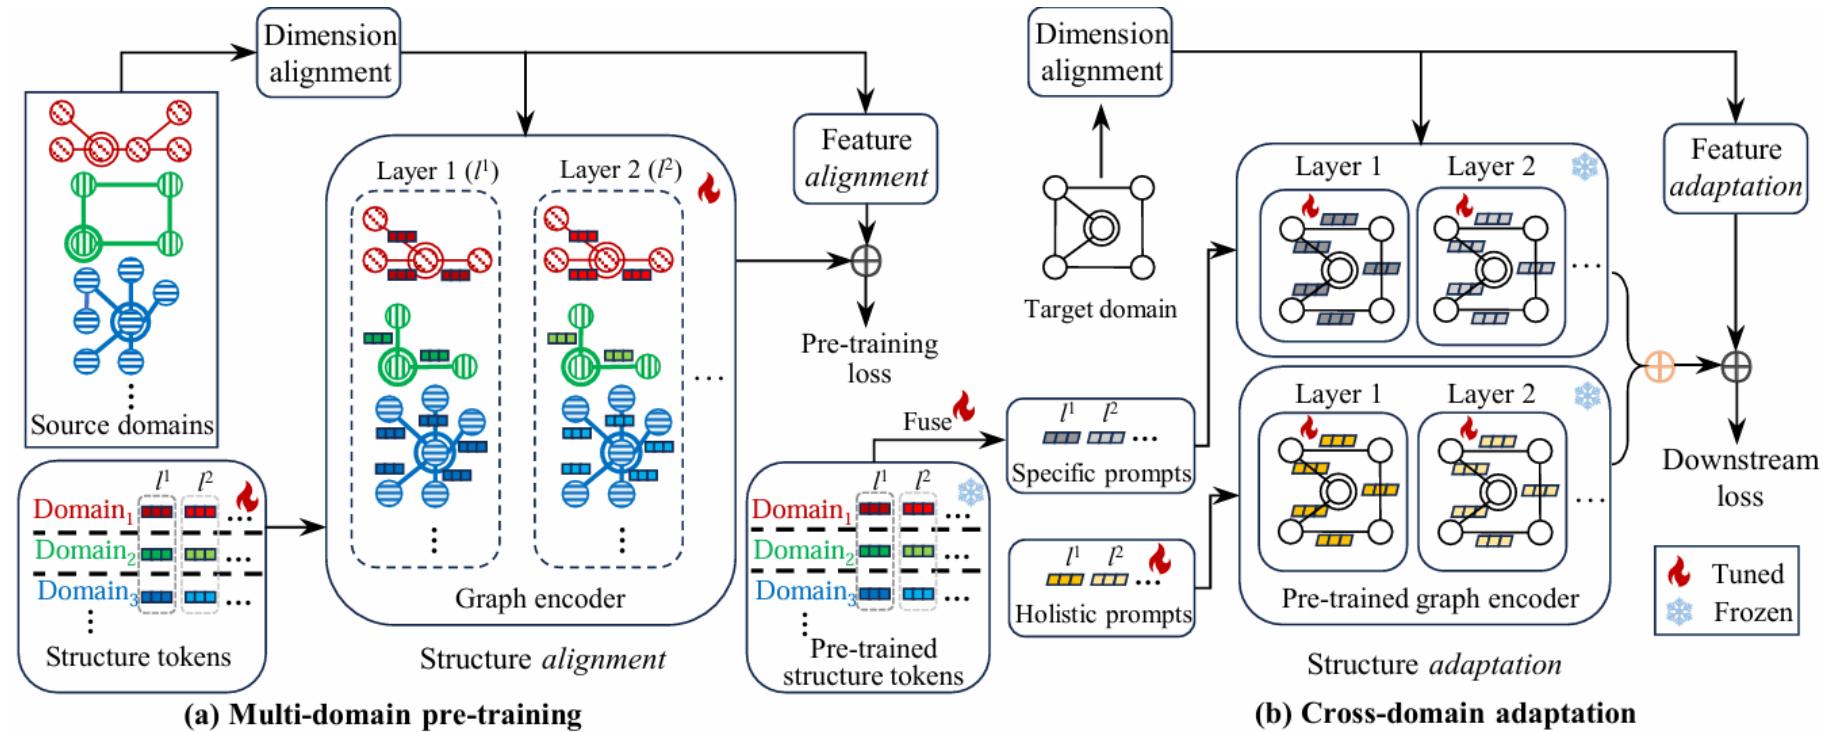


Figure 2: Overall framework of SAMGPT.

## Multi-domain pre-training

### Structural alignment

$$\mathbf{h}_v^l = \text{Aggr}(\mathbf{h}_v^{l-1}, \{\mathbf{t}_{S_i}^l \odot \mathbf{h}_u^{l-1} : u \in \mathcal{N}_v\}; \theta^l), \forall v \in V_i,$$

### Holistic prompt

$$\mathbf{h}_v^l = \text{Aggr}(\mathbf{h}_v^{l-1}, \{\mathbf{p}_{\text{hol}}^l \odot \mathbf{h}_u^{l-1} : u \in \mathcal{N}_v\}; \theta_{\text{pre}}^l), \mathbf{p}_{\text{spe}}^l = \sum_{i=1}^K \lambda_i^l \mathbf{t}_{S_i}^l,$$

## Downstream adaptation

### Specific prompt

Yu, et al. "SAMGPT: Text-free Graph Foundation Model for Multi-domain Pre-training and Cross-domain Adaptation." WWW'25.

# Summary

- Existing research often focus on **text-free** graphs, fail to leverage the vast amount of **textual data** to learn a more comprehensive knowledge
- LLMs have achieved significant performance

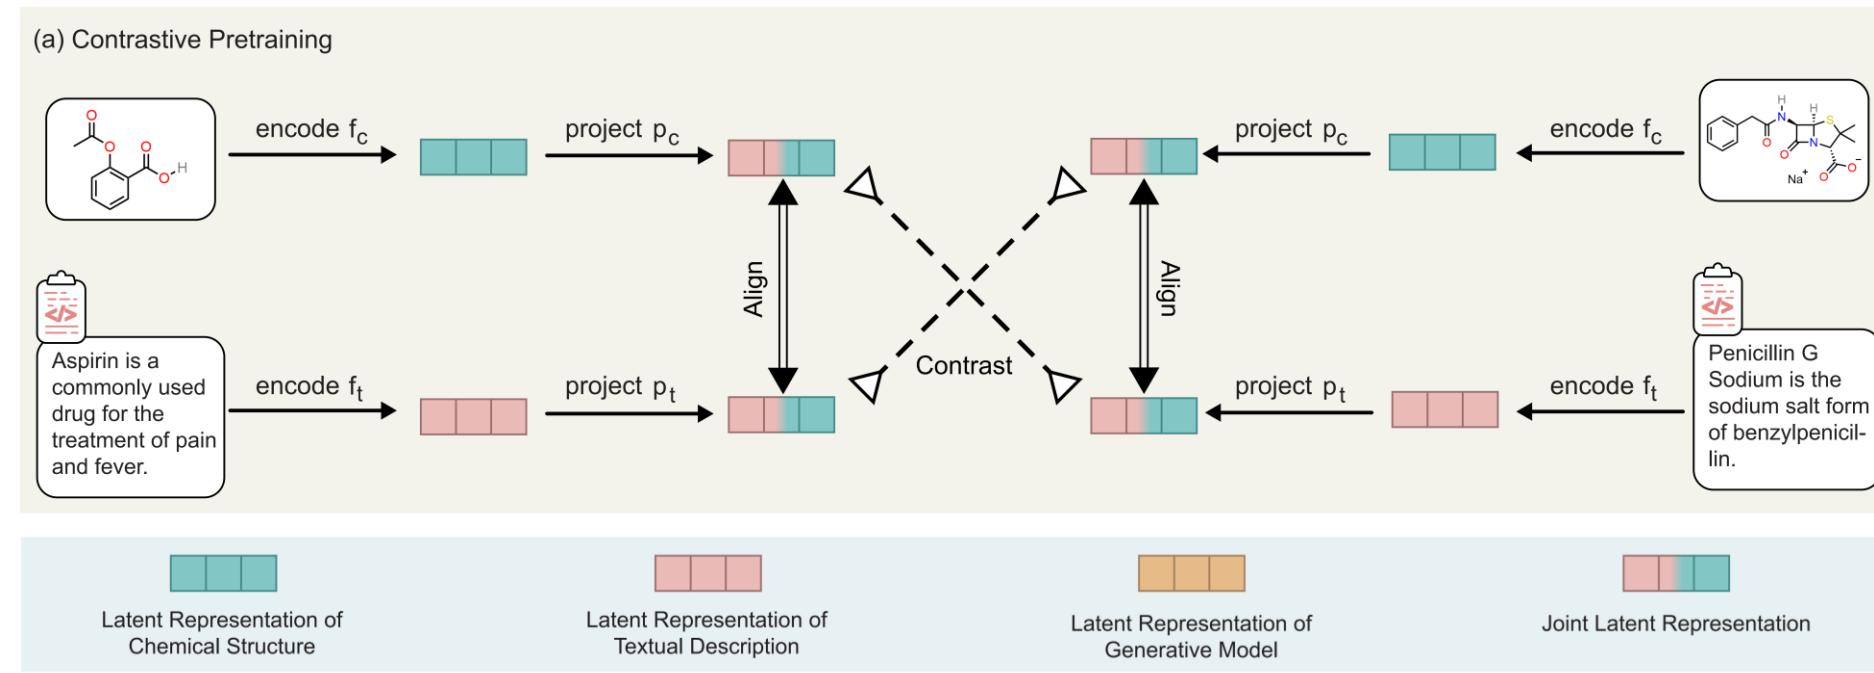
Can we leverage LLMs to integrate textual data and thereby improve the performance of graph few-shot learning?

# Graph + PLM

- Pre-training:
  - Contrastive pre-training
  - Language modeling
- Adaptation:
  - Prompt-tuning
  - Parameter-efficient fine-tuning (PEFT)

# Graph + PLM: Pre-training

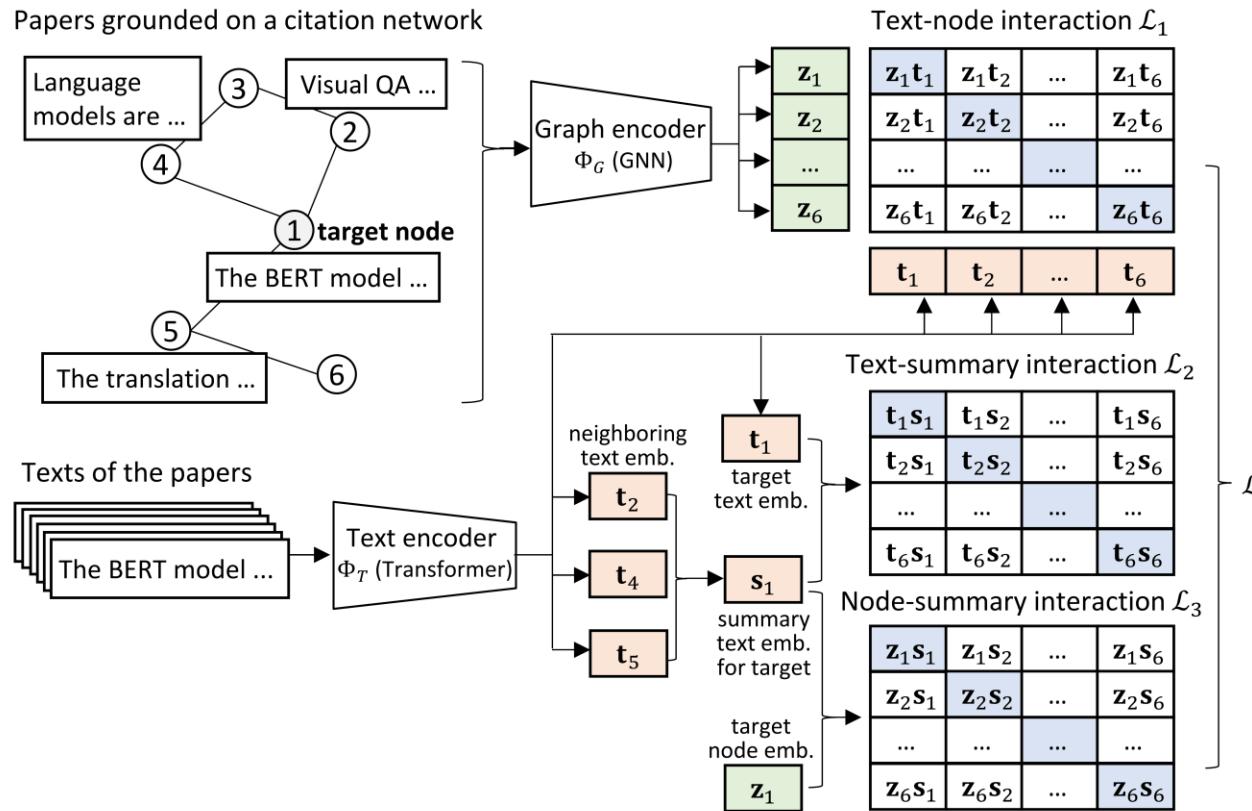
- MoleculeSTM: Contrastive pre-training



Graph-Text contrastive learning: Graph encoder + Text encoder → projector layers

# Graph + PLM: Pre-training

- G2P2: Contrastive pre-training



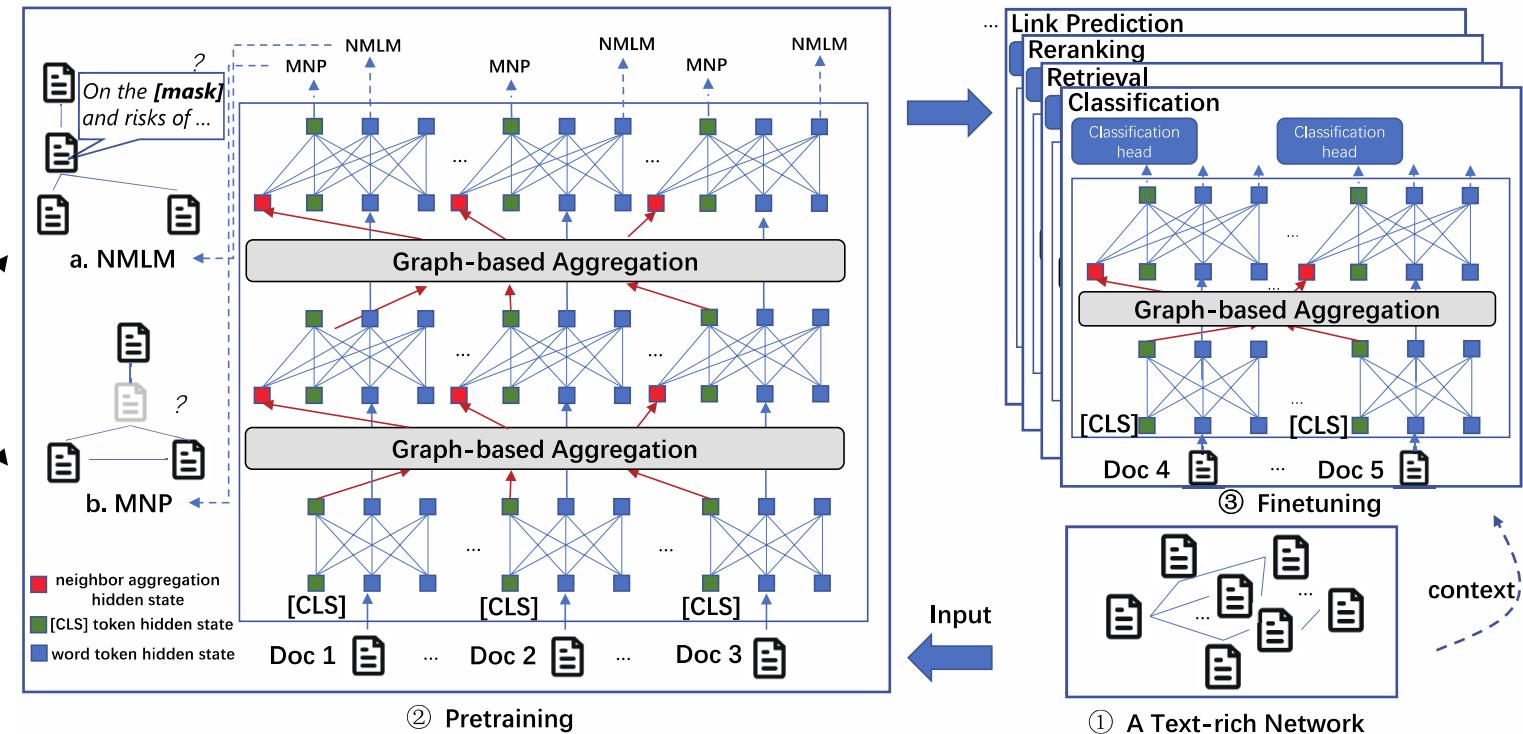
Graph encoder + Text encoder →  
Three contrastive loss:

- Text-Node
- Text summary-Text
- Text summary-Node

# Graph + PLM: Pre-training

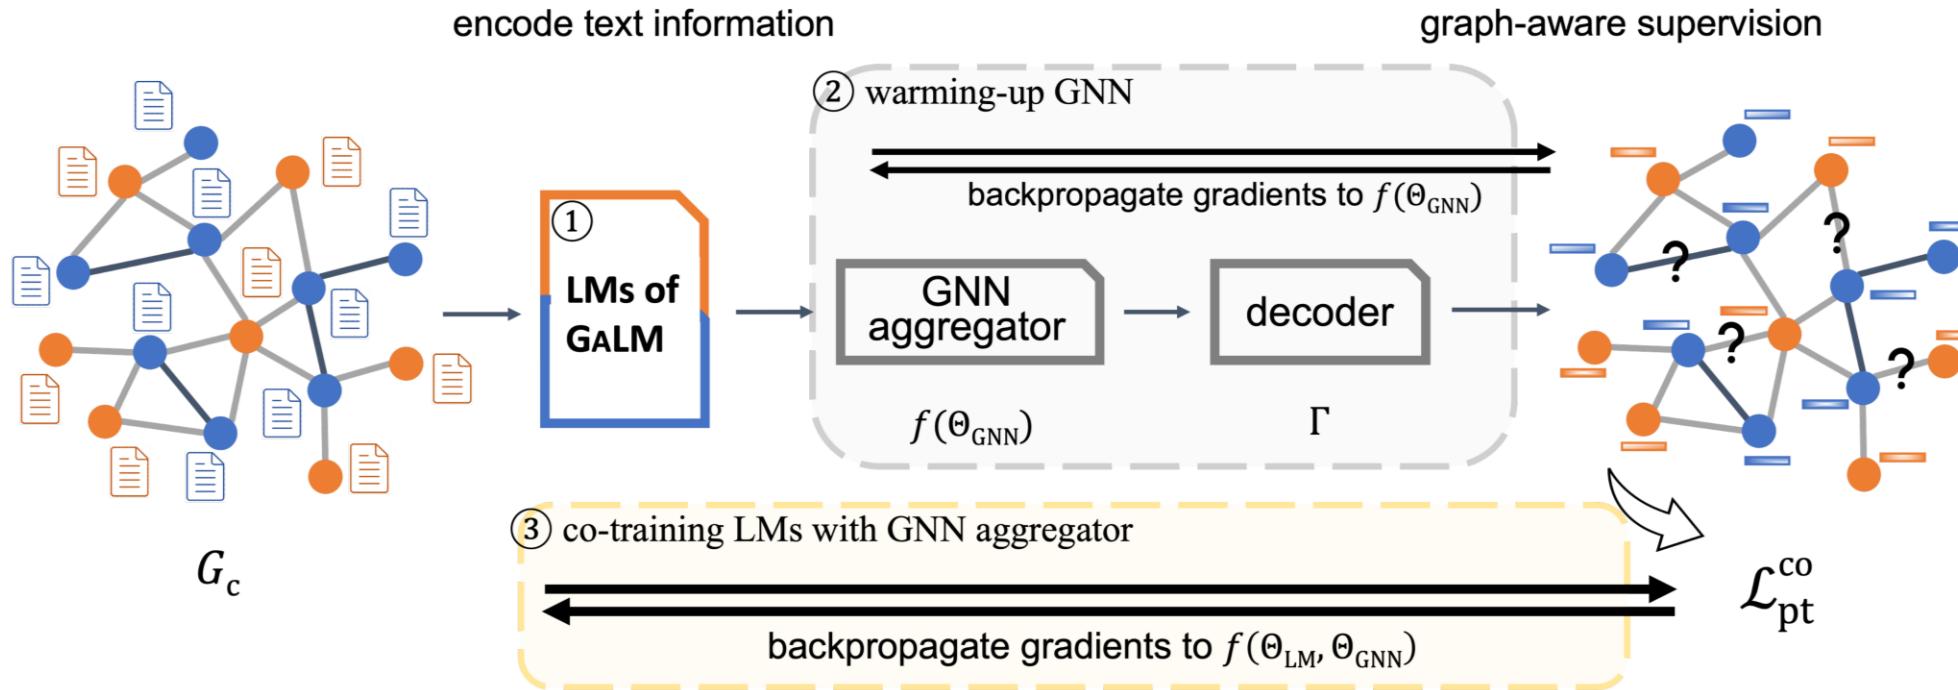
- PATTON: Masked language modeling + Masked node prediction

Two pretraining strategies



# Graph + PLM: Pre-training

- GaLM: Graph-aware language model pre-training



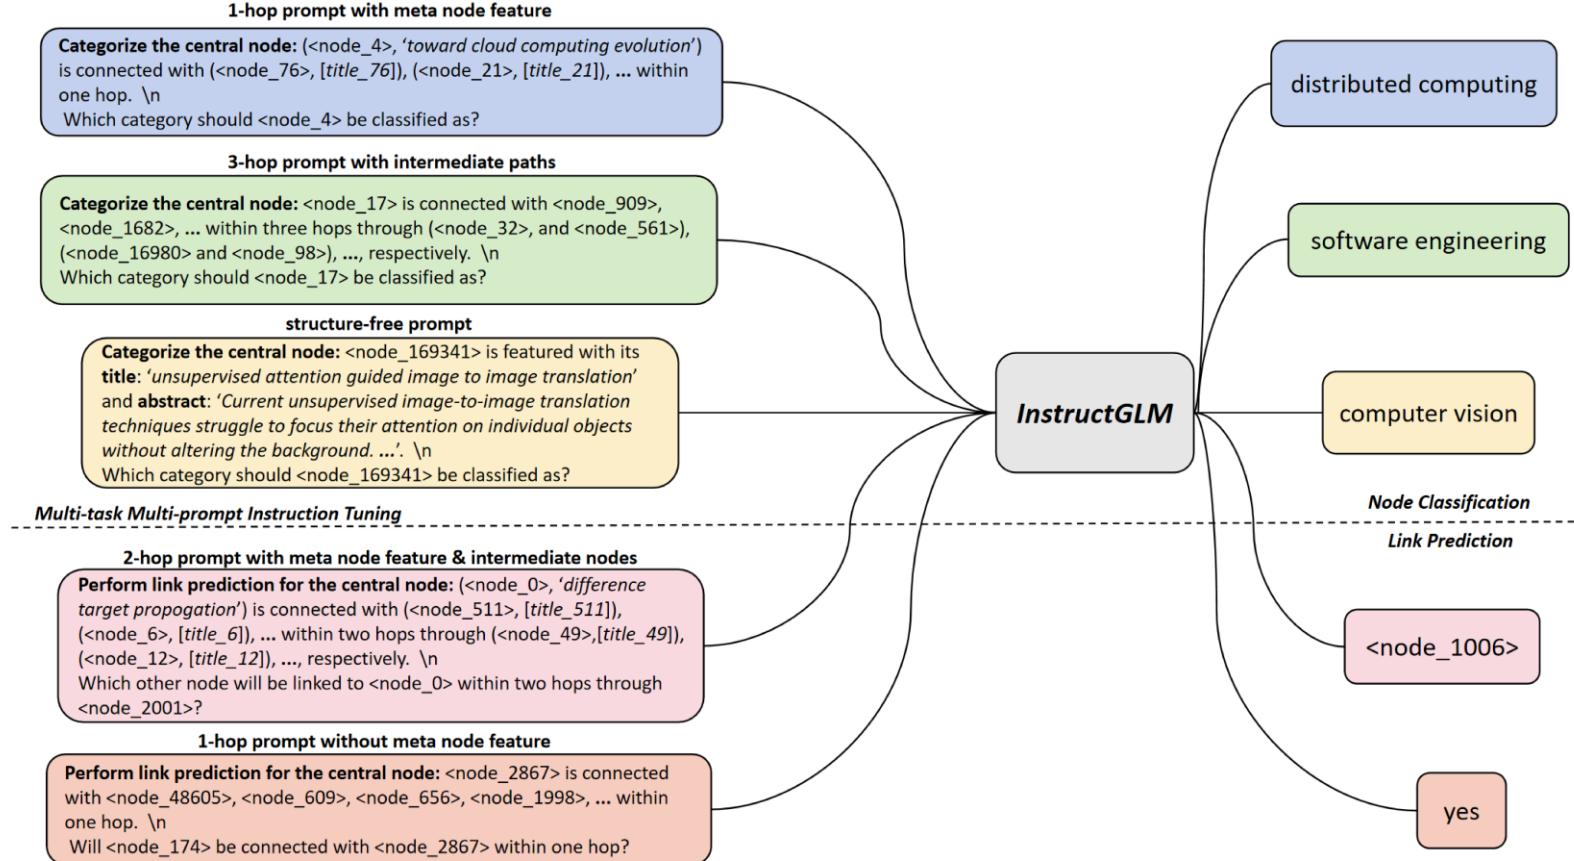
LLM-encoded node embeddings → Graph encoder

Link prediction task → Pre-train both the LLM and the graph encoder

# Graph + PLM: Pre-training

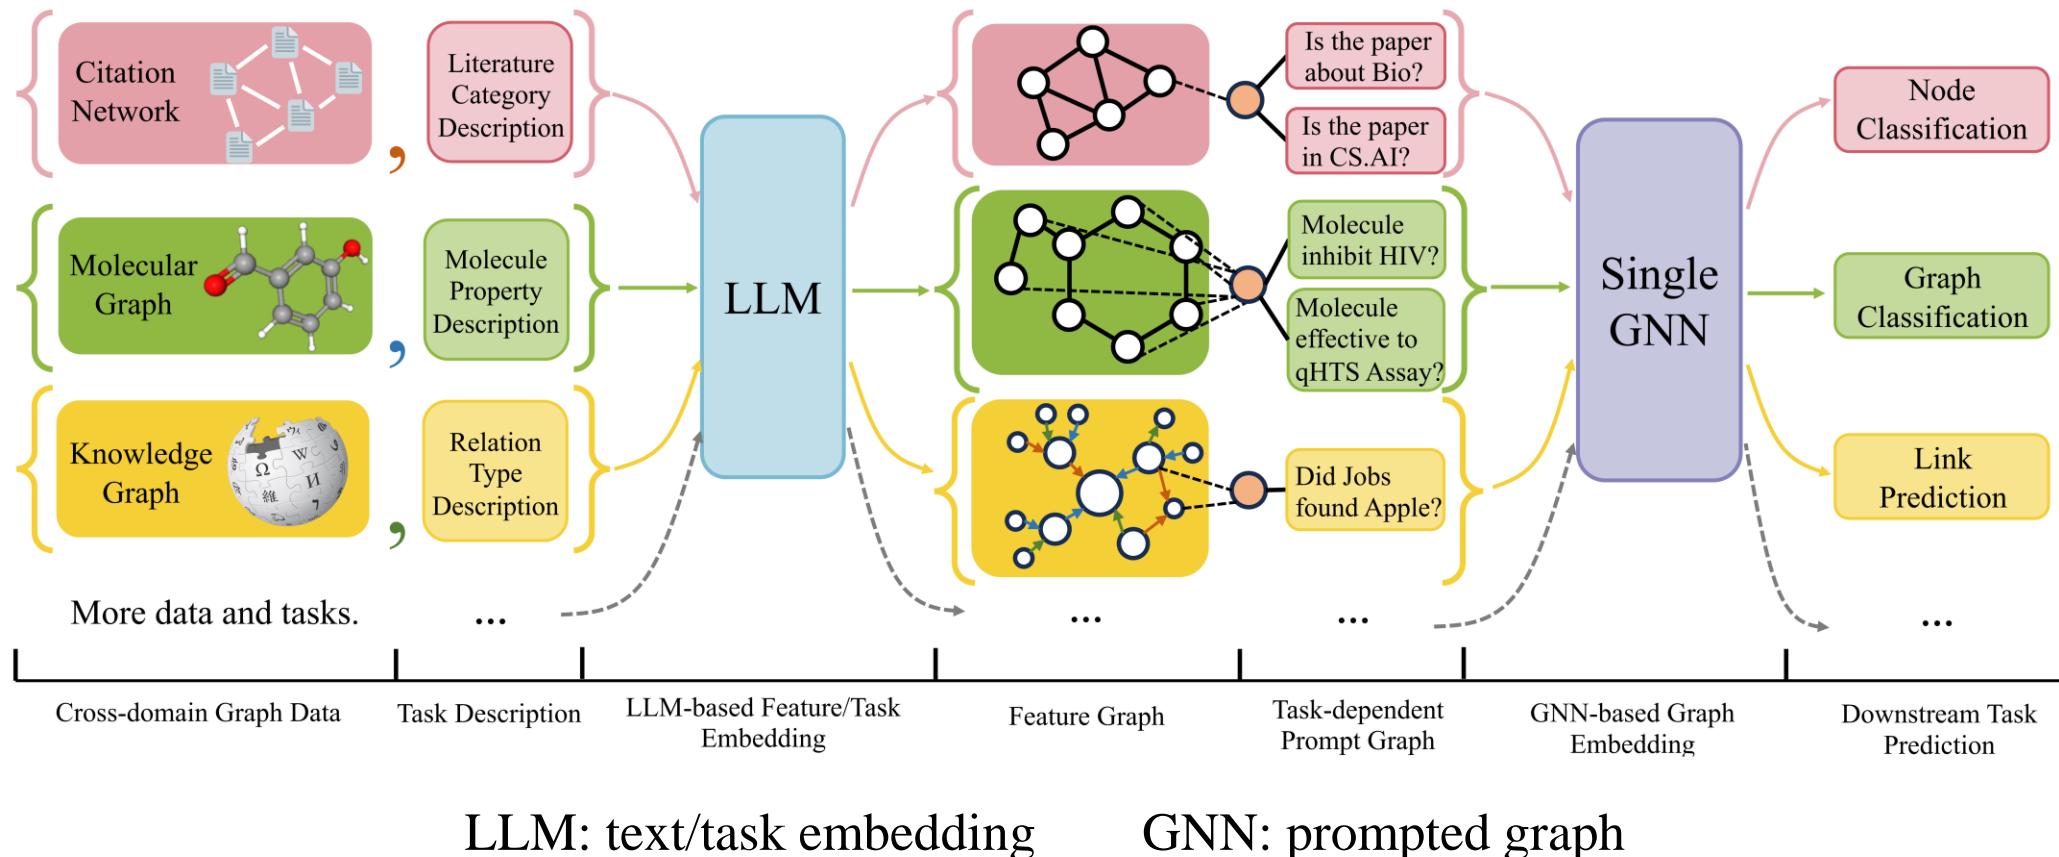
- InstructGLM: Language model pre-training

Natural language →  
Describe graph structures



# Graph + PLM: Pre-training

- One for all: LLM + GNN pre-training



H. Liu, et al. "One for all: Towards training one graph model for all classification tasks." ICLR'24.

# Graph + PLM: Adaptation

- Prompt-tuning
- Parameter-efficient fine-tuning (PEFT)

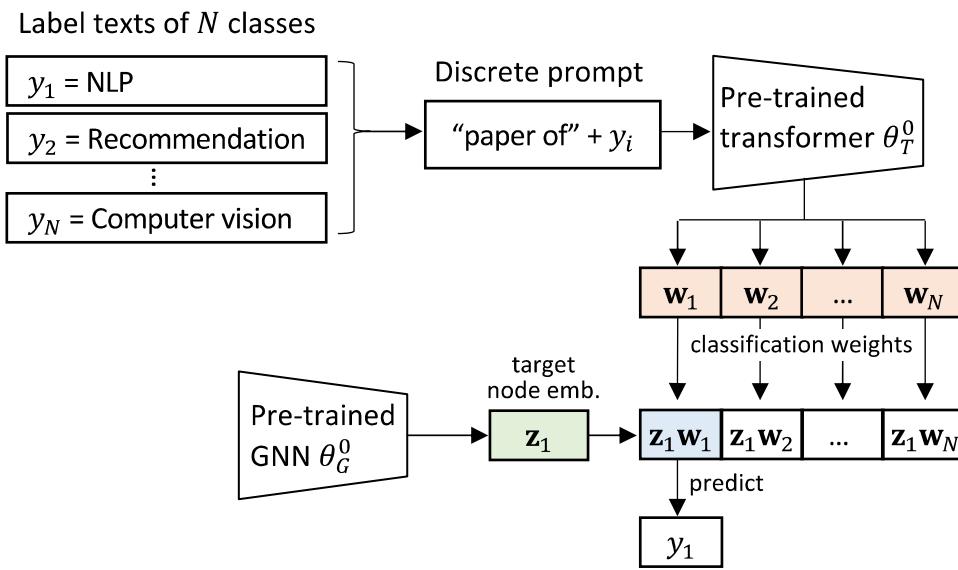
TABLE IX: Summary of prompt tuning on text-attributed graphs.

Paper	Instruction		Learnable prompt	Downstream Task		
	Text	Graph		Node	Edge	Graph
G2P2 [44]	✓	✗	vector	✓	✗	✗
G2P2* [82]	✓	✗	condition-net	✓	✗	✗
GraphGPT [170]	✓	✓	✗	✓	✗	✗
InstructGLM [175]	✓	✓	✗	✓	✗	✗
GIMLET [185]	✓	✓	✗	✗	✗	✓
OFA [45]	✓	✓	✗	✓	✓	✓
HiGPT [186]	✓	✓	✗	✓	✗	✗

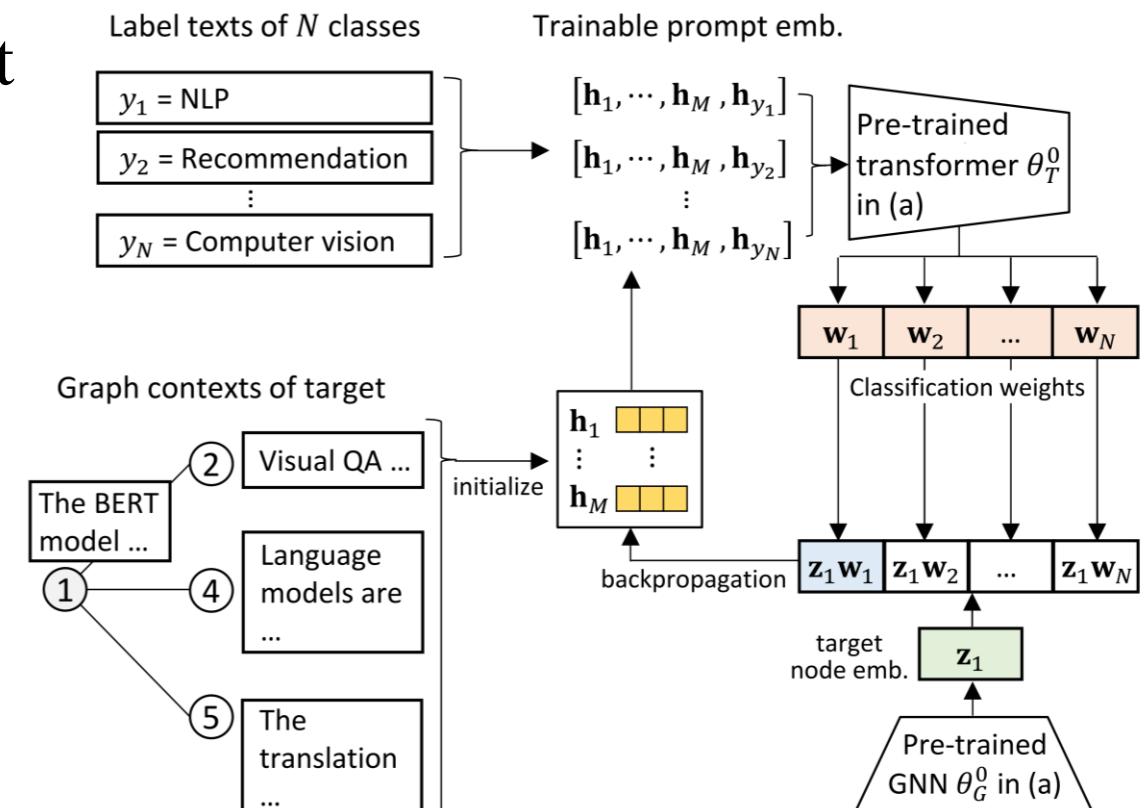
# Graph + PLM: Prompt-Tuning

- G2P2

- Discrete prompt-tuning: zero-shot
- Trainable prompt-tuning: few-shot



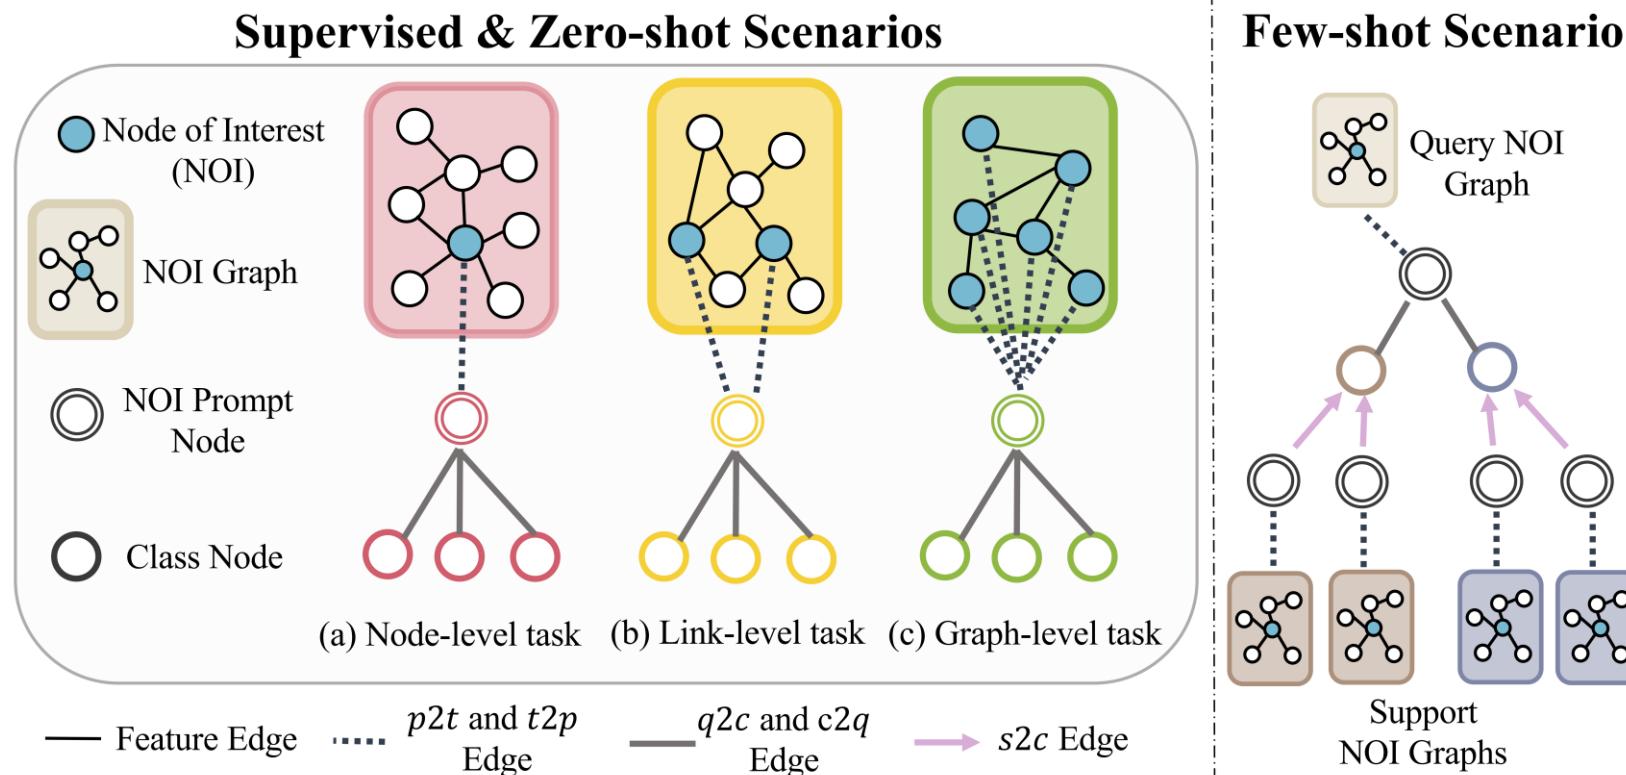
(a) Zero-shot



(b) Few-shot

# Graph + PLM: Prompt-Tuning

- One for all



# Graph + PLM: PEFT

- MolCA

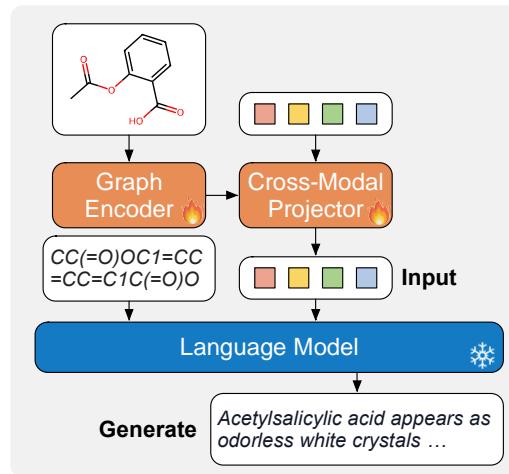


Figure 4: MolCA’s pretrain stage 2 by molecule captioning.

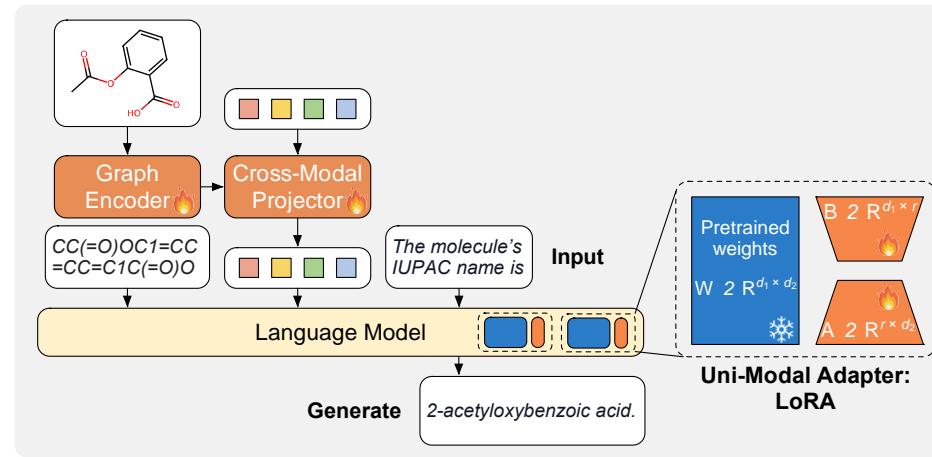


Figure 5: MolCA’s fine-tune stage for molecule-to-text generation. The example shows the prediction of a molecule’s IUPAC name.

Cross-Modal Projector: bridge the gap between graph structural and textual representations  
Uni-Modal LoRA Adapter: efficient downstream adaptation

# Graph + PLM: PEFT

- GraphGPT: only fintune projector  
➤ Align graph to LLM

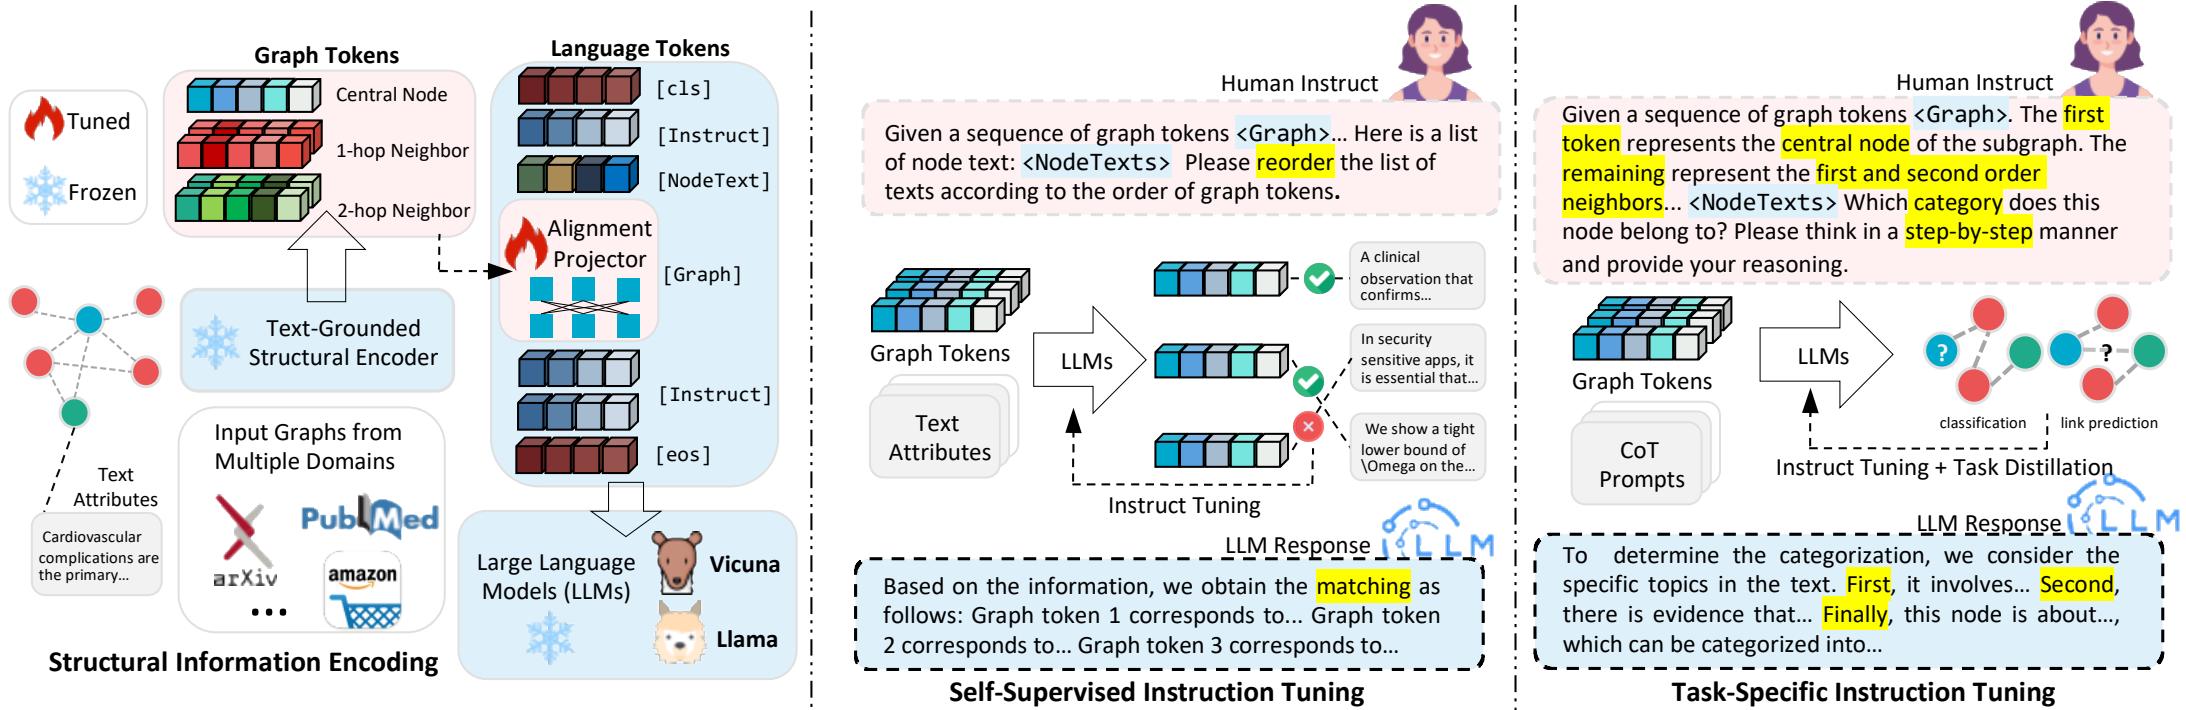


Figure 2: The overall architecture of our proposed GraphGPT with graph instruction tuning paradigm.

# Graph + PLM: PEFT

- GraphTranslator

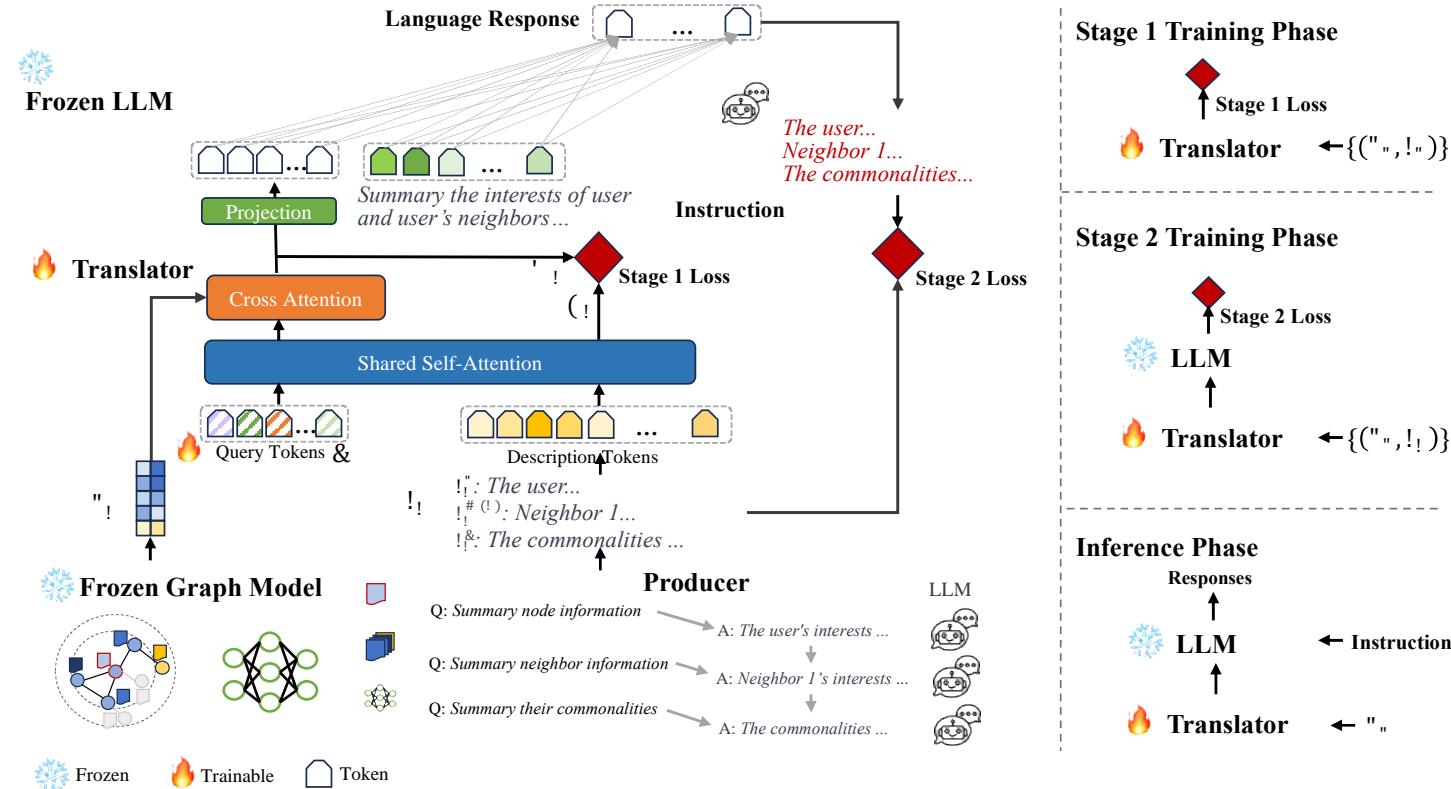


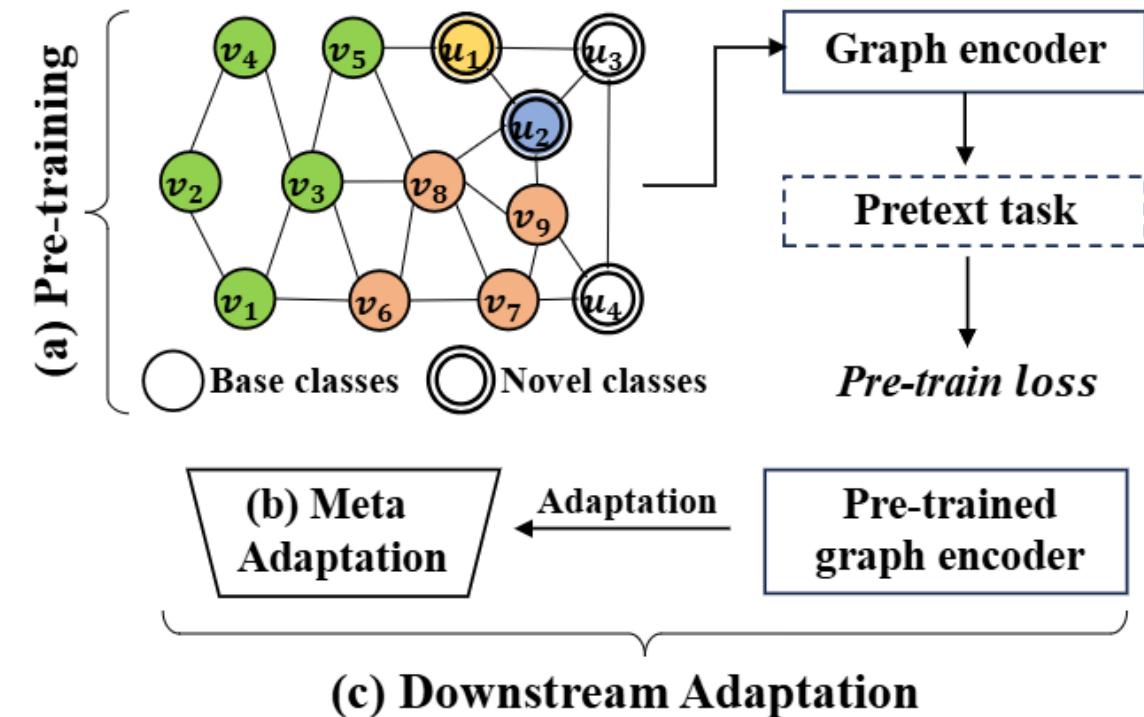
Figure 2: The overall framework of our *GraphTranslator*, which aligns GM to LLM by Translator for open-ended tasks. We train the lightweight Translator module following a two-stage paradigm, with the alignment data generated by our Producer.

# Summary

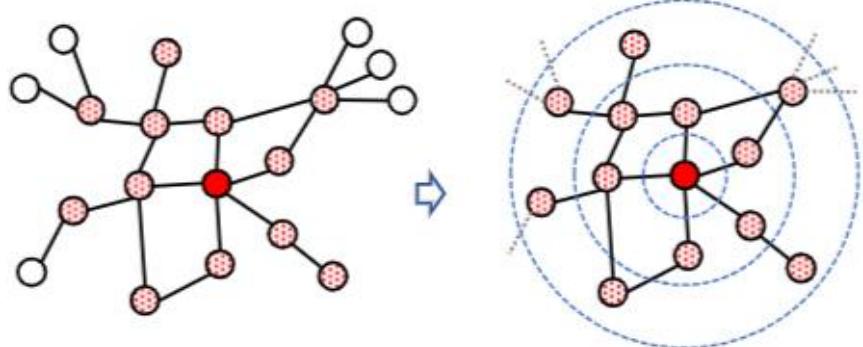
- Pre-training approaches employ self-supervised pretext tasks on unlabeled data
- Pre-training approaches are more effective in scenarios where labeled data are limited to novel tasks without a pre-existing set of annotated tasks.
- When a large annotated base set is available, meta-learning tends to perform better as it can leverage related meta-training tasks derived from the base set.
- Parameter-efficient adaptation strategies, including prompt tuning, adapter tuning and LoRA, present a more promising direction for few-shot learning on graphs.

# Hybrid Approaches

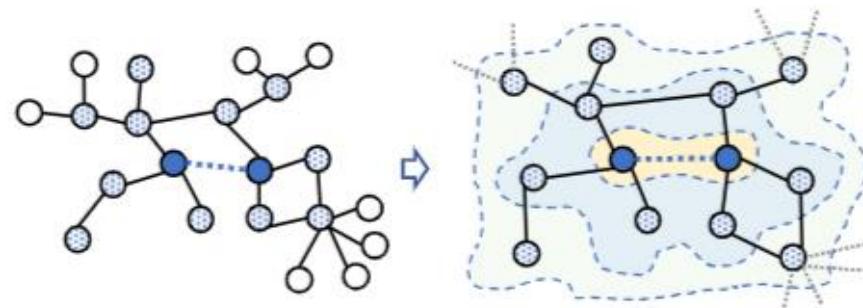
- Adopt pretext tasks to pre-train a graph encoder
  - Unlabeled data for pre-training
- The pre-trained model is adapted in conjunction with meta-learning
  - Annotated base set for meta-learning



# ProG



(a) Induced graphs for nodes



(b) Induced graphs for edges

Unify node level and edge level tasks as graph level tasks

## Prompt graph

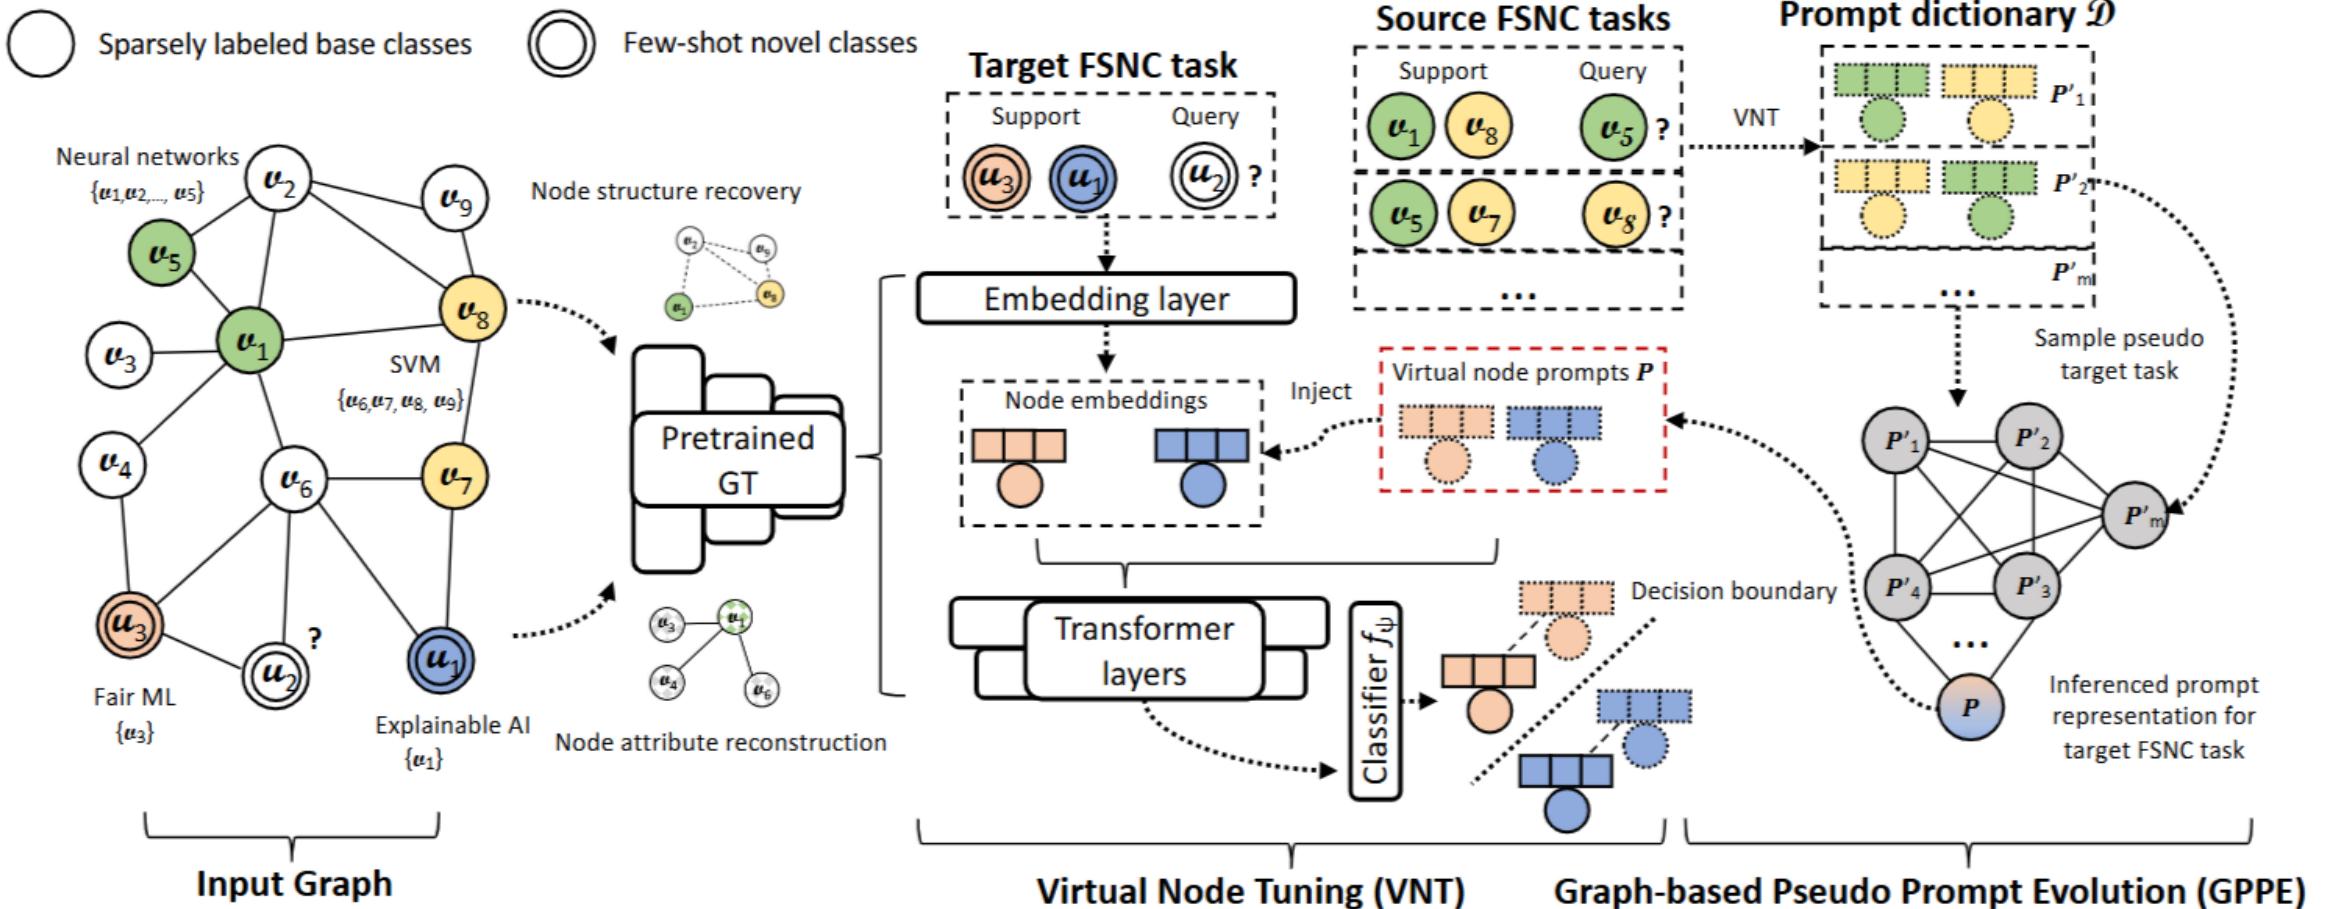
$$\mathcal{G}_P = (\mathcal{P}, \mathcal{S}) \quad \mathcal{P} = \{p_1, p_2, \dots, p_{|\mathcal{P}|}\}$$

## Prompt modification

$$\hat{\mathbf{x}}_i = \mathbf{x}_i + \sum_{k=1}^{|\mathcal{P}|} w_{ik} \mathbf{p}_k$$

$$w_{ik} = \begin{cases} \sigma(\mathbf{p}_k \cdot \mathbf{x}_i^T), & \text{if } \sigma(\mathbf{p}_k \cdot \mathbf{x}_i^T) > \delta \\ 0, & \text{otherwise} \end{cases}$$

First pre-train a graph encoder, then apply meta-learning to the prompting phase



First pre-train a graph transformer

Prompts

$$P = [p_1; \dots; p_p; \dots; p_P]$$

Prompt modification

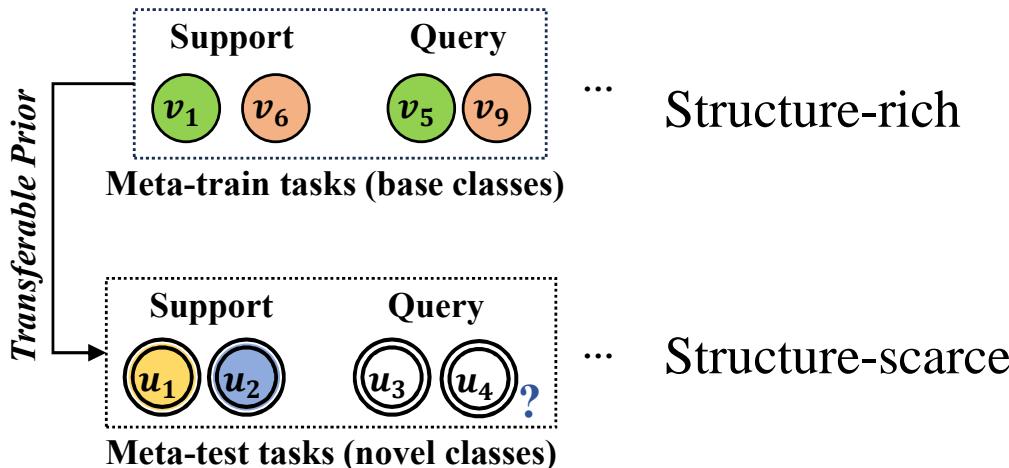
$$[E^1 || Z^1] = L^1([E^0 || P]) \in \mathbb{R}^{(V+P) \times F}$$

Tan, et al. "Virtual node tuning for few-shot node classification." SIGKDD'23.

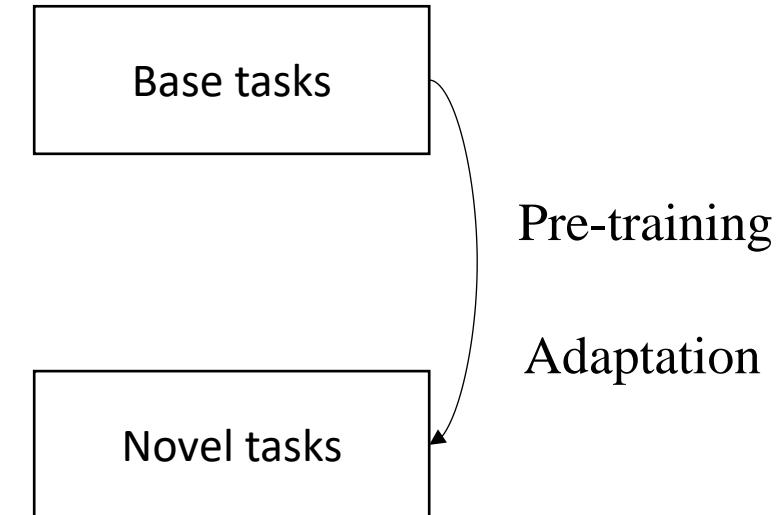
# Future Avenues in Problem Settings

- Structure scarcity learning on graphs

Existing solution:



Future direction:



Constrain:

Independent and Identically Distributed (i.i.d.)

Bridge the gap between base and novel tasks

# Future Avenues in Problem Settings

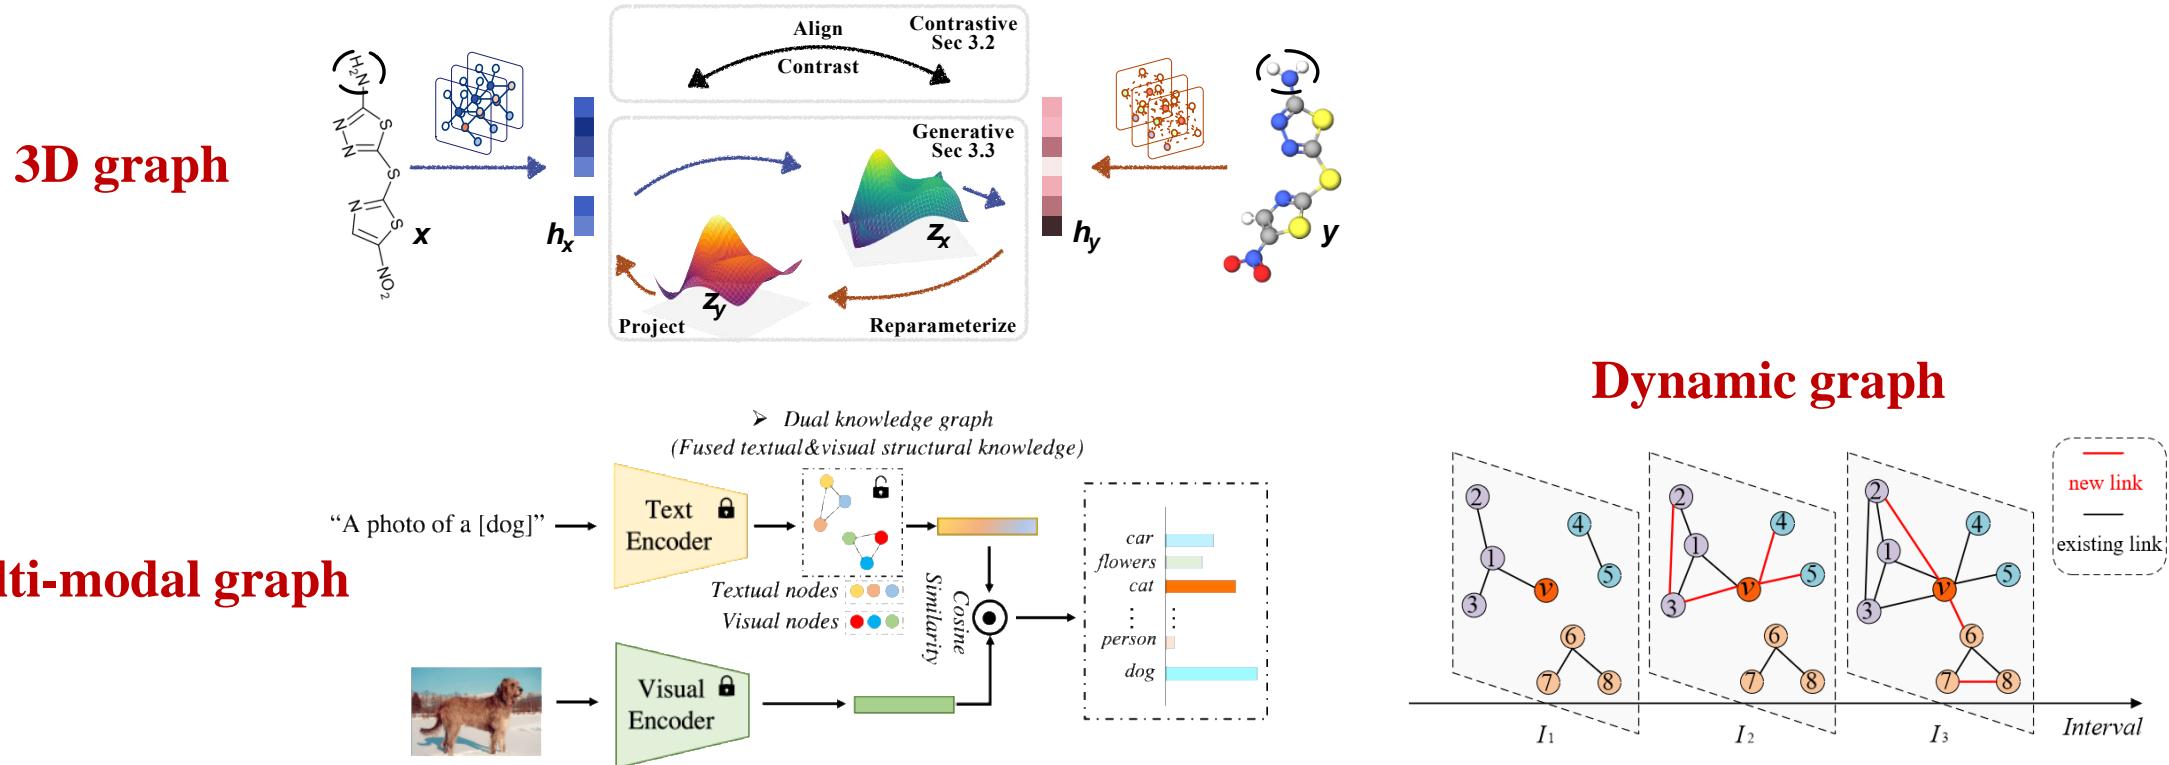
- Few-shot learning on large-scale graphs

Category	Name	Scale	#Graphs	Average #Nodes	Average #Edges	Average Node Deg.	Average Clust. Coeff.	MaxSCC Ratio	Graph Diameter
Node ogbn-	products	medium	1	2,449,029	61,859,140	50.5	0.411	0.974	27
	proteins	medium	1	132,534	39,561,252	597.0	0.280	1.000	9
	arxiv	small	1	169,343	1,166,243	13.7	0.226	1.000	23
	papers100M	large	1	111,059,956	1,615,685,872	29.1	0.085	1.000	25
	mag	medium	1	1,939,743	25,582,108	21.7	0.098	1.000	6
Link ogbl-	ppa	medium	1	576,289	30,326,273	73.7	0.223	0.999	14
	collab	small	1	235,868	1,285,465	8.2	0.729	0.987	22
	ddi	small	1	4,267	1,334,889	500.5	0.514	1.000	5
	citation	medium	1	2,927,963	30,561,187	20.7	0.178	0.996	21
	wikikg	medium	1	2,500,604	17,137,181	12.2	0.168	1.000	26
	biokg	small	1	93,773	5,088,434	47.5	0.409	0.999	8
Graph ogbg-	molhiv	small	41,127	25.5	27.5	2.2	0.002	0.993	12.0
	molpcba	medium	437,929	26.0	28.1	2.2	0.002	0.999	13.6
	ppa	medium	158,100	243.4	2,266.1	18.3	0.513	1.000	4.8
	code	medium	452,741	125.2	124.2	2.0	0.0	1.000	13.5
Dataset	Nodes	Edges	Classes	splitting (Train/Validation/Test)			Task		
Flickr	89,250	899,756	7	0.50 / 0.25 / 0.25			Multi-Class Classification		
Reddit	232,965	11,606,919	41	0.66 / 0.10 / 0.24			Multi-Class Classification		
ogbn-products	2,449,029	61,859,140	47	0.10 / 0.02 / 0.88			Multi-Class Classification		

Challenges: Finer-grained adaptation strategies to deal with potential variations among distant localities on a large graph.

# Future Avenues in Problem Settings

- Few-shot learning on complex graphs



S. Liu, et al. "Pre-training Molecular Graph Representation with 3D Geometry." ICLR'22

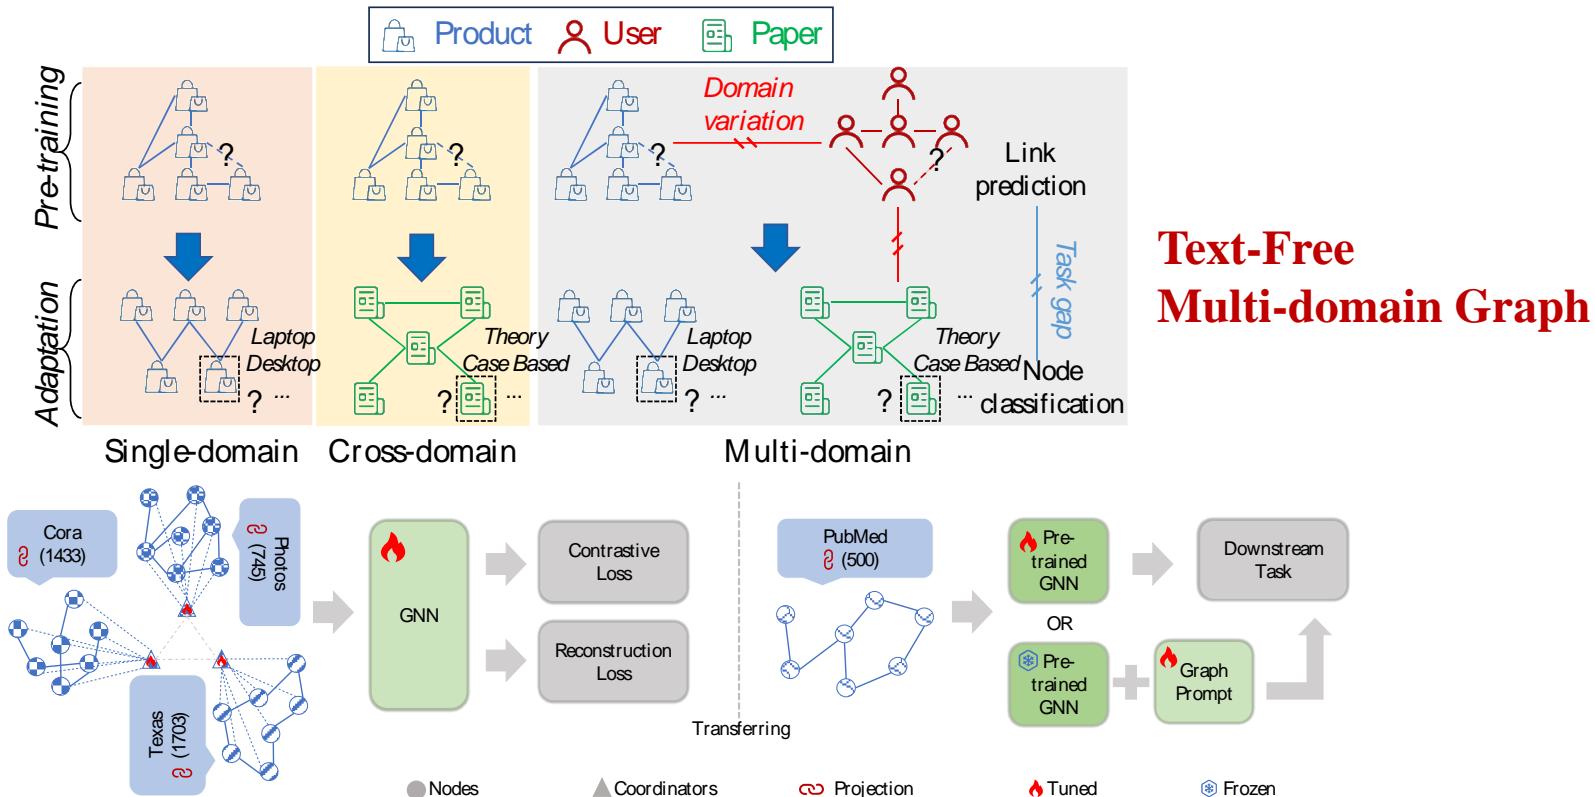
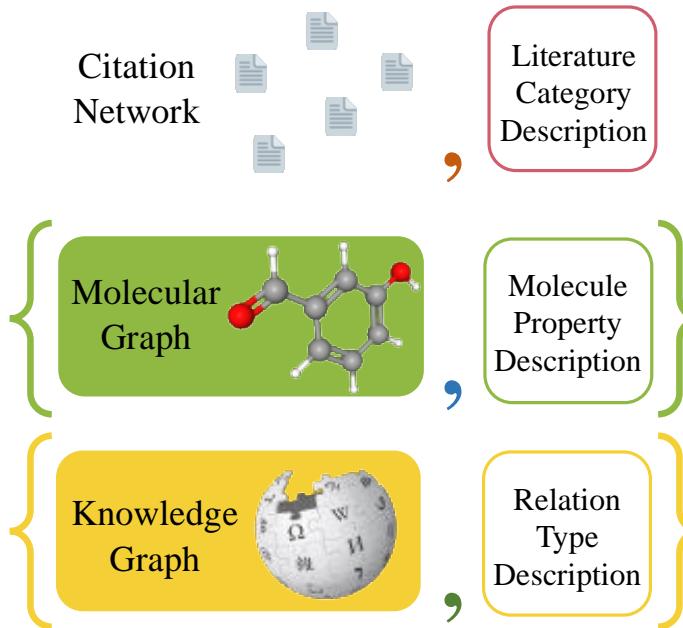
X. Li, et al. "GraphAdapter: Tuning Vision-Language Models With Dual Knowledge Graph." NeurIPS'23

C. Yang, et al. "Few-shot Link Prediction in Dynamic Networks." WSDM'22

# Future Avenues in Problem Settings

- Few-shot learning on cross-domain graphs

## Text-Attribute Multi-domain Graph



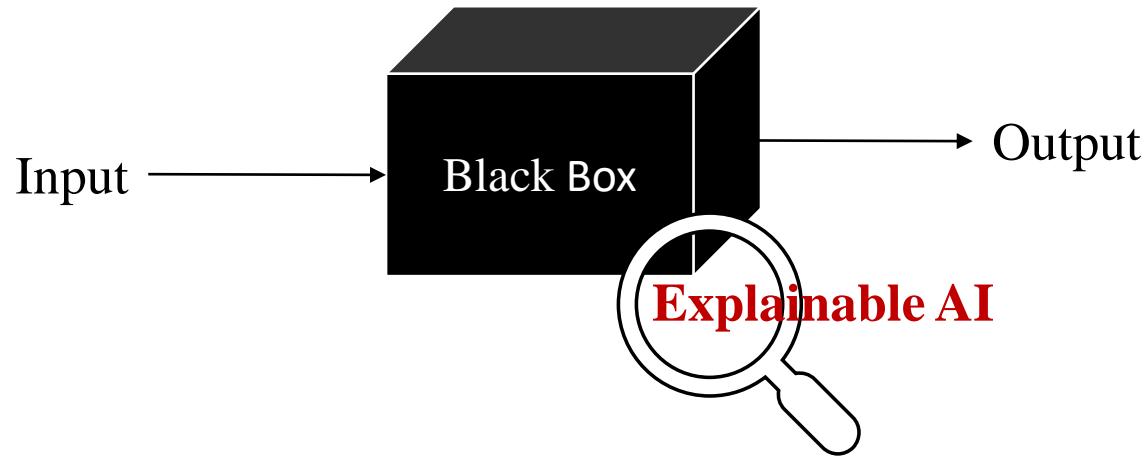
H. Liu, et al. “One for All: Towards Training One Graph Model for All Classification Tasks.” ICLR’24

X. Yu, et al. “Text-Free Multi-domain Graph Pre-training: Toward Graph Foundation Models.” ArXiv 2024

H. Zhao, et al, “All in One and One for All: A Simple yet Effective Method towards Cross-domain Graph Pretraining.” KDD’24

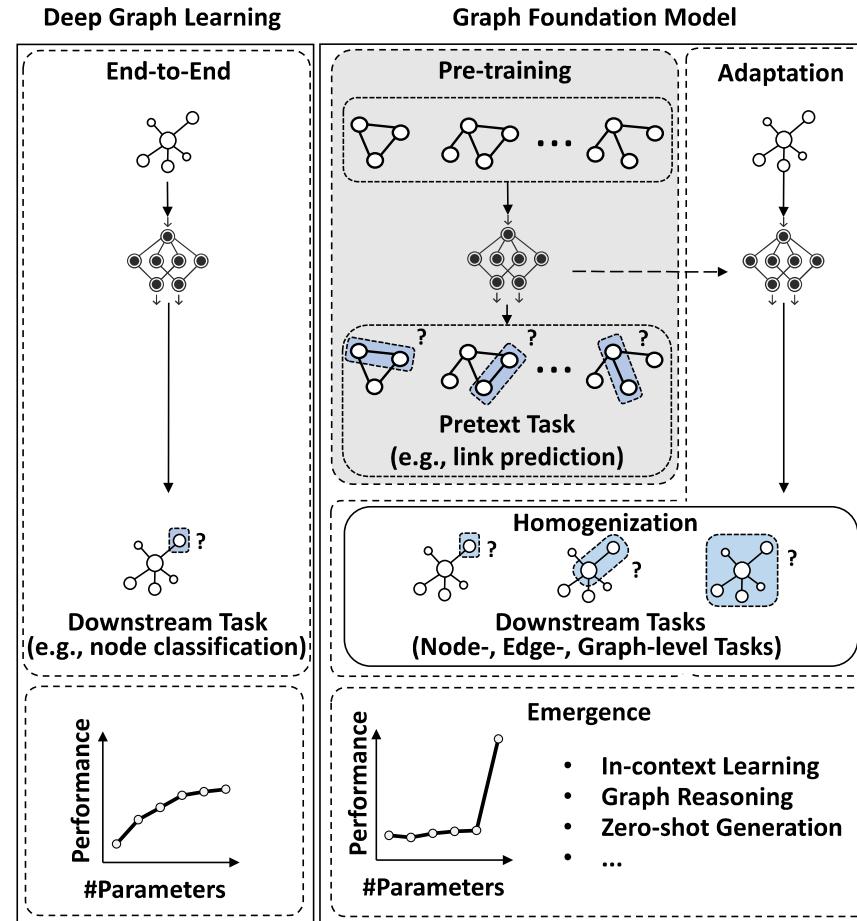
# Future Avenues in Techniques

- Improving interpretability



# Future Avenues in Techniques

- Foundation models on graphs



J. Liu, et al. "Graph Foundation Models: Concepts, Opportunities and Challenges." TPAMI'25